



Reflections on Programming Environments and Productivity (based on experiences with HPCS and Chapel)

Brad Chamberlain, Chapel Team, Cray Inc.
ASCR Exascale Computing Systems Productivity Workshop
June 3rd, 2014



COMPUTE

STORE

ANALYZE

Safe Harbor Statement

This presentation may contain forward-looking statements that are based on our current expectations. Forward looking statements may include statements about our financial guidance and expected operating results, our opportunities and future potential, our product development and new product introduction plans, our ability to expand and penetrate our addressable markets and other statements that are not historical facts. These statements are only predictions and actual results may materially vary from those projected. Please refer to Cray's documents filed with the SEC from time to time concerning factors that could affect the Company and these forward-looking statements.

Programming Environments

coupled closely
with applications

- **Programming Notations / Programming Models**

- languages, language extensions, DSLs, pragmas, libraries, ...

my bias

- **Tools**

- debuggers, profilers, autotuners, IDEs, ...

- **Runtime Support**

- communication, tasking, memory, I/O

coupled closely
with OS

Chapel (part of the reason for my bias)

- An emerging parallel programming language
 - Design and development led by Cray Inc.
 - in collaboration with academia, labs, industry
- Goal: Improve productivity of parallel programming
- A work-in-progress

Productivity: Traditional, pre-Exascale Concerns

What does “Productivity” mean to you?

Recent Graduates:

“something similar to what I used in school: Python, Matlab, Java, ...”

Seasoned HPC Programmers:

“that sugary stuff that I don’t need because I ~~was born to suffer~~
want full control
to ensure performance”

Computational Scientists:

“something that lets me express my parallel computations
without having to wrestle with architecture-specific details”

Chapel Team:

“something that lets computational scientists express what they want,
without taking away the control that HPC programmers want,
implemented in a language as attractive as recent graduates want.”

Productivity: My nightmare scenario

Scenario:

- A mainstream computing buddy wants to do some scalable parallelism
 - Accustomed to using Python, Matlab, or Java, say
 - Also IDEs with auto-completion, refactoring, integrated debugging, ...
- Knowing you're an expert in the field, wants recommendations

The source of my fears:

“As the HPC community, do we have anything we can recommend as a productive solution to such a person with a straight face?”

“Do any of us even recognize what productivity means to most programmers anymore? Would we know it if it bit us on the leg?”

How to attract/retain HPC programmers?



One Answer: “Decent” Parallel Languages

What was the last parallel notation you used that felt:


- productive?
- high-level?
- powerful?
- flexible?
- effective?
- modern?
- fun?
- (...all the things we judge good software by...)?




(Because that's what we're competing against when it comes to attracting new talent)

Productivity in the Exascale Era


Prototypical Exascale Processor Technologies




Intel MIC



AMD APU



Nvidia Echelon







Tilera Tile-Gx

http://download.intel.com/pressroom/images/Aubrey_Isle_die.jpg <http://www.zdnet.com/amds-trinity-processors-take-on-intels-ivy-bridge-3040155225/>

<http://insidehpc.com/2010/11/26/nvidia-reveals-details-of-echelon-gpu-designs-for-exascale/> <http://tilera.com/sites/default/files/productbriefs/Tile-Gx%203036%20SB012-01.pdf>

Exascale: Programmer Productivity Challenges



Emerging processor designs...

...are increasingly locality-sensitive

...potentially have multiple processor/memory types

- ⇒ Exascale programmers will have a lot more to think about at the node level than in the past
- ⇒ What will it take to keep the productivity bar level (to say nothing of improving it?)

Summarizing: Three Productivity Challenges

1. How to improve productivity relative to current practice?
2. How to improve productivity to entice new users?
3. How to maintain productivity in the face of exascale?
 - let alone improve it?

Happily, #3 gives us a renewed excuse to work on #'s 1 & 2:

If we're going to have to switch to something new anyway, it's a great opportunity to change to something truly productive

Productivity, HPCS, and Cray:

A Brief History/Review

Productivity, as defined by HPCS

Productivity (10x improvement goal) =

- performance
- + programmability
(readability, writeability, maintainability, modifiability, tunability, ...)
- + portability
- + robustness

A reasonable starting point...

...yet, how to combine four areas down to a single metric?

- particularly given that most of them are hard to measure individually?

Also some unreasonable (IMO) goals/expectations:

- Initially, a stated desire to see the establishment of Moore's Law-style productivity improvements year after year

Productivity: Played Out?

- **There's some sense that productivity isn't “hot” anymore**
 - “Didn't we [solve | fail to solve] that in HPCS?”
- **Arguably analogous to “peace”**
 - not particularly “new” or “hip” as a concept
 - reasonable causes for skepticism about our ability to achieve it
 - but clearly something to desire over the alternatives
- **Personally, I prefer not to throw in the towel (in either case)**

My “Zany Metrics” (a brainstorming exercise)



The DARPA logo is in the top left corner, featuring the acronym "DARPA" in white letters inside a blue oval with a globe graphic.

“Zany” Metrics



The HPCS logo is in the top right corner, with the letters "HPCS" in a stylized, handwritten font with a blue arrow pointing upwards to the right.

- ◆ **Abstractness of Code**
 - how much code must change if we...
 - ◆ change number of processors, shape of processor set?
 - ◆ change problem size?
 - ◆ make processors not divide problem size evenly?
 - ◆ make processor dimensions, problem size non- 2^k ?
 - ◆ switch dense arrays to sparse?
 - ◆ change an array's rank?
- ◆ **Portability of Code**
 - how much code must change...
 - ◆ to run on another vendor's machine?
 - ◆ to get performance satisfactory to that vendor?

“Language Bingo”

DARPA

Language Comparison


HPCS ↗

	MPI	SHMEM	Java	UPC	CAF	HPF	OpenMP	Fortran/C
Performs Well	O	O	?	?	O	?	?	~
Portable	O	?	O	?	?	~	X	~
Performance Model	O	O	O	O	O	X	X	X
Global View	X	X	X	X	X	O	O	O
Post-scalar	~\X	~\X	O	X	~	~	~\X	~\X
Abstractions	X	X	O	~	~	X	X	X
Succinct	X	X	X	X	X	~	~	~
General	O	O	O	~	X	X	X	O

-- = no comment O = good ~ = so-so X = poor ? = unproven

CRAY

Sterling's Model of Productivity



Sterling's Model of Productivity



General Model of Productivity

HPCS

R_i \equiv i^{th} result product

T_i \equiv time to compute result R_i

T_L \equiv total lifetime of machine

T_V \equiv total overhead time of machine

T_Q \equiv quiescent time of machine

T_R \equiv working time of machine

N_R \equiv total number of result products during T_L

C_L \equiv all costs associated with machine during T_L

C_{LS} \equiv application software costs during T_L

C_{LO} \equiv costs of ownership during T_L

C_M \equiv cost of procurement and initial installation

C_{Si} \equiv cost of application software for result R_i

Ψ \equiv productivity



...coincidence?!? I think not...

$$T_R = \sum_i^{N_R} T_i$$

$$T_L = T_R + T_V + T_Q$$

$$R_L = \sum_i^{N_R} R_i$$

$$C_L = C_{LS} + C_M + C_{LO}$$

$$C_{LS} = \sum_i^{N_R} C_{Si}$$

$$\Psi = \frac{R_L}{C_L \times T_L}$$

CRAY


Various Others' Models of Productivity

Sign In to gain access to subscriptions and/or My Tools.

[Sign In](#) | [My Tools](#) | [Contact Us](#) | [HELP](#)

SAGE journals  [Search all journals](#)  [Advanced Search >](#)  [Search History >](#)  [Browse Journals >](#)

International Journal of HIGH PERFORMANCE COMPUTING APPLICATIONS



[Home](#) [OnlineFirst](#) [All Issues](#) [Subscribe](#) [RSS !\[\]\(b65a4359c057803ec5d5456d155c4974_img.jpg\)](#) [Email Alerts](#)

[« Return to Search Results | Edit My Last Search](#)  [Search this journal](#)  [Advanced Journal Search »](#)

Impact Factor: 1.295 | **Ranking:** 15/50 in Computer Science, Hardware & Architecture | 27/100 in Computer Science, Theory & Methods | 53/99 in Computer Science, Interdisciplinary Applications

Source: 2012 Journal Citation Reports® (Thomson Reuters, 2013)

Table of Contents

Winter 2004; 18 (4)

Articles

- Jeremy Kepner
HPC Productivity: An Overarching View
 International Journal of High Performance Computing Applications Winter 2004
 18: 393-397, doi:10.1177/1094342004048533
[Abstract](#) [Full Text \(PDF\) !\[\]\(adc4b29c7532d08fefd48c8d44095417_img.jpg\)](#) [Request Permissions](#)
- D. E. Post and R. P. Kendall
Software Project Management and Quality Engineering Practices for Complex, Coupled Multiphysics, Massively Parallel Computational Simulations: Lessons Learned From ASCI
 International Journal of High Performance Computing Applications Winter 2004
 18: 399-416, doi:10.1177/1094342004048534
[Abstract](#) [Full Text \(PDF\) !\[\]\(54768eee398db0160ed86d2944893c26_img.jpg\)](#) [References](#) [Request Permissions](#)
- Marc Snir and David A. Bader
A Framework for Measuring Supercomputer Productivity
 International Journal of High Performance Computing Applications Winter 2004
 18: 417-429, doi:10.1177/1094342004048535
[Abstract](#) [Full Text \(PDF\) !\[\]\(cde76982820855aff29d029b89d231fa_img.jpg\)](#) [References](#) [Request Permissions](#)

[« Previous | Next Issue »](#)

This Issue

Winter 2004; 18 (4)



[Index By Author](#)

[Articles](#)

Find articles in this issue containing these words:
 [Search Issue](#) [Search](#)

[Submit a Manuscript](#)


[Free Sample Copy](#)

[Email Alerts](#)

[RSS feed](#)

[More about this journal](#)

The Application Kernel Matrix



The Application Kernel Matrix Website (R.I.P.)

Cascade: Application Kernel Matrix

info the kernels the matrix programmer's log kernel submission form discussion forum

Kernel Specs & Solutions:

- NASPB Conjugate Gradient
- Sweep3D
- NASPB Unstructured Adaptive
- Connected Components
- Chip Floorplan Design
- NASPB Fourier Transform
- NASPB Multigrid Benchmark
- Protein Sequence Matching
- Sparse Matrix Triangular Backsolve
- Vector Max and Prefix Sums

Links:

- Cray, Inc.
- The Cascade Project
- HPCS: High Productivity Computing Systems program
- DARPA: Defense Advance Research Projects Agency

Contacts:

- David Mizell
- John Feo
- John Lewis
- Brad Chamberlain
- Justin Garcia

Kernel Matrix:

The matrix is a graphical representation of all the submissions that we have received and confirmed. Programmers can submit a "generic" solution, or one that is tuned for high performance on a specific computer system. If you submit a kernel solution in a language that hasn't been used in the matrix before, a new row gets added to the matrix. Hover over a row, column, or cell for more information about already-submitted solutions.

Select a target system:		Generic Implementation							
CG	S3D	UA	CCG	CFD	NFT	NMG	PSM	SMB	VMP
Fortran			1						
Unified Parallel C									
Chapel									
ZPL					1				

Most recent submission: January 7 2005 @ 19:27:39

Submitted August 17 2004 @ 15:34:39
Submitted by: Justin Garcia of Rice University
Kernel: Connected Components
Language: Fortran
Line count: 1000
Token count: 1000
Execution time: 1.2 seconds

Compiled with `gcc -O3` on a Powerbook G4 running Mac OS 10.3.

[Download the source code](#)

Last Modified: October 27 2004 03:51:26 PM

W3C HTML W3C CSS


(While the AKM did not catch on, other similar comparison sites have, at least in the mainstream

One of my favorites (content and form):


- The Computer Language Benchmarks Game
<http://benchmarksgame.alioth.debian.org/>

The HPC Challenge, Berkeley Dwarfs, and Graph-500 arguably play similar roles).


HPCS Workflows



Timed Markov Models



User Studies: Quantitative Evaluation



User Studies: Qualitative Evaluation

“The biggest feature from a broad perspective for me was domains. Especially for scientific codes, it is invaluable to be able to define the couple problem domains you're working with. It makes it trivial to change the size or layout or distribution if you decide you need to, it helps guarantee that all of your different arrays match up. **A 3D rectangular grid is infinitely more clear in Chapel with domains than in C**, where you have to figure out how they laid it out (is it one giant array? what is the major dimension? x? z? y?).”

“I loved not having to think as hard about offsets and counts for the parallel version of the code in Chapel, as opposed to the MPI version, where I almost always had to chase down two or three indexing errors.”

“Lastly, I'm a huge huge fan of the type inference used in Chapel. I like that I don't have to specify types everywhere--they can just be inferred from how I'm using them, but if I mess something up, the compiler catches it.”

Summary: Many Useful Concepts/Techniques...

BUT...

The Catch with Productivity Metrics

My suspicion: If I were to tell you that any of these metrics...

...demonstrated that Chapel was 10x better than Fortran+MPI

...showed that X10 was 2x better than Chapel

...indicated that Perl was 5x better than Python

...proved that emacs was 8.5x better than vim

...dared suggest that alpine is 7x better than Outlook

...would you believe me? Enough to change your practices?

Maxim #1: “I can’t define productivity, but I know it when I see it”

- “And seeing is believing”

Maxim #2: “Productivity is in the eye of the beholder”


- “To each his/her own productivity solutions”

If Not Metrics, Then What?

For many of us, case studies can be compelling...

Here are three examples from Chapel

Case 1: LULESH in Chapel



COMPUTE

STORE

ANALYZE

LULESH in Chapel

1288 lines of source code

plus 266 lines of comments
 487 blank lines

(the corresponding C+MPI+OpenMP version is nearly 4x bigger)

This can be found in Chapel v1.9 in examples/benchmarks/lulesh/*.chpl

LULESCH in Chapel

This is the only representation-dependent code.
It specifies:

- data structure choices
 - structured vs. unstructured mesh
 - local vs. distributed data
 - sparse vs. dense materials arrays

A good poster child for my “zany metrics”

LULESH's Technical Productivity Lesson

- Put hardware-mapping-specific choices in declarations
 - Makes computation independent of key decisions like:
 - memory layouts
 - distributions
 - sparse vs. dense
 - # of dimensions
 - Supports switching between options easily
- This is a design trend that (happily) seems to be growing
 - Several similar designs reported on at PADAL 2014 workshop



Workshop on
Programming
Abstractions for Data
Locality (PADAL)

Workshop on Programming Abstractions for Data Locality (PADAL)

Abstracts
Logistics
Panels
Participants
Presentations
Registration
Schedule

The cost of data movement has become the dominant factor of a high performance computing system both in terms of energy consumption and performance. To minimize data movement, applications have to be optimized both for vertical data movement in the memory hierarchy and horizontal data movement between processing units. While microarchitectural technology trends allow the scaling of the number of cores per chip, cache coherence will likely not scale to the large number of cores due to the traffic overhead of maintaining coherence. In the future, software-managed memory and incoherent caches or scratchpad memory will be prevalent. Thus, application developers need a set of programming abstractions to describe data locality on the new computing ecosystems.

Architectural trends break our existing programming paradigm because the current software tools optimize for floating point operations not memory traffic. They ignore the incurred cost of



LULESH's Social Productivity Lesson

- Written by intern in final 2 weeks as a “bonus” project
 - productive!
- Illustrated Chapel use in a familiar setting to scientists
 - ones who'd heard many Chapel talks previously, and yet...
- Served as a medium for discussion, collaboration
- Demonstrated productivity of features thought not to be
 - global indexing, sparse domains
- Provided bidirectional feedback/knowledge transfer
 - co-design!


(a nice win for DOE's proxy apps effort)

Case 2: Chapel Rank-Independent AMR Framework

UW applied mathematician wrote one code that could be used to produce results in 2D, 3D, 6D, 17D...



Case 3: Chapel's Appeal to Educators/Students



http://prezi.com/wp13iqmsl1di/summer-scholars/?utm_campaign=share&utm_medium=copy

SIGCSE
ATLANTA
2014

OpenConf Peer Review & Conference Management System

[OpenConf Home](#) [Email Chair](#)

[Full Program »](#)

Chapel: A versatile tool for teaching undergraduates parallel programming

Chapel is a programming language being developed for high-performance applications. It is well-suited for teaching parallelism in a wide variety of undergrad courses. Chapel is easy to learn since it supports a low-overhead style like a scripting language as well as a full OO style. It is easier to learn than C or MPI because it does not require a developer to launch an asynchronous task, run a parallel loop, or perform a reduction.

<http://faculty.knox.edu/dbunde/teaching/chapel/SIGCSE14/>

March 5-8, 2014
Atlanta, Georgia

CSEP524: Parallel Computation

Software

Pthreads: (included with the above)

OpenMP: (included with the above)

Chapel: chapel-fedora17-1.6.1.1.tar.gz

(This is a pre-release of the Chapel 1.6.1 sources, pre-compiled for the Fedora 17 VM; and skipping step 2 if you're using the VM)

MPI: [Setup MPI](#)

This describes how to install MPI and how to run it locally and on our course VM cluster.

<http://courses.cs.washington.edu/courses/csep524/13wi/>

COMPUTE

STORE

ANALYZE

Chapel Wrap-up

- (Many other productive demonstrations in addition...)
- Chapel's Productivity Scorecard
 - Performance
 - Programmability
 - Portability
 - Robustness
- Work is ongoing to improve Chapel's weak areas
 - Productivity often requires long-term investment and patience

Summary / Takeaways

- **Productivity still matters**
 - even if the term is well-worn
- **Exascale brings new productivity challenges**
 - but also an opportunity to improve upon past approaches
- **Not convinced we can measure productivity**
 - nor that it matters whether or not we can
- **Proxy apps are a useful medium for productivity studies**
 - serve as a place to demonstrate productivity features
 - serve as meeting place for distinct communities
- **Educators and Students are a good resource**
 - can't stand in for today's experts, but may be tomorrow's
- **Productivity gains may not happen overnight**

Funding Productivity: Personal Opinions

- **Invest in software to compensate for hardware challenges**
- **Pursue a combination of evolution and revolution**
 - can't afford to just do one
- **Evaluate based on user (or prospective user) viewpoints**
 - computer science viewpoints are not necessarily as relevant
- **Must budget for evaluation/study of new technologies**
 - can't expect such studies to be spare time activities
 - requires time from experts, not simply novices
 - again, well-designed proxy apps can be useful here

Legal Disclaimer

Information in this document is provided in connection with Cray Inc. products. No license, express or implied, to any intellectual property rights is granted by this document.

Cray Inc. may make changes to specifications and product descriptions at any time, without notice.

All products, dates and figures specified are preliminary based on current expectations, and are subject to change without notice.

Cray hardware and software products may contain design defects or errors known as errata, which may cause the product to deviate from published specifications. Current characterized errata are available on request.

Cray uses codenames internally to identify products that are in development and not yet publically announced for release. Customers and other third parties are not authorized by Cray Inc. to use codenames in advertising, promotion or marketing and any use of Cray Inc. internal codenames is at the sole risk of the user.

Performance tests and ratings are measured using specific systems and/or components and reflect the approximate performance of Cray Inc. products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance.

The following are trademarks of Cray Inc. and are registered in the United States and other countries: CRAY and design, SONEXION, URIKA, and YARCDATA. The following are trademarks of Cray Inc.: ACE, APPRENTICE2, CHAPEL, CLUSTER CONNECT, CRAYPAT, CRAYPORT, ECOPHLEX, LIBSCI, NODEKARE, THREADSTORM. The following system family marks, and associated model number marks, are trademarks of Cray Inc.: CS, CX, XC, XE, XK, XMT, and XT. The registered trademark LINUX is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a worldwide basis. Other trademarks used in this document are the property of their respective owners.

Copyright 2014 Cray Inc.



CRAY
THE SUPERCOMPUTER COMPANY

<http://chapel.cray.com>

chapel_info@cray.com

<http://sourceforge.net/projects/chapel/>