

# LLVM-based Communication Optimizations for Chapel

Chapel Lightning Talks BoF session at SC '14,  
New Orleans, LA

Akihiro Hayashi, Jisheng Zhao (Rice University)  
Michael Ferguson (Laboratory for Telecommunication Sciences  
Vivek Sarkar(Rice University)



# A Big Picture



Habanero-C, ...



©Oak Ridge National Lab.



© Argonne National Lab.





© RIKEN AICS

Pictures borrowed from <http://chapel.cray.com/logo.html>, <http://llvm.org/Logo.html>, <http://upc.lbl.gov/>, <http://commons.wikimedia.org/>, <https://www.olcf.ornl.gov/titan/>



# LLVM-based Chapel compiler



- Use of address space feature of LLVM offers more opportunities for **communication optimization** than C generation

```
// C-Code generation  
chpl_comm_get(&x, ...);
```

Backend Compiler's Optimizations  
(e.g. gcc -O3)

```
// Chapel  
x = possibly_remoteData;
```

```
// LLVM IR Generation  
%x = load i64 addrspace(100)* %xptr
```

LLVM Optimizations  
(e.g. LICM, scalar replacement)



Pictures borrowed from 1) <http://chapel.cray.com/logo.html>  
2) <http://llvm.org/Logo.html>



# An optimization Example : Communication Optimization with the existing LLVM passes

(Pseudo-Code: Before LICM)

```
for i in 1..N {  
    // POSSIBLY REMOTE GET  
%x = load i64 addrspace(100)* %xptr  
    A(i) = %x;  
}
```

LICM by  
LLVM

Remote data access per  
each iteration

(Pseudo-Code: After LICM)

```
// POSSIBLY REMOTE GET  
%x = load i64 addrspace(100)* %xptr  
for i in 1..N {  
    A(i) = %x;  
}
```

Hoisted out of the loop!



# An optimization Example : Bulk Transformation (Coalescing)

(Pseudo-Code: Before Bulk Transformation)

```
for i in 1..N {  
    // POSSIBLY REMOTE GET  
    ... = A(i);
```

```
}
```

Bulk  
Transformation

Remote array access per  
each iteration

(Pseudo-Code: After LICM)

```
var localA: [1..N] int;  
localA = A; // Bulk Transfer  
for i in 1..N {  
    ... = localA(i);  
}
```

Create Local  
Buffer &  
Perform bulk transfer

Converted to Definitely-  
Local Access!



# An Optimization Example : Locality Inference for avoiding runtime affinity checking

```
proc habanero(ref x, ref y, ref z) {  
    var p: int = 0;  
var A:[1..N] int;  
    if (x == 0) {  
        p = y;  
    } else {  
        local { p = z; }  
    }  
z = A(0) + z;  
}
```


A is definitely-local

p and z are  
definitely local

Definitely-local access!  
(avoid runtime affinity checking)




# Results on Cray XC-30 (LLVM-unopt vs. LLVM-allopt)



*4.6x performance improvement on average  
(6 applications, 1, 2, 4, 8, 16, 32, 64 locales)*



# Results on Westmere Cluster (LLVM-unopt vs. LLVM-allopt)



*4.4x performance improvement on average  
(6 applications, 1, 2, 4, 8, 16, 32 locales)*



# Conclusions


- LLVM-based Communication optimizations for Chapel
- Preliminary Evaluation with 6 applications
  - Cray-XC30 Supercomputer
    - 4.6x average performance improvement
  - Westmere Cluster
    - 4.4x average performance improvement
- Future Work
  - Extend for other languages



# Backup slides



# LLVM IR Generation from Chapel



# LLVM-based Communication Optimizations for Chapel

- 1) Wide pointer optimization (`--llvm-wide-opt`)
  - Utilize the existing optimization passes such as loop invariant code motion for the purpose of communication optimization (The Existing LLVM Passes)
  - Combine sequences of loads/stores on adjacent memory locations into a single memcpy (Aggregation Pass)
- 2) *Bulk Transformation (Coalescing data accesses)*
  - Create locale-local buffer
  - Insert bulktransfer and replace remote accesses with local buffer access
- 3) *Locality optimization (Locality-Inference)*
  - Transform *possibly-remote* access to *definitely-local* access at compile-time to avoid runtime affinity checking



# Performance Evaluations: Platforms

- Cray-XC30 Supercomputer @ NERSC
  - Per Node information
    - Intel Xeon E5-2695 @2.40GHz x 24 cores
    - 64GB of RAM
  - Interconnect
    - Cray Aries interconnect with Dragonfly topology
- Westmere Cluster @ Rice
  - Per Node information
    - Intel Xeon CPU X5660@2.80GHz x 12 cores
    - 48GB of RAM
  - Interconnect
    - Quad-data rated Infiniband
    - Mellanox FCA support



# Performance Evaluations: Details of Compiler & Runtime

## □ Compiler:

**Chapel version 1.9.0.23154 (Apr. 2014)**

- LLVM 3.3

## □ Runtime:

- GASNet-1.22.0

- Cray-XC30 : aries
  - Westmere Cluster : ibv-conduit

- qthreads-1.10

- Cray-XC30 : 2 shepherds, 24 workers/shepherd
  - Westmere Cluster : 2 shepherds, 6 workers/shepherd




Benchmark	Comm Kind	<b>Cray XC-30</b>	
		LLVM-gopt	LLVM-allopt
<b>Smith-Waterman</b> Note : obtained with 18,560x19,200 input	LOCAL_GET	63.6%	75.5%
	REMOTE_GET	36.4%	36.7%
	LOCAL_PUT	58.0%	58.0%
	REMOTE_PUT	0.0%	0.0%
<b>Cholesky</b> Note : obtained with 2,000x2,000 input	LOCAL_GET	77.6%	87.9%
	REMOTE_GET	84.7%	99.8%
	LOCAL_PUT	10.3%	10.8%
	REMOTE_PUT	0.0%	0.0%
<b>NPB EP</b>	LOCAL_GET	58.6%	58.6%
	REMOTE_GET	39.7%	39.7%
	LOCAL_PUT	29.5%	58.8%
	REMOTE_PUT	0.0%	0.0%
<b>Sobel</b> Note : obtained with CLASS=B	LOCAL_GET	74.6%	95.2%
	REMOTE_GET	0.0%	0.0%
	LOCAL_PUT	35.8%	68.3%
	REMOTE_PUT	0.0%	0.0%
<b>SSCA2</b>	LOCAL_GET	55.6%	56.2%
	REMOTE_GET	60.9%	60.8%
	LOCAL_PUT	5.6%	3.8%
	REMOTE_PUT	0.0%	0.0%
<b>Stream-EP</b>	LOCAL_GET	70.6%	70.6%
	REMOTE_GET	35.7%	35.7%
	LOCAL_PUT	17.3%	17.3%
	REMOTE_PUT	0.0%	0.0%

**Table 3.** The amount of Chapel Comm APIs calls made by LLVM-gopt and LLVM-allopt relative to LLVM-unopt (Cray-XC30, 16 locales)



# Future Work: A compiler that can uniformly optimize PGAS Programs

- Extend LLVM IR to support parallel programs with PGAS and explicit task parallelism
  - Two parallel intermediate representations(PIR) as extensions to LLVM IR  
(Runtime-Independent, Runtime-Specific)



**RICE** e.g. Task Parallel Construct e.g. GASNet API

