



Communication Optimization for the Chapel Programming Language

Michael Ferguson

Cray Inc.

March 24, 2016



Safe Harbor Statement

This presentation may contain forward-looking statements that are based on our current expectations. Forward looking statements may include statements about our financial guidance and expected operating results, our opportunities and future potential, our product development and new product introduction plans, our ability to expand and penetrate our addressable markets and other statements that are not historical facts. These statements are only predictions and actual results may materially vary from those projected. Please refer to Cray's documents filed with the SEC from time to time concerning factors that could affect the Company and these forward-looking statements.



Talk Outline

- 15 m **Chapel Background**
- 5 m **Communication Optimization Motivation**
- 5 m **Memory Consistency Models constrain optimization**
- 15 m **Sequential Consistency for Data Race Free Programs**
- 5 m **Optimizing communication with a cache for remote data**
- 5 m **Using LLVM to optimize communication**



What is Chapel?



COMPUTE

|

STORE

|

ANALYZE

Chapel's Origins: HPCS

DARPA HPCS: High Productivity Computing Systems

- **Goal:** improve productivity by a factor of 10x
- **Timeframe:** summer 2002 – fall 2012
- Cray developed a new system architecture, network, software, ...
 - this became the very successful Cray XC30™ Supercomputer Series



...and a new programming language: Chapel



Chapel Motivation

**Q: Why doesn't parallel programming have an equivalent to
Python / Matlab / Java / C++ / (your favorite programming language here) ?**

- one that makes it easy to quickly get codes up and running
- one that is portable across system architectures and scales
- one that bridges the HPC, data analysis, and mainstream communities

A: We believe this is due not to any particular technical challenge, but rather a lack of sufficient...

- ...long-term efforts
- ...resources
- ...community will
- ...co-design between developers and users
- ...patience

Chapel is our attempt to change this

Chapel's Implementation

- **Being developed as open source at GitHub**
 - Licensed as Apache v2.0 software
- **Portable design and implementation, targeting:**
 - multicore desktops and laptops
 - commodity clusters and the cloud
 - HPC systems from Cray and other vendors
 - *in-progress:* manycore processors, CPU+accelerator hybrids, ...



Chapel is a Collaborative, Community Effort



Colorado
State
University



RICE®



ETH Zürich



Lawrence Berkeley
National Laboratory



(and many others as well...)



<http://chapel.cray.com/collaborations.html>

COMPUTE

| STORE |

ANALYZE

Sustained Performance Milestones

1 GF – 1988: Cray Y-MP; 8 Processors

- Static finite element analysis
- Fortran77 + Cray autotasking + vectorization



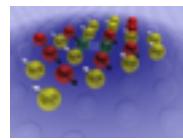
1 TF – 1998: Cray T3E; 1,024 Processors

- Modeling of metallic magnet atoms
- Fortran + MPI (Message Passing Interface)



1 PF – 2008: Cray XT5; 150,000 Processors

- Superconductive materials
- C++/Fortran + MPI + vectorization



1 EF – ~20__: Cray ____; ~10,000,000 Processors

- TBD
- TBD: C/C++/Fortran + MPI + OpenMP/OpenACC/CUDA/OpenCL?

Or, perhaps something completely different?

COMPUTE

STORE

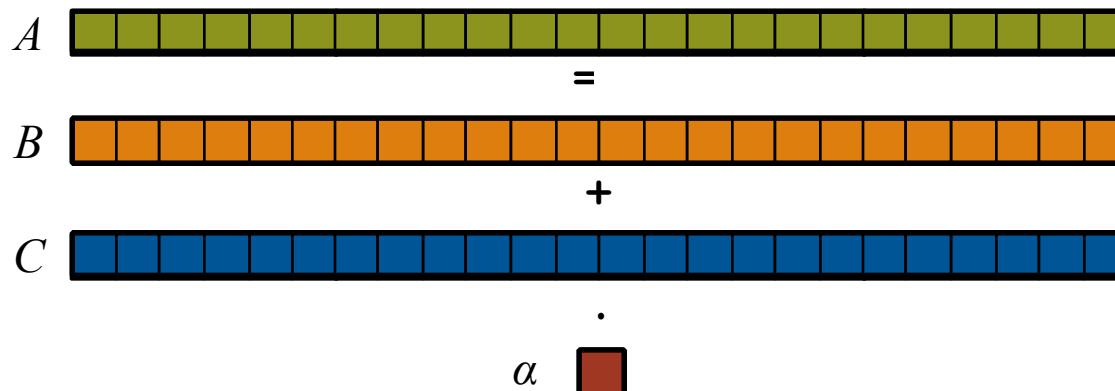
ANALYZE

STREAM Triad: a trivial parallel computation

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures:

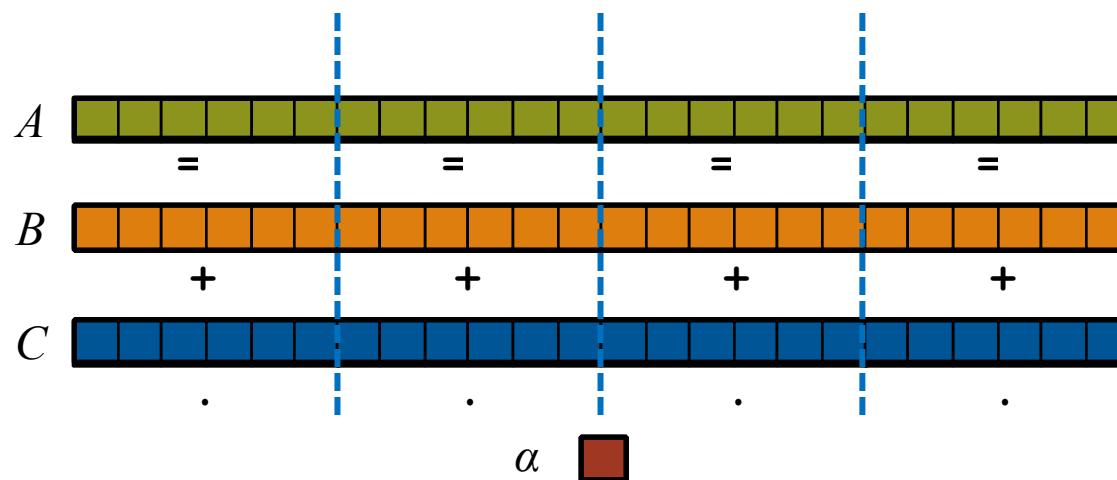


STREAM Triad: a trivial parallel computation

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel:

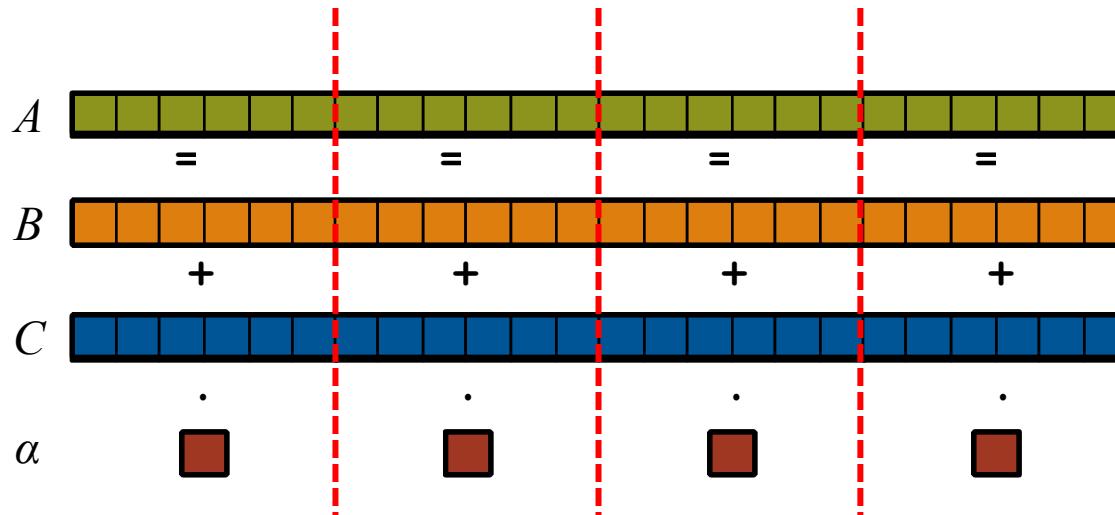


STREAM Triad: a trivial parallel computation

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel (distributed memory):

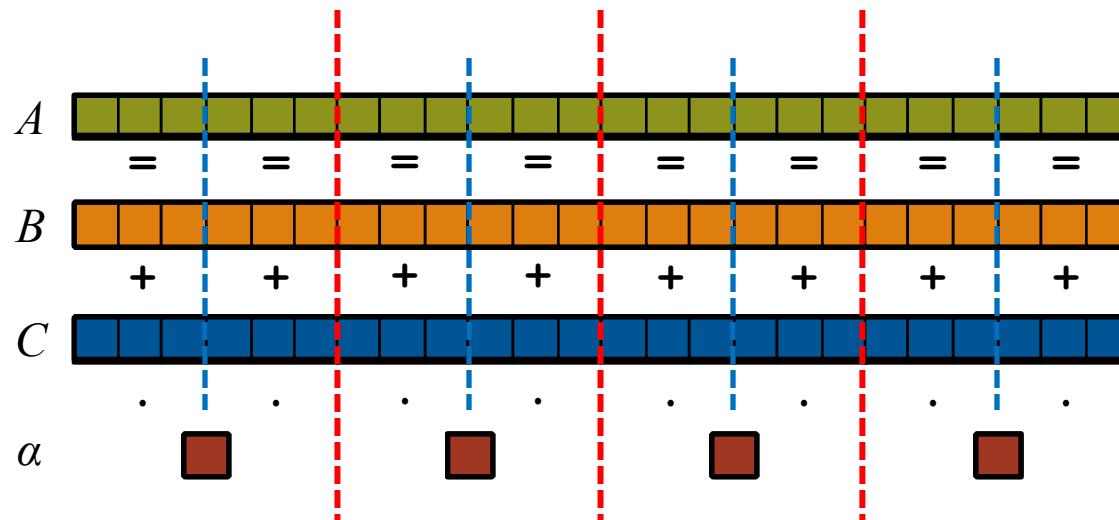


STREAM Triad: a trivial parallel computation

Given: m -element vectors A, B, C

Compute: $\forall i \in 1..m, A_i = B_i + \alpha \cdot C_i$

In pictures, in parallel (distributed memory multicore):



STREAM Triad: MPI



MPI

```
#include <hpcc.h>

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Parms *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0,
                comm );
}

return errCount;
}

int HPCC_Stream(HPCC_Parms *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3,
                                       sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );
```

```
if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory (%d). \n",
                 VectorSize );
        fclose( outFile );
    }
    return 1;
}

for (j=0; j<VectorSize; j++) {
    b[j] = 2.0;
    c[j] = 1.0;
}

scalar = 3.0;

for (j=0; j<VectorSize; j++)
    a[j] = b[j]+scalar*c[j];

HPCC_free(c);
HPCC_free(b);
HPCC_free(a);
```



COMPUTE

|

STORE

|

ANALYZE



STREAM Triad: MPI+OpenMP

MPI + OpenMP

```
#include <hpcc.h>
#ifndef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Parms *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;

    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );

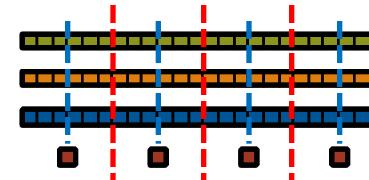
    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0,
                comm );
}

return errCount;
}

int HPCC_Stream(HPCC_Parms *params, int doIO) {
    register int j;
    double scalar;

    VectorSize = HPCC_LocalVectorSize( params, 3,
                                       sizeof(double), 0 );

    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );
```



```
if (!a || !b || !c) {
    if (c) HPCC_free(c);
    if (b) HPCC_free(b);
    if (a) HPCC_free(a);
    if (doIO) {
        fprintf( outFile, "Failed to allocate memory (%d). \n",
                 VectorSize );
        fclose( outFile );
    }
    return 1;
}

#ifndef _OPENMP
#pragma omp parallel for
#endif
for (j=0; j<VectorSize; j++) {
    b[j] = 2.0;
    c[j] = 1.0;
}

scalar = 3.0;

#ifndef _OPENMP
#pragma omp parallel for
#endif
for (j=0; j<VectorSize; j++)
    a[j] = b[j]+scalar*c[j];

HPCC_free(c);
HPCC_free(b);
HPCC_free(a);
```



COMPUTE

|

STORE

|

ANALYZE

STREAM Triad: MPI+OpenMP vs. CUDA

MPI + OpenMP

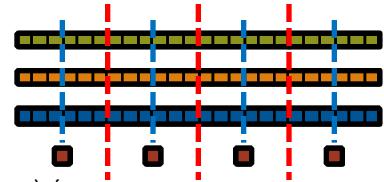
```
#ifdef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;
    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );
    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0, comm );
    return errCount;
}

int HPCC_Stream(HPCC_Params *params, int doIO) {
    int i, j, k;
    FILE *outFile;
    VectorSize = HPCC_LocalVectorSize( params, 3, sizeof(double), 0 );
    a = HPCC_XMALLOC( double, VectorSize );
    b = HPCC_XMALLOC( double, VectorSize );
    c = HPCC_XMALLOC( double, VectorSize );
    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {
            fprintf( outFile, "Failed to allocate memory (%d).\n", VectorSize );
            fclose( outFile );
        }
        return 1;
    }

#ifdef _OPENMP
#pragma omp parallel for
#endif
    for (j=0; j<VectorSize; j++) {
        b[j] = 2.0;
        c[j] = 1.0;
    }
    scalar = 3.0;
#ifdef _OPENMP
#pragma omp parallel for
#endif
    for (j=0; j<VectorSize; j++) {
        a[j] = b[j]+scalar*c[j];
    }
    HPCC_free(c);
    HPCC_free(b);
    HPCC_free(a);
    return 0;
}
```



CUDA

```
#define N      2000000

int main() {
    float *d_a, *d_b, *d_c;
    float scalar;

    cudaMalloc((void**)&d_a, sizeof(float)*N);
    cudaMalloc((void**)&d_b, sizeof(float)*N);
    cudaMalloc((void**)&d_c, sizeof(float)*N);

    dim3 dimBlock(128);
    dim3 dimGrid(N/dimBlock.x);

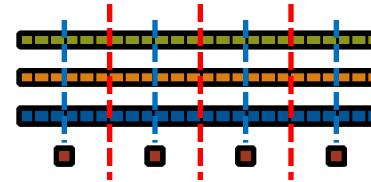
    set_array<<<dimGrid, dimBlock>>>(d_b, .5f, N);
    set_array<<<dimGrid, dimBlock>>>(d_c, .5f, N);

    scalar=3.0f;
    STREAM_Triad<<<dimGrid, dimBlock>>>(d_b, d_c, d_a, scalar, N);
    cudaThreadSynchronize();

    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);

__global__ void set_array(float *a, float value, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) a[idx] = value;
}

__global__ void STREAM_Triad( float *a, float *b, float *c,
                             float scalar, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) c[idx] = a[idx]+scalar*b[idx];
}
```



COMPUTE

STORE

ANALYZE

Why so many programming models?

HPC has traditionally given users...

- ...low-level, *control-centric* programming models
- ...ones that are closely tied to the underlying hardware
- ...ones that support only a single type of parallelism

Type of HW Parallelism	Programming Model	Unit of Parallelism
<i>Inter-node</i>	<i>MPI</i>	<i>executable</i>
<i>Intra-node/multicore</i>	<i>OpenMP / pthreads</i>	<i>iteration/task</i>
<i>Instruction-level vectors/threads</i>	<i>pragmas</i>	<i>iteration</i>
<i>GPU/accelerator</i>	<i>Open[MP CL ACC] / CUDA</i>	<i>SIMD function/task</i>

benefits: lots of control; decent generality; easy to implement
downsides: lots of user-managed detail; brittle to changes

Rewinding a few slides...

MPI + OpenMP

```
#ifdef _OPENMP
#include <omp.h>
#endif

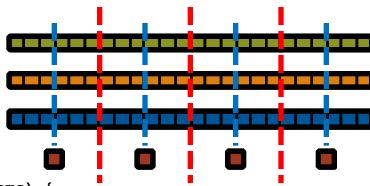
static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params) {
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;
    MPI_Comm_size( comm, &commSize );
    MPI_Comm_rank( comm, &myRank );
    rv = HPCC_Stream( params, 0 == myRank );
    MPI_Reduce( &rv, &errCount, 1, MPI_INT, MPI_SUM, 0, comm );
    return errCount;
}

int HPCC_Stream(HPCC_Params *params, int doIO) {
    FILE *outFile;
    if (doIO) {
        outFile = fopen("output.dat", "w");
        if (!outFile) {
            fprintf(outFile, "Failed to allocate memory (%d).\n", VectorSize);
            fclose(outFile);
        }
        return 1;
    }

#ifdef _OPENMP
#pragma omp parallel for
#endif
    for (j=0; j<VectorSize; j++) {
        b[j] = 2.0;
        c[j] = 1.0;
    }
    scalar = 3.0;
}

#ifdef _OPENMP
#pragma omp parallel for
#endif
    for (j=0; j<VectorSize; j++)
        a[j] = b[j]+scalar*c[j];
    HPCC_free(c);
    HPCC_free(b);
    HPCC_free(a);
    return 0;
}
```



CUDA

```
#define N 2000000

int main() {
    float *d_a, *d_b, *d_c;
    float scalar;

    cudaMalloc((void**)&d_a, sizeof(float)*N);
    cudaMalloc((void**)&d_b, sizeof(float)*N);
    cudaMalloc((void**)&d_c, sizeof(float)*N);

    dim3 dimBlock(128);
    dim3 dimGrid(N/dimBlock.x);

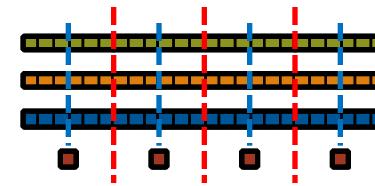
    set_array<<<dimGrid, dimBlock>>>(d_b, .5f, N);
    set_array<<<dimGrid, dimBlock>>>(d_c, .5f, N);

    scalar=3.0f;
    STREAM_Triad<<<dimGrid, dimBlock>>>(d_b, d_c, d_a, scalar, N);
    cudaThreadSynchronize();

    cudaFree(d_a);
    cudaFree(d_b);
    cudaFree(d_c);

__global__ void set_array(float *a, float value, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) a[idx] = value;
}

__global__ void STREAM_Triad( float *a, float *b, float *c,
                             float scalar, int len) {
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    if (idx < len) c[idx] = a[idx]+scalar*b[idx];
}
```



COMPUTE

STORE

ANALYZE

STREAM Triad: Chapel

MPI + OpenMP

```
#include <hpcc.h>
#ifndef _OPENMP
#include <omp.h>
#endif

static int VectorSize;
static double *a, *b, *c;

int HPCC_StarStream(HPCC_Params *params,
    int myRank, commSize;
    int rv, errCount;
    MPI_Comm comm = MPI_COMM_WORLD;
    MPI_Comm_size(comm, &commSize);
    MPI_Comm_rank(comm, &myRank);
    rv = HPCC_Stream(params, 0 == myRank);
    MPI_Reduce(&rv, &errCount, 1, MPI_
    return errCount;
}

int HPCC_Stream(HPCC_Params *params,
    register int j;
    double scalar;
    VectorSize = HPCC_LocalVectorSize();
    a = HPCC_XMALLOC(double, VectorSize);
    b = HPCC_XMALLOC(double, VectorSize);
    c = HPCC_XMALLOC(double, VectorSize);
    if (!a || !b || !c) {
        if (c) HPCC_free(c);
        if (b) HPCC_free(b);
        if (a) HPCC_free(a);
        if (doIO) {
            fprintf(outFile, "Failed to al
            fclose(outFile);
        }
    }
    return 1;
}
```

Chapel

```
config const m = 1000,
alpha = 3.0;

const ProblemSpace = {1..m} dmapped ...;

var A, B, C: [ProblemSpace] real;

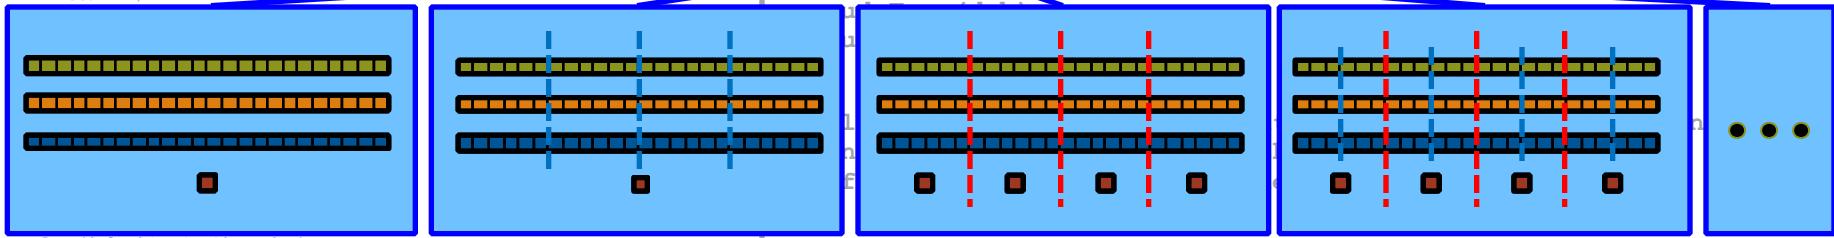
B = 2.0;
C = 1.0;

A = B + alpha * C;
```

dmapped ...;

the special
sauce

cudaFree(d_a);



```
for (j=0; j<VectorSize; j++)
    a[j] =
```

```
HPCC
HPCC
HPCC_
return
```

Philosophy: Good language design can tease details of locality and parallelism away from an algorithm, permitting the compiler, runtime, applied scientist, and HPC expert to each focus on their strengths.

COMPUTE

STORE

ANALYZE

Motivating Chapel Themes

- 1) General Parallel Programming**
- 2) Global-View Abstractions**
- 3) Multiresolution Design**
- 4) Control over Locality/Affinity**
- 5) Reduce HPC ↔ Mainstream Language Gap**

Motivating Chapel Themes

- 1) General Parallel Programming**
- 2) Global-View Abstractions**
- 3) Multiresolution Design**
- 4) PGAS: Control over Locality/Affinity**
- 5) Reduce HPC ↔ Mainstream Language Gap**

1) General Parallel Programming

With a unified set of concepts...

...express any parallelism desired in a user's program

- **Styles:** data-parallel, task-parallel, concurrency, nested, ...
- **Levels:** model, function, loop, statement, expression

...target any parallelism available in the hardware

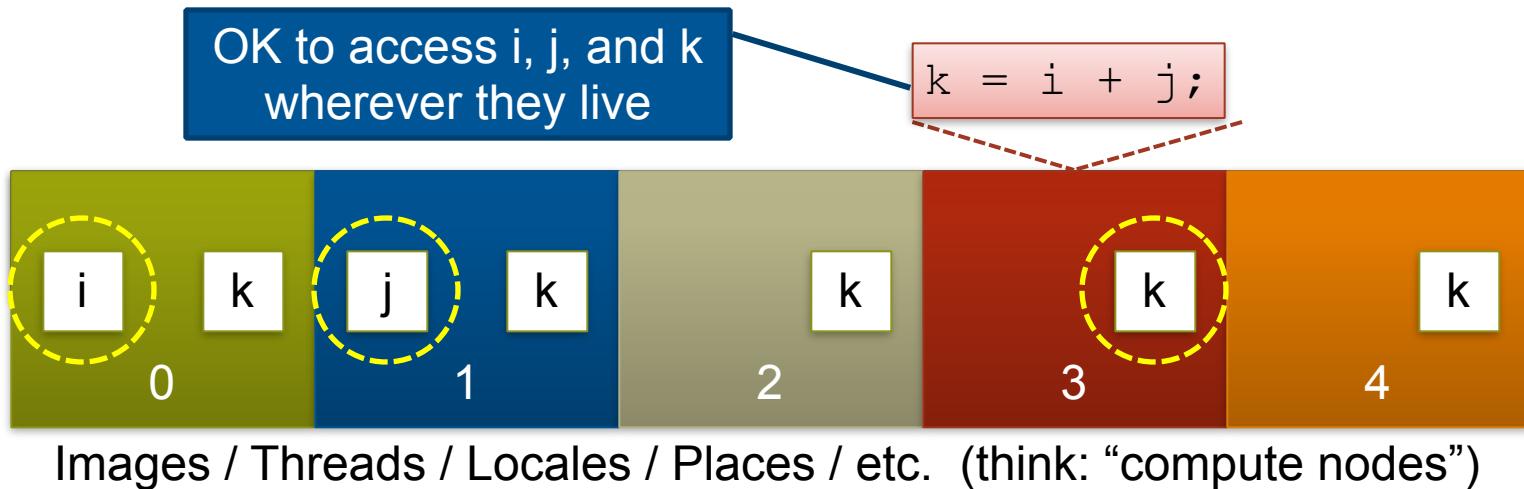
- **Types:** machines, nodes, cores, instruction

Type of HW Parallelism	Programming Model	Unit of Parallelism
<i>Inter-node</i>	<i>Chapel</i>	<i>task(or executable)</i>
<i>Intra-node/multicore</i>	<i>Chapel</i>	<i>iteration/task</i>
<i>Instruction-level vectors/threads</i>	<i>Chapel</i>	<i>iteration</i>
<i>GPU/accelerator</i>	<i>Chapel</i>	<i>SIMD function/task</i>

PGAS Programming in a Nutshell

Global Address Space:

- permit parallel tasks to access variables by naming them
 - regardless of whether they are local or remote
 - compiler / library / runtime will take care of communication



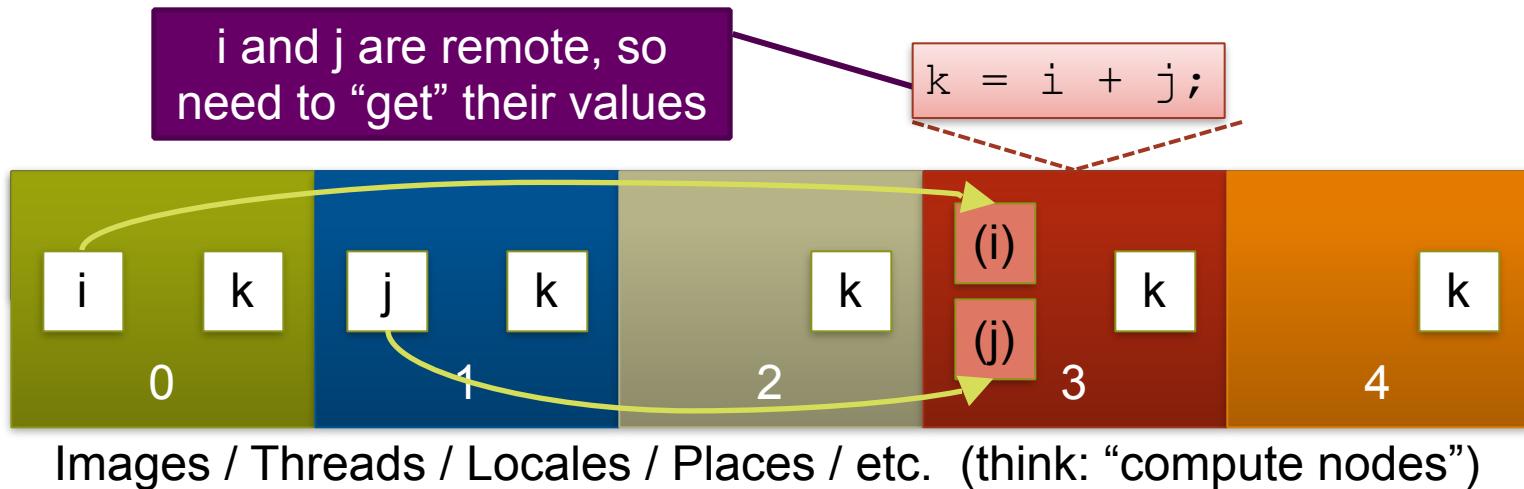
PGAS Programming in a Nutshell

Global Address Space:

- permit parallel tasks to access variables by naming them
 - regardless of whether they are local or remote
 - compiler / library / runtime will take care of communication

Partitioned:

- establish a strong model for reasoning about locality
 - every variable has a well-defined location in the system
 - local variables are typically cheaper to access than remote ones



COMPUTE

|

STORE

|

ANALYZE

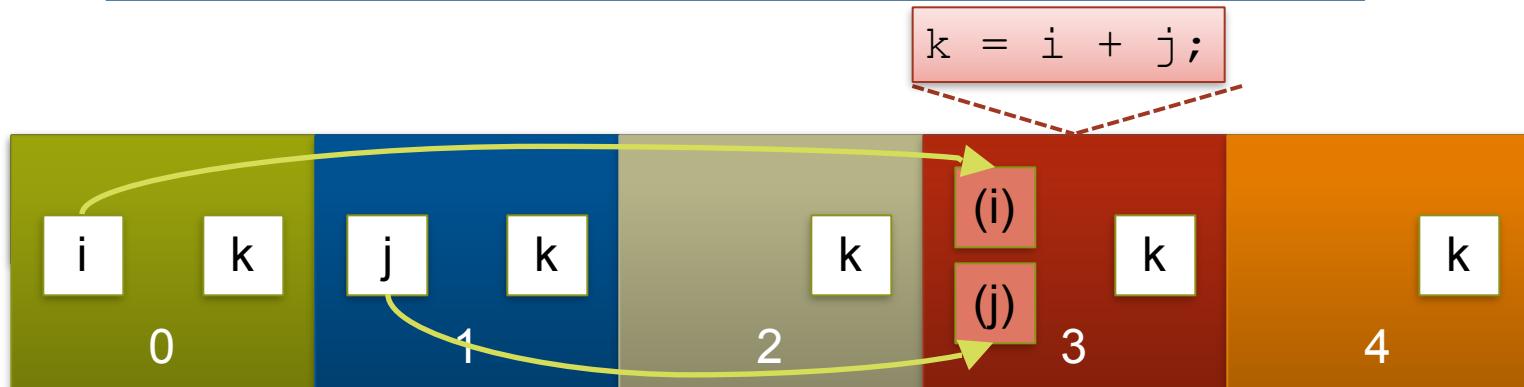
PGAS Programming in a Nutshell

Global Address Space:

- permit parallel tasks to access variables by naming them
 - regardless of whether they are local or remote
 - compiler / library / runtime will take care of communication

Partitioned:

- establish partition boundaries
 - even
 - local
- Communication is implicit!
One sided GET and PUT.



Images / Threads / Locales / Places / etc. (think: “compute nodes”)

WHY COMMUNICATION OPTIMIZATION?



COMPUTE

|

STORE

|

ANALYZE



CC Flickr/Daniel Jolivet

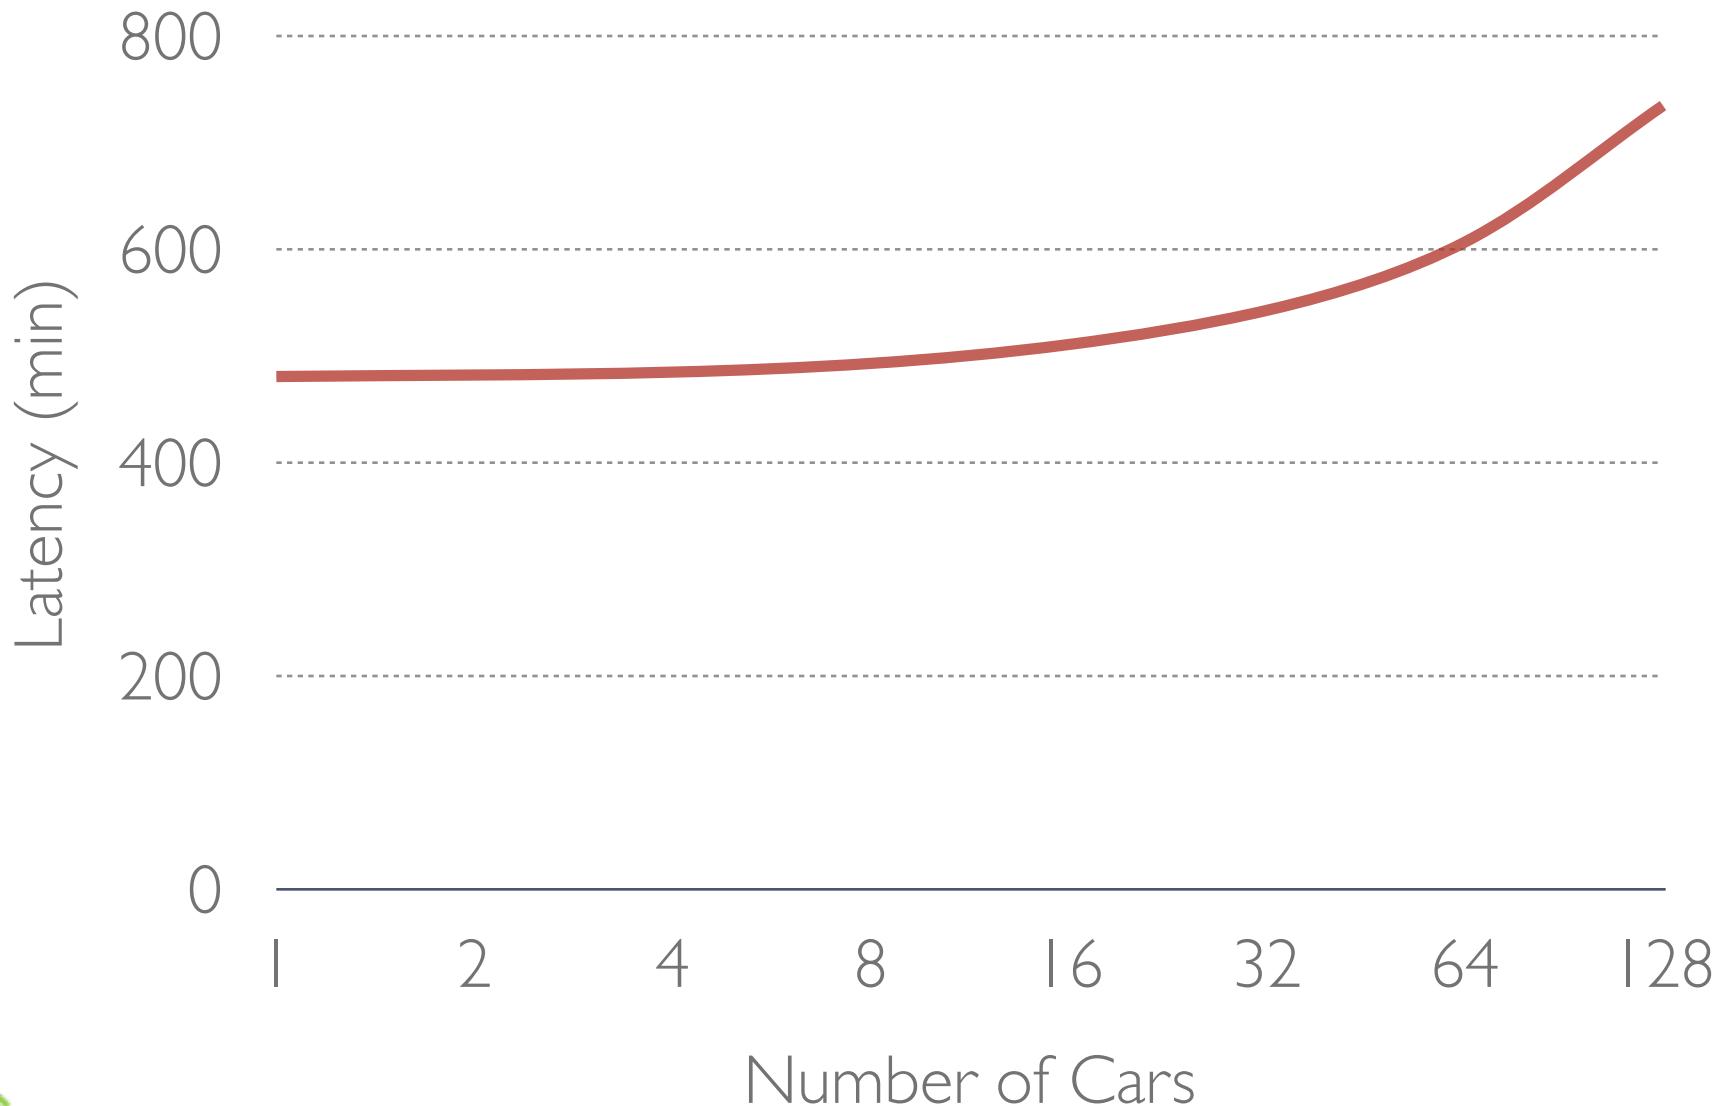


CC Flickr/Ben Salter

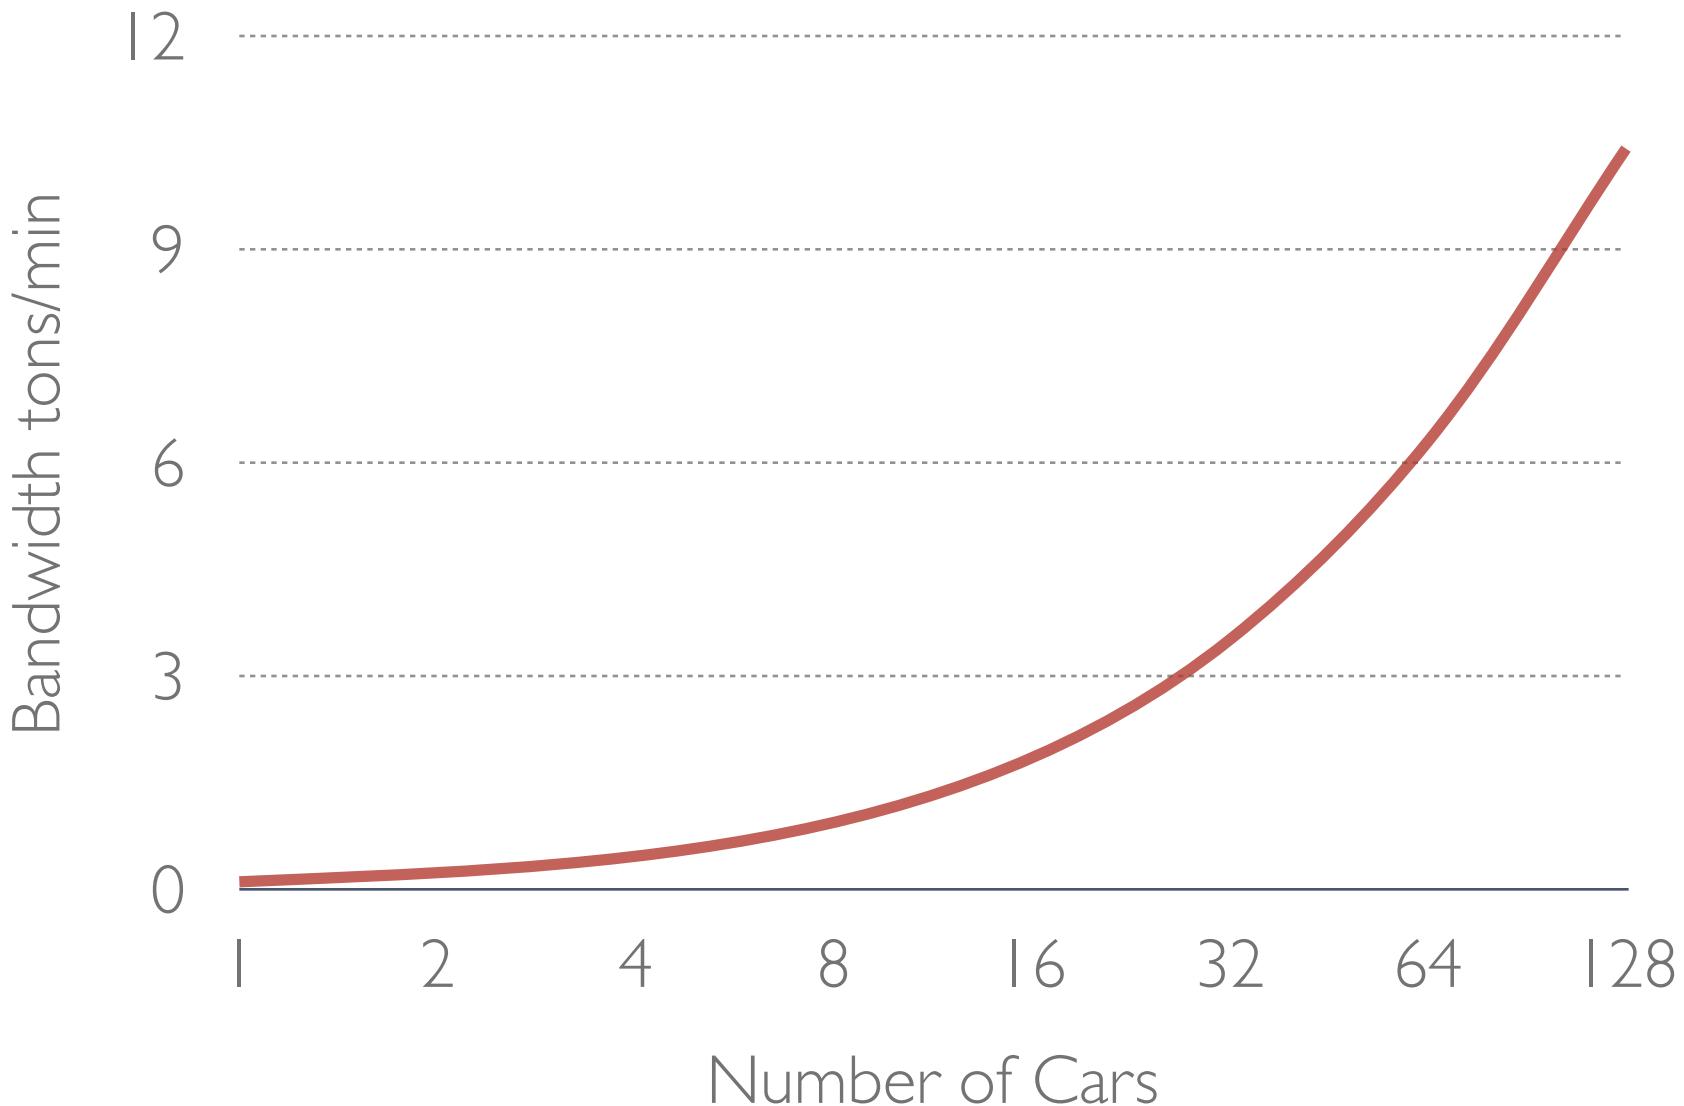


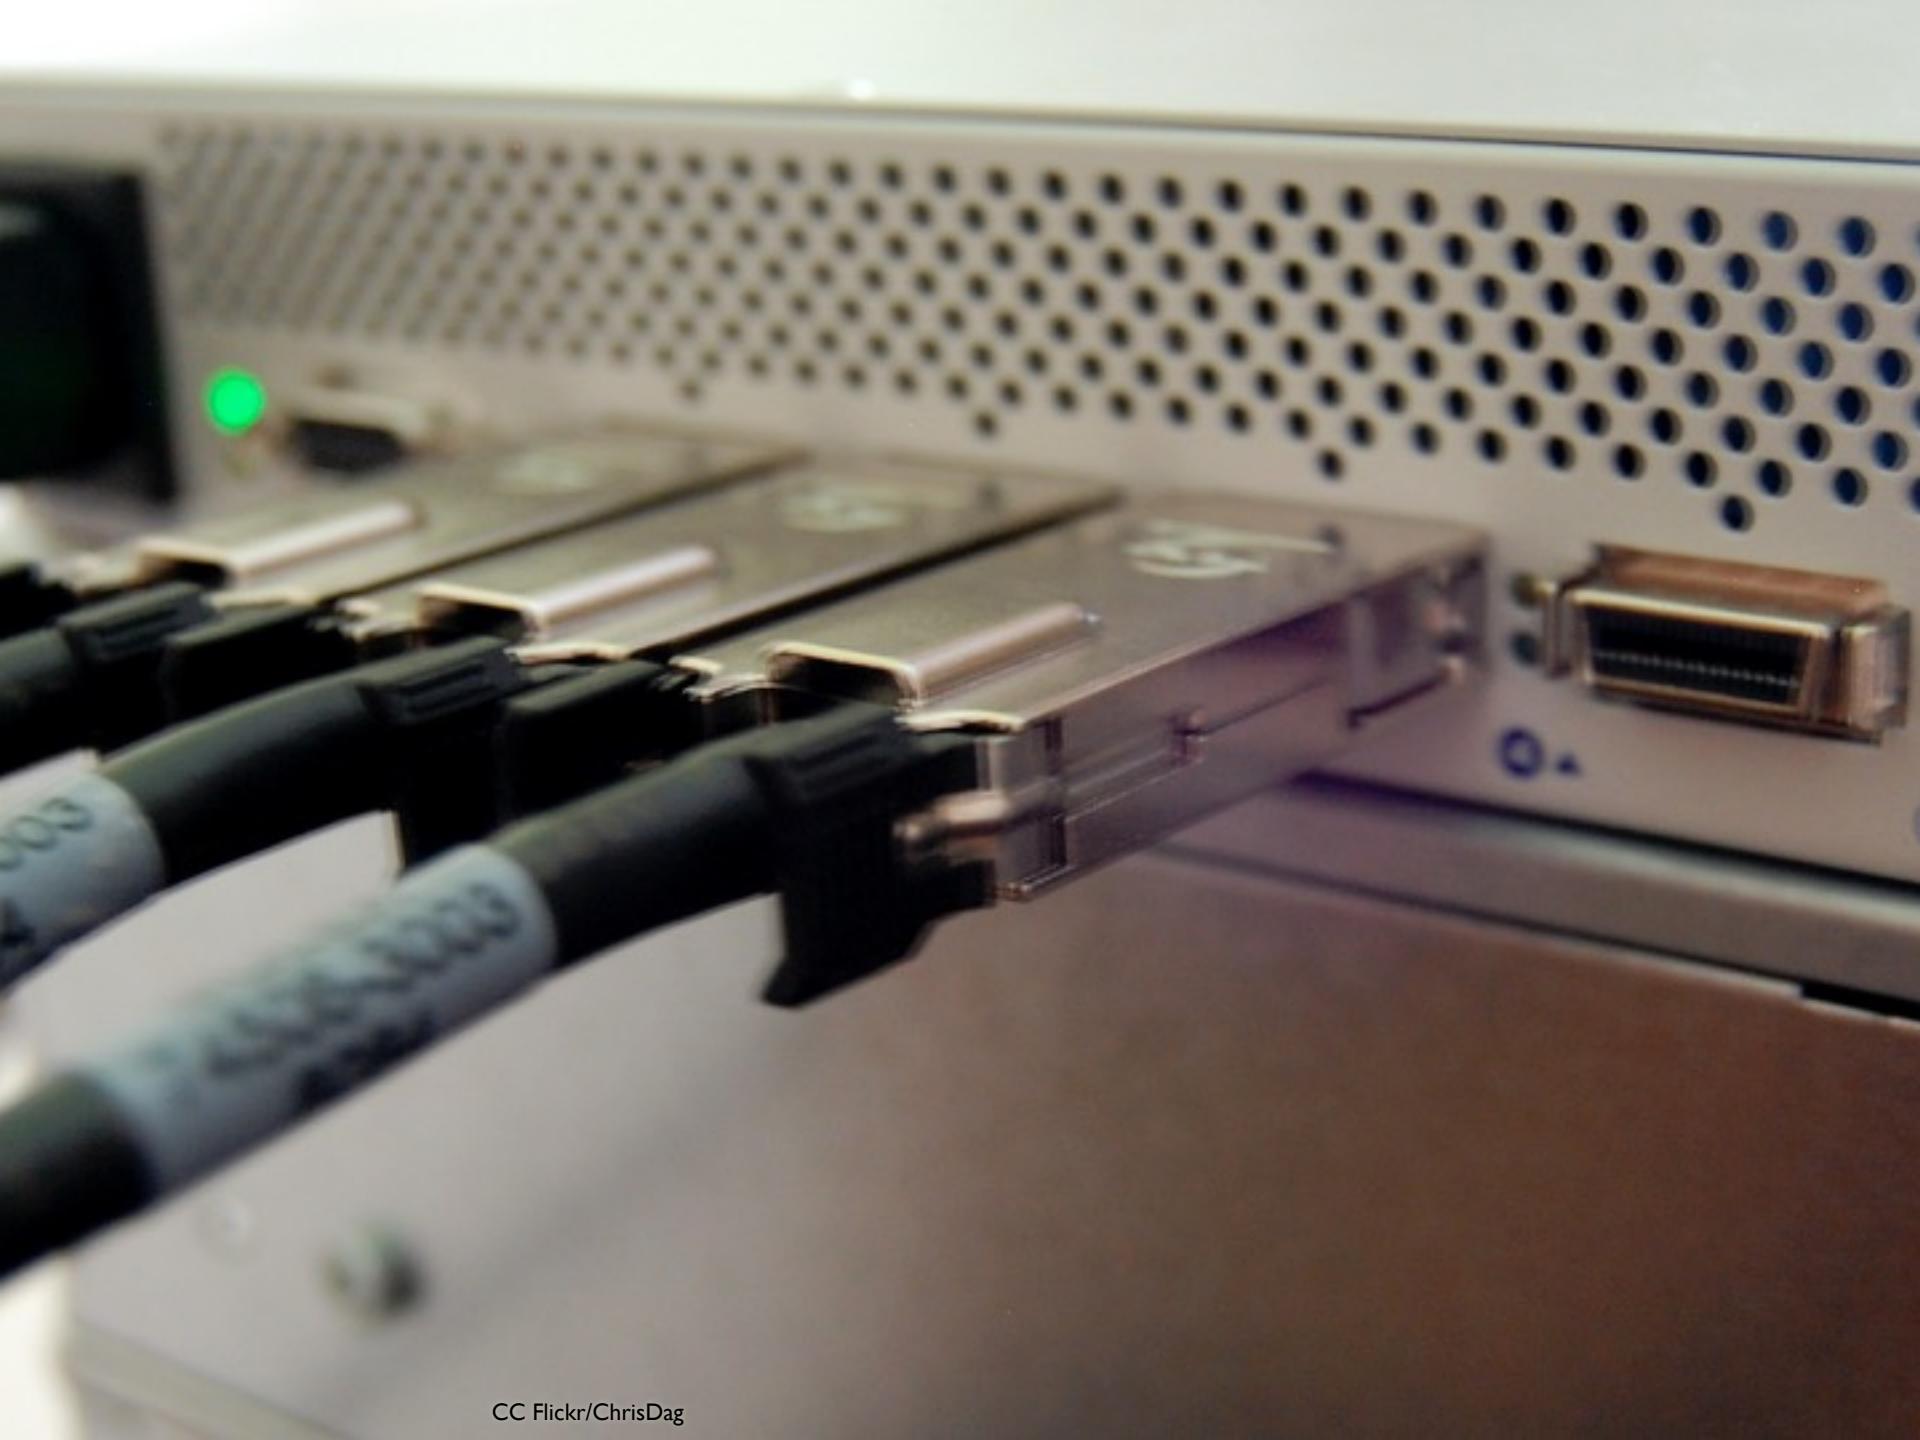
TRAIN LATENCY

(8 HOUR TRIP, 60 TON CARS, 60 SEC/CAR)



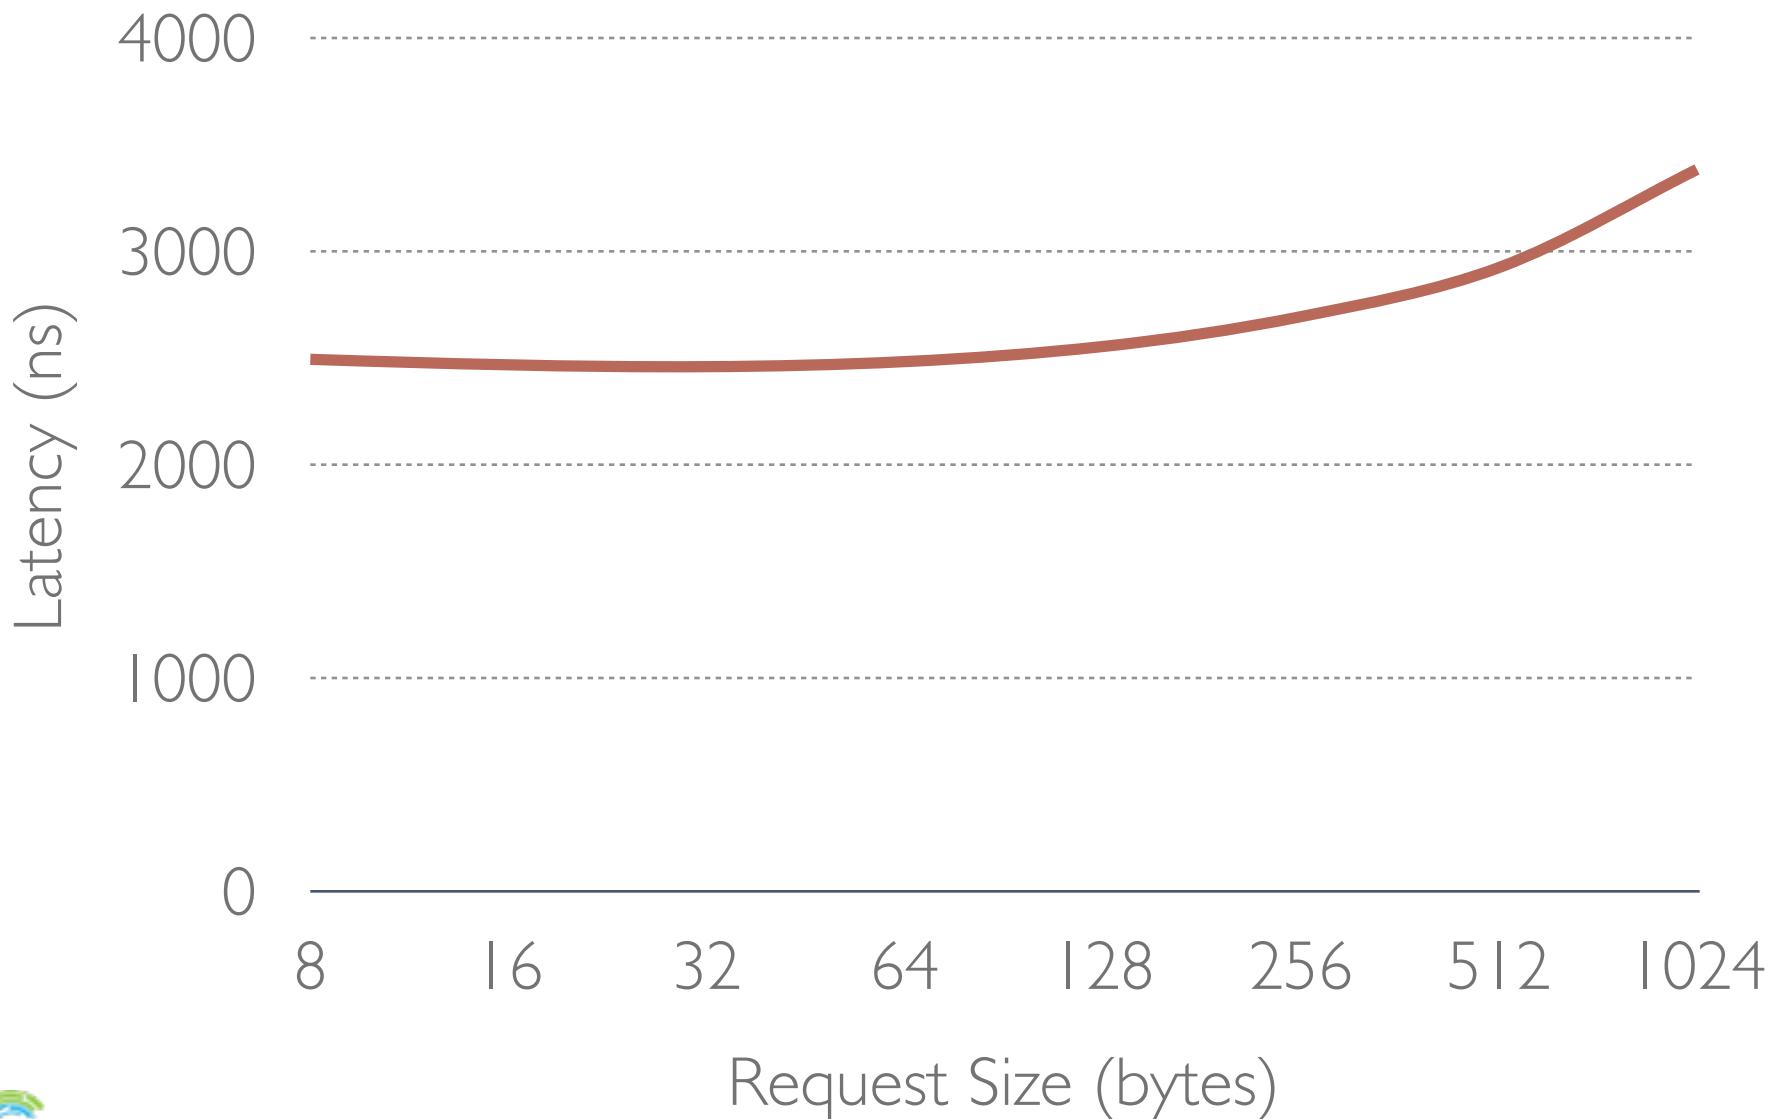
TRAIN BANDWIDTH





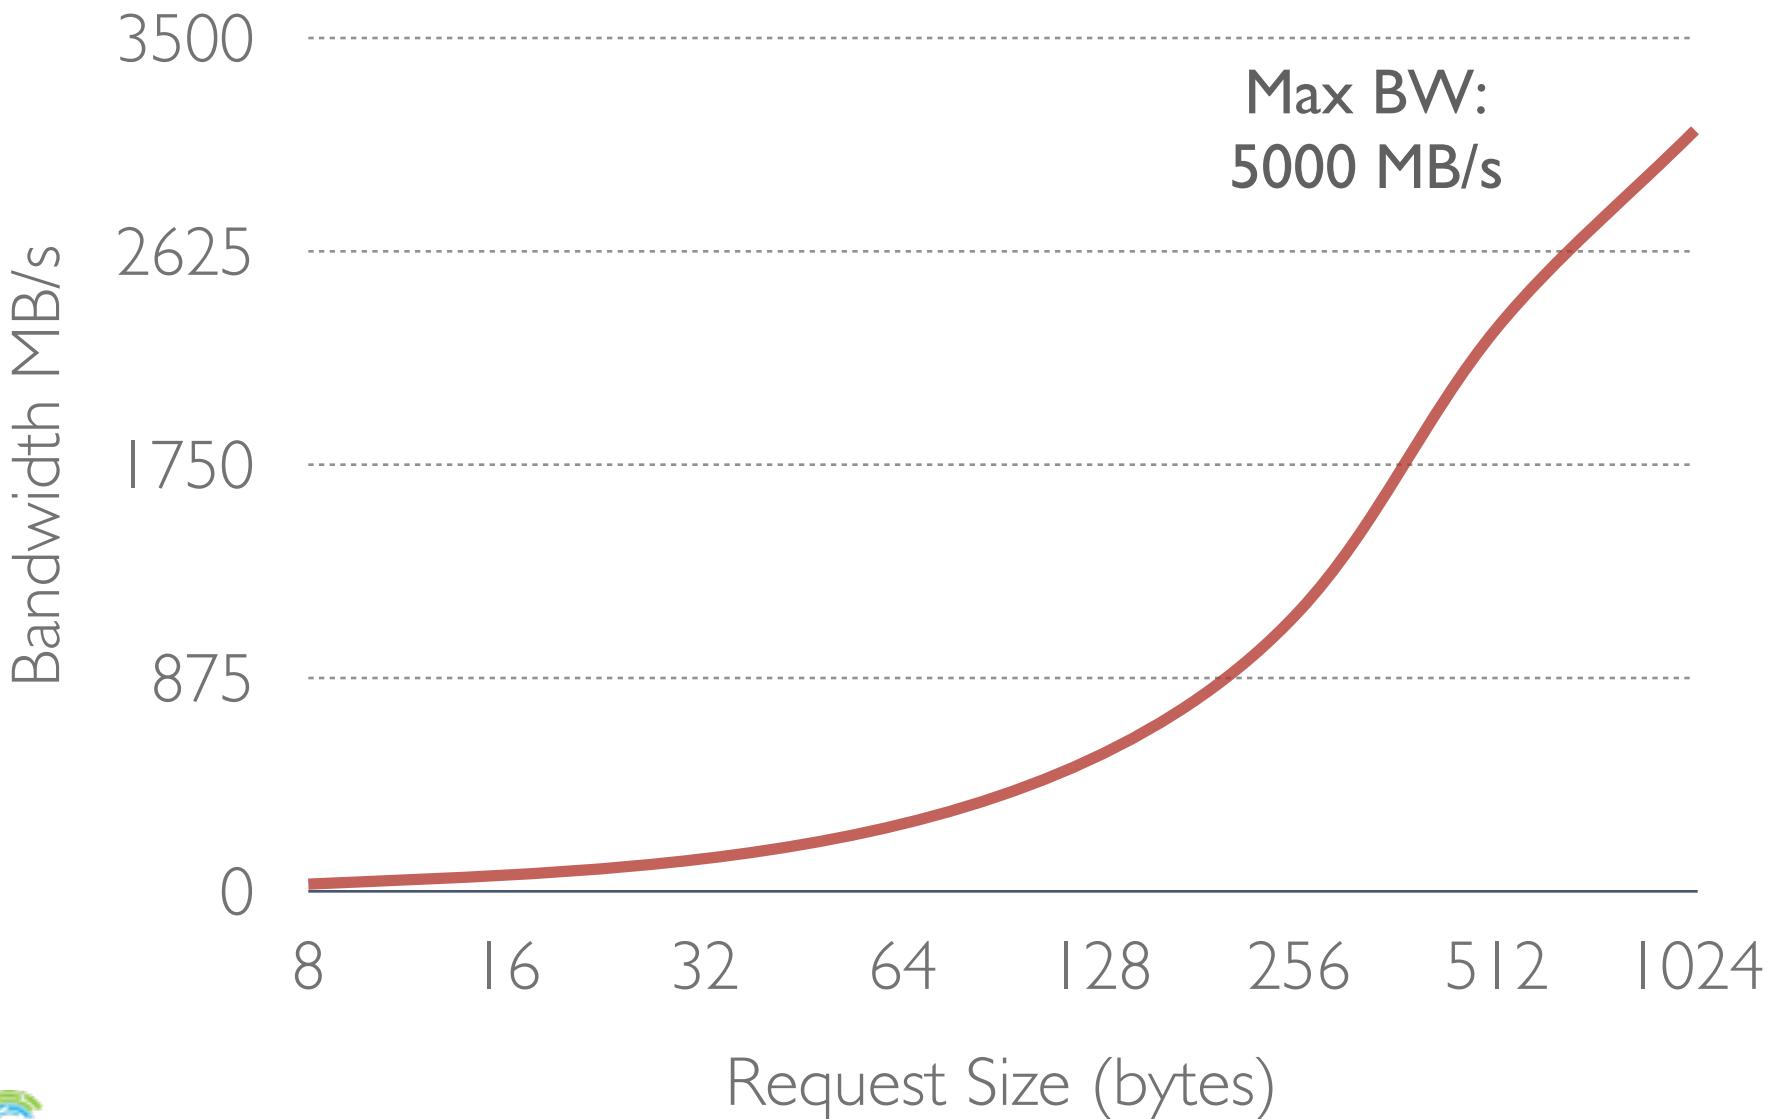
INFINIBAND (IB) LATENCY

* with small 10-node cluster, QDR IB



INFINIBAND (IB) BANDWIDTH

* with small 10-node cluster, QDR IB



MEMORY MODELS CONSTRAIN PREFETCH AND WRITE-BEHIND



AGGREGATION



CC Flickr/ajmexico

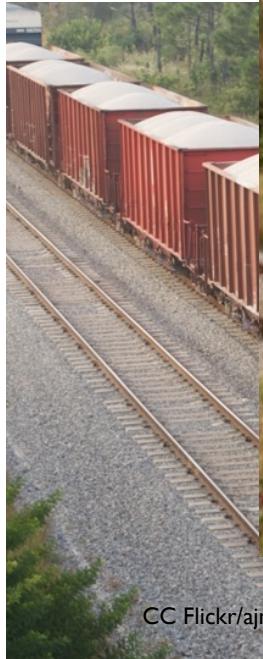
OVERLAP



CC Flickr/Barry Lewis

AGC

AP



CC Flickr/ajmexico

CC Flickr/Ben Salter

CC Flickr/Barry Lewis

AGC

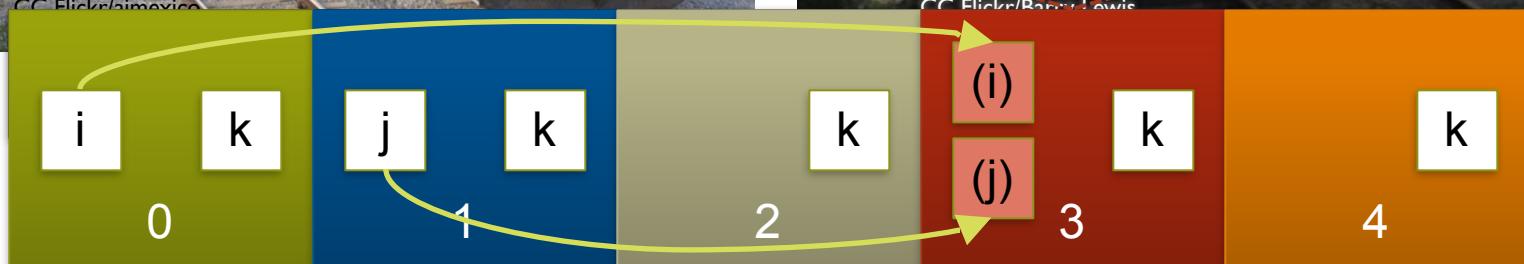
AP



CC Flickr/Ben Salter

$$k = i + j;$$

CC Flickr/Randy Lewis



A RACY PROGRAM

Thread 1

```
x = 42;  
notify = 1;
```

Thread 2

```
while 0 == notify { /* wait */  
compute_with(x);
```

A RACY PROGRAM

Thread 1

```
x = 42;
```

```
notify = 1;
```

Thread 2

```
while 0 == notify { /* wait */ }
```

```
compute_with(x);
```

Thread 1

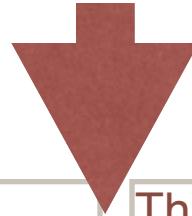
```
r1 = 42;
```

```
notify = 1; x = r1;
```

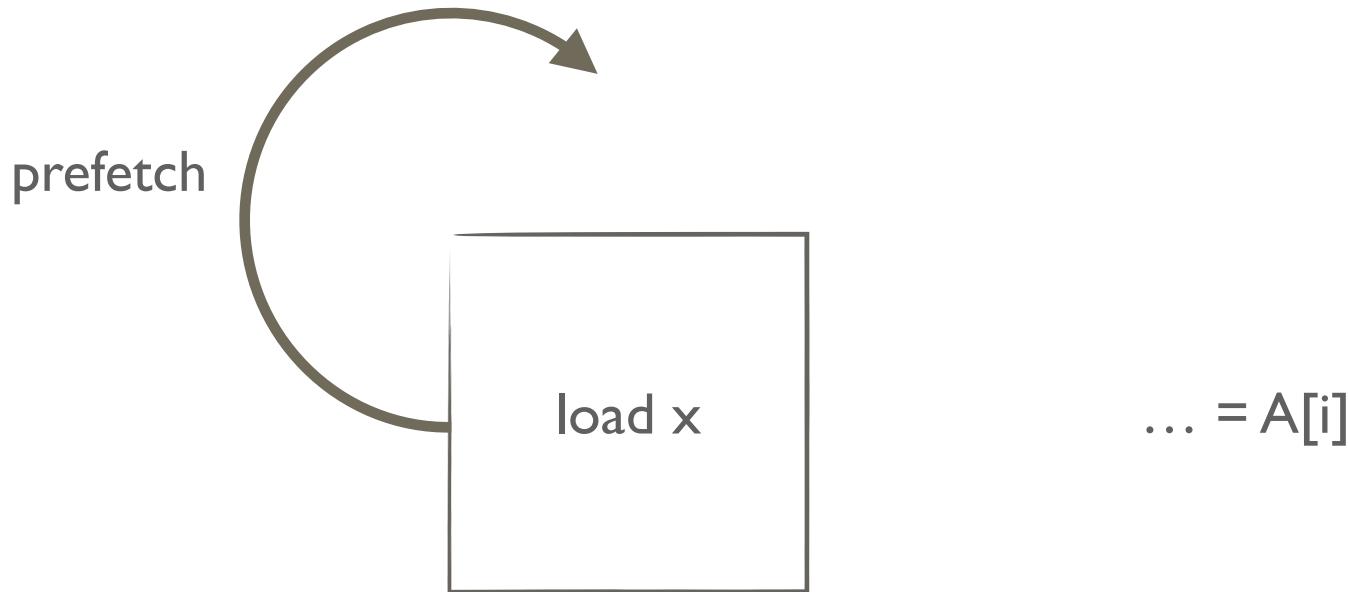
Thread 2

```
r2 = notify; while 0 == r2 { /* wait */ }
```

```
compute_with(x);
```



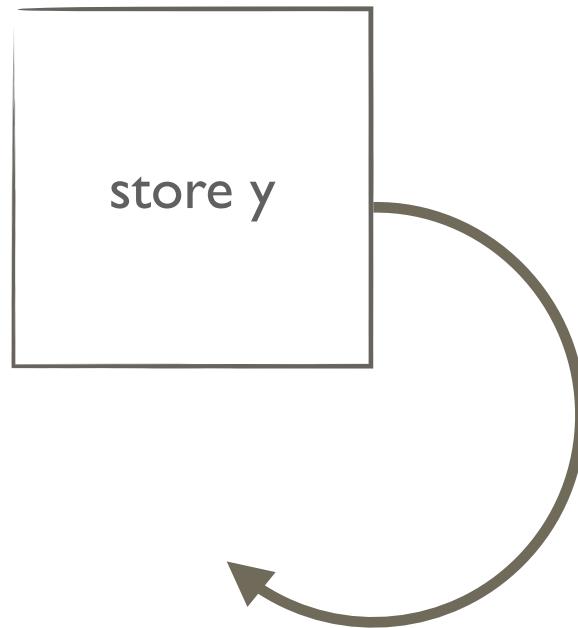
compiler or processor



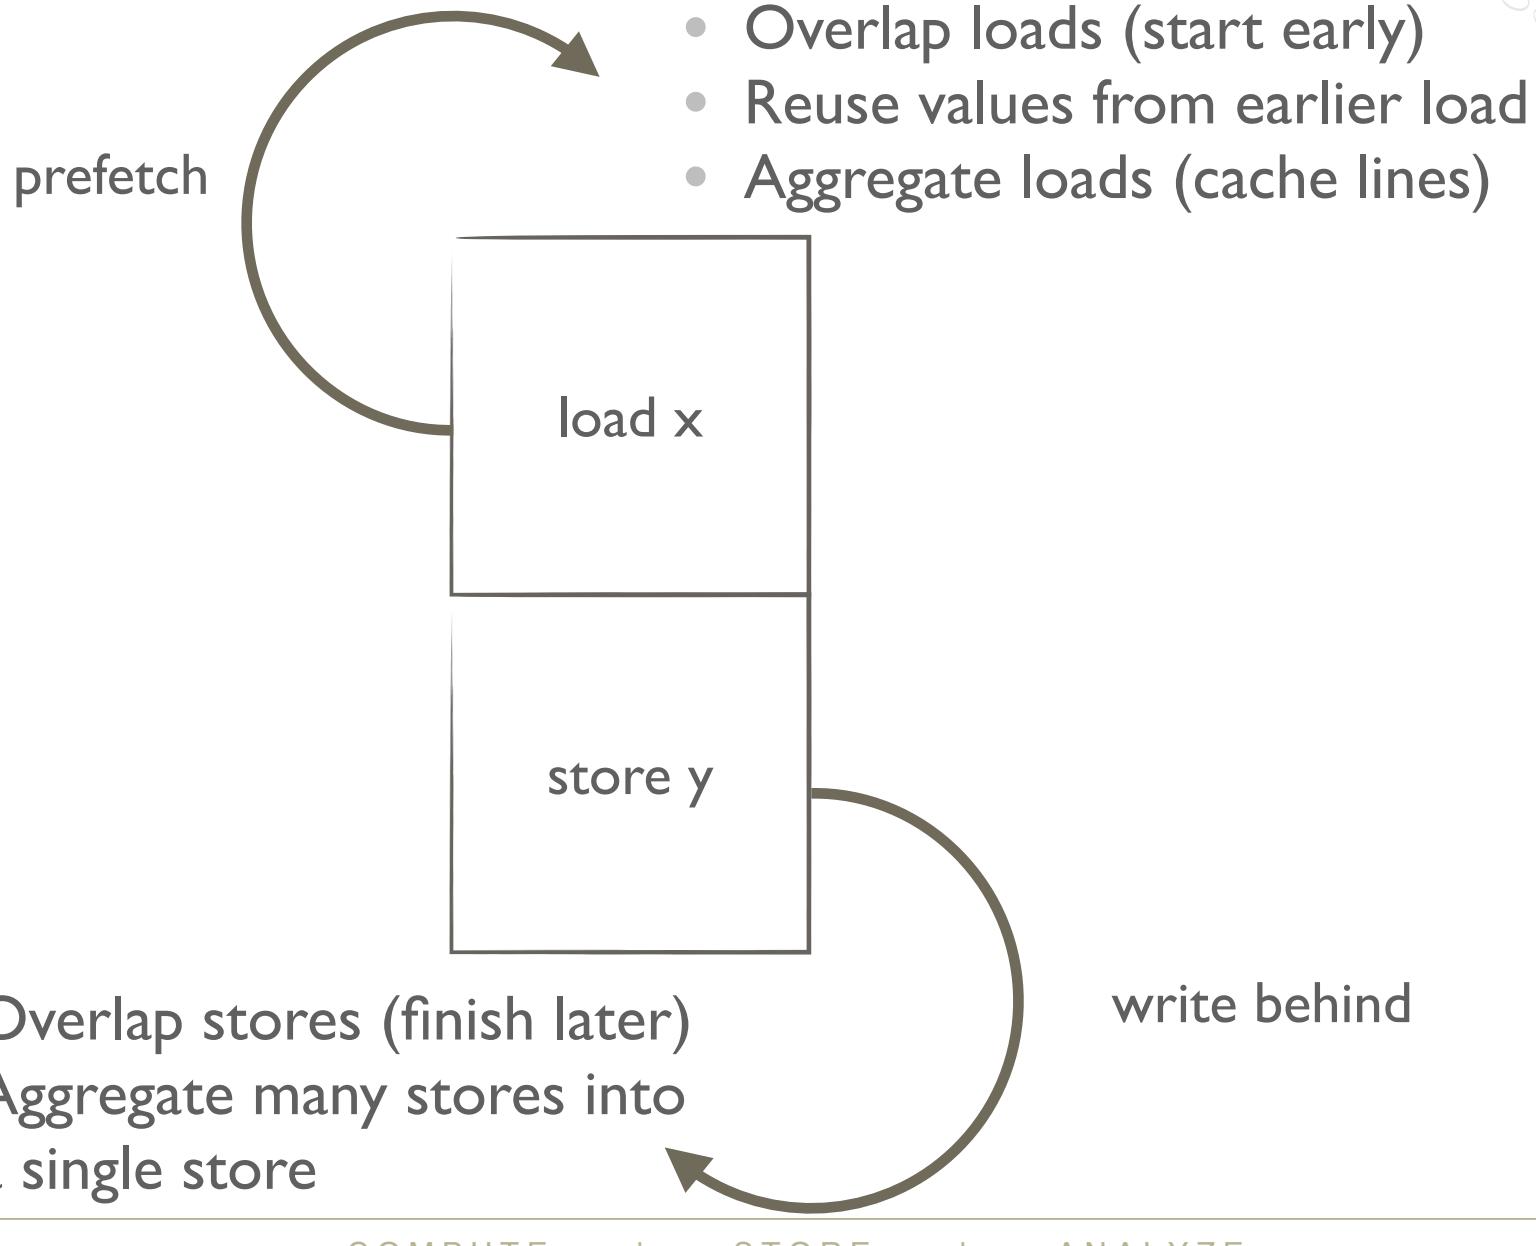
Compiler and processor would like to start loads earlier in order to hide memory latency. We'll call that *prefetch*.

Compiler and processor would like to complete stores later in order to hide memory latency. We'll call that *write behind*.

$B[i] = \dots$



write behind



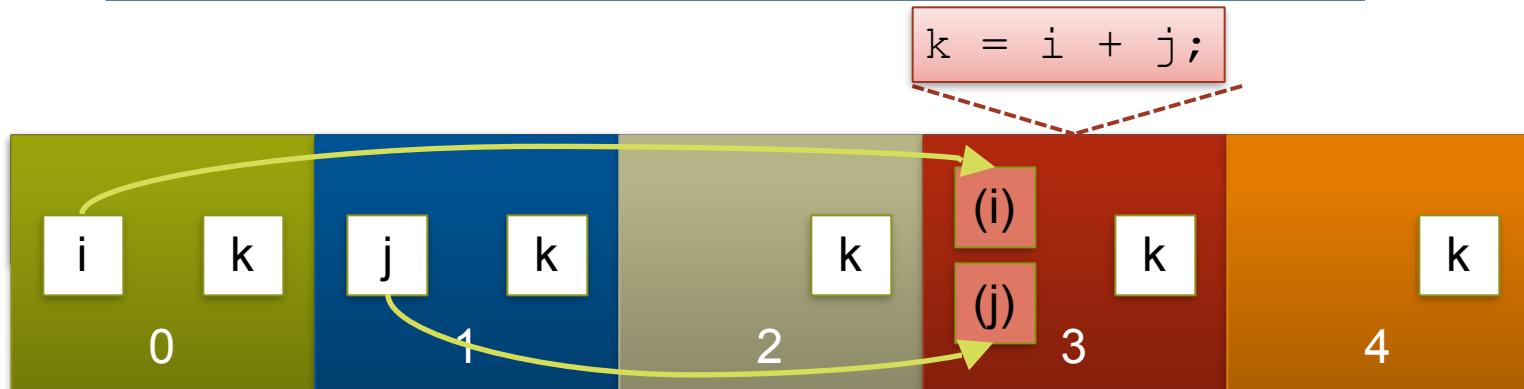
PGAS Programming in a Nutshell

Global Address Space:

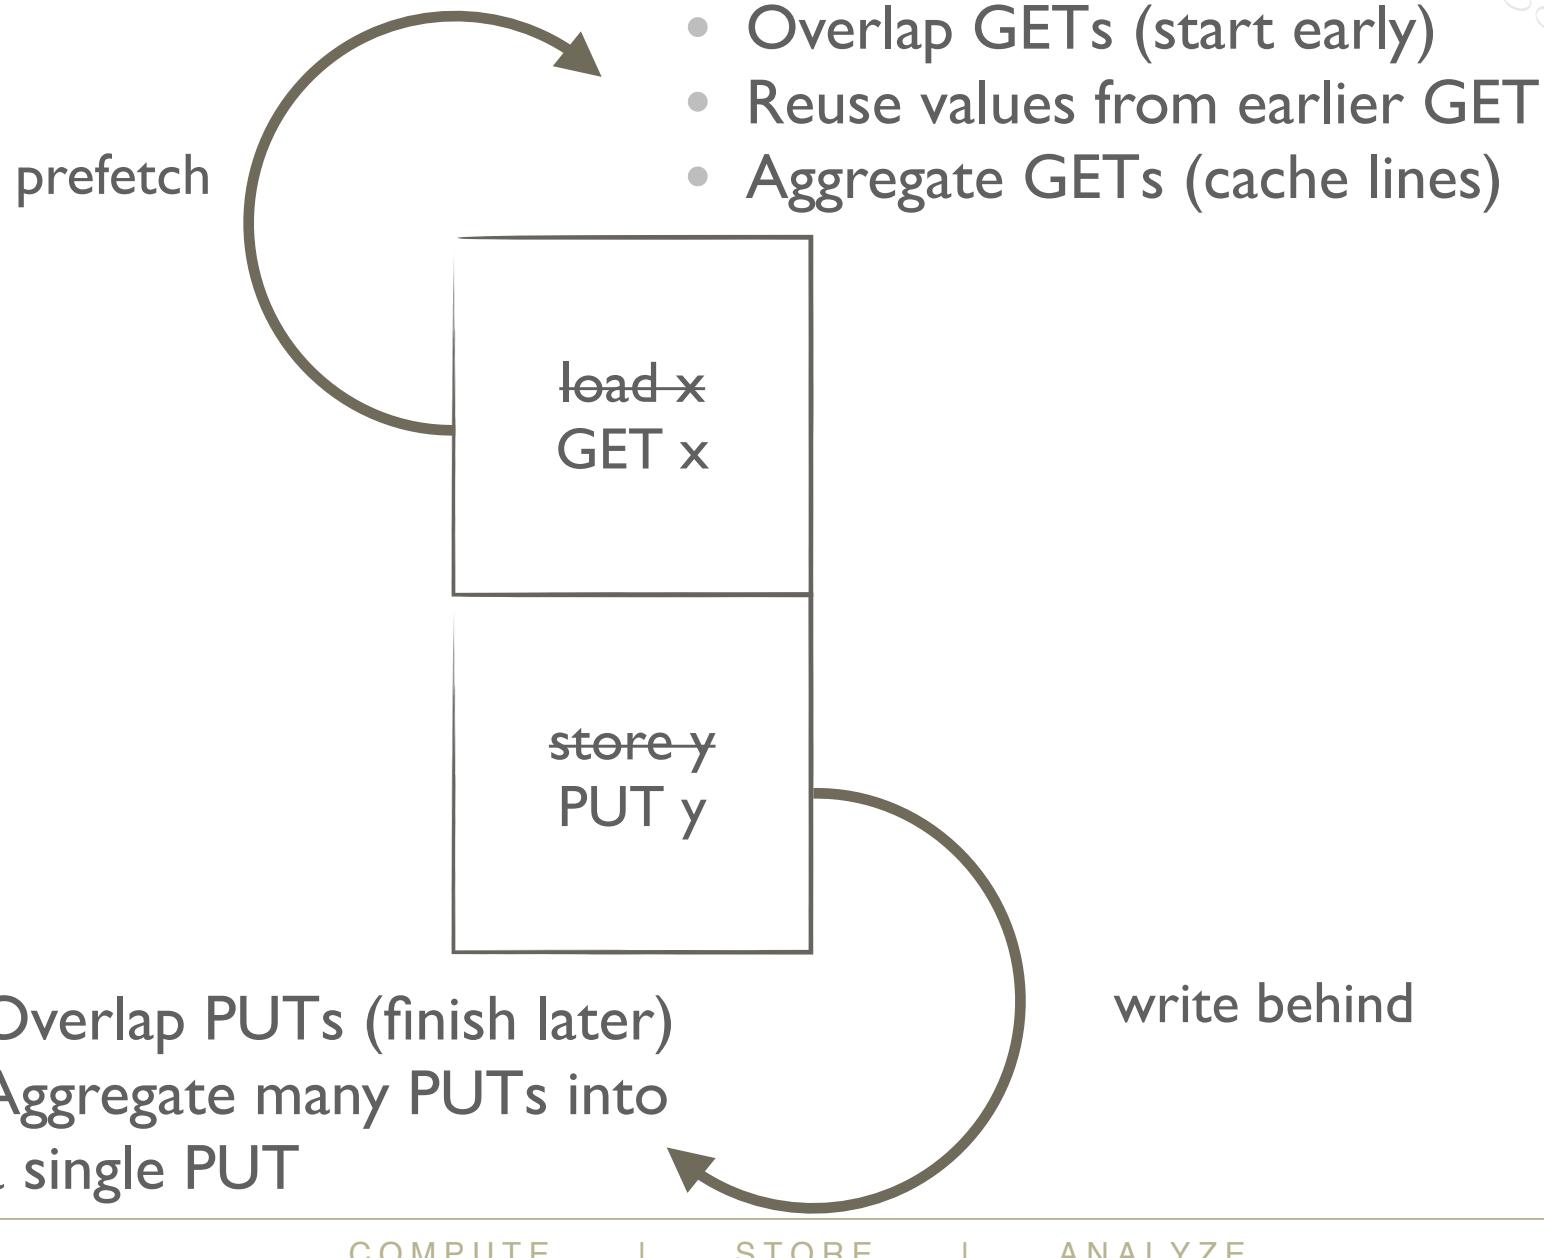
- permit parallel tasks to access variables by naming them
 - regardless of whether they are local or remote
 - compiler / library / runtime will take care of communication

Partitioned:

- establish partition boundaries
 - even
 - local
- Communication is implicit!
One sided GET and PUT.



Images / Threads / Locales / Places / etc. (think: “compute nodes”)



REMEMBER THE RACY PROGRAM?



Thread 1

```
x = 42;
```

```
notify = true;
```

Thread 2

```
while 0 == notify { /* wait */ }
```

```
compute_with(x);
```

Thread 1

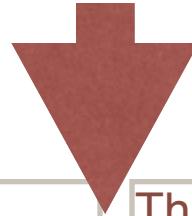
```
r1 = 42;
```

```
notify = 1; x = r1;
```

Thread 2

```
r2 = notify; while 0 == r2 { /* wait */ }
```

```
compute_with(x);
```



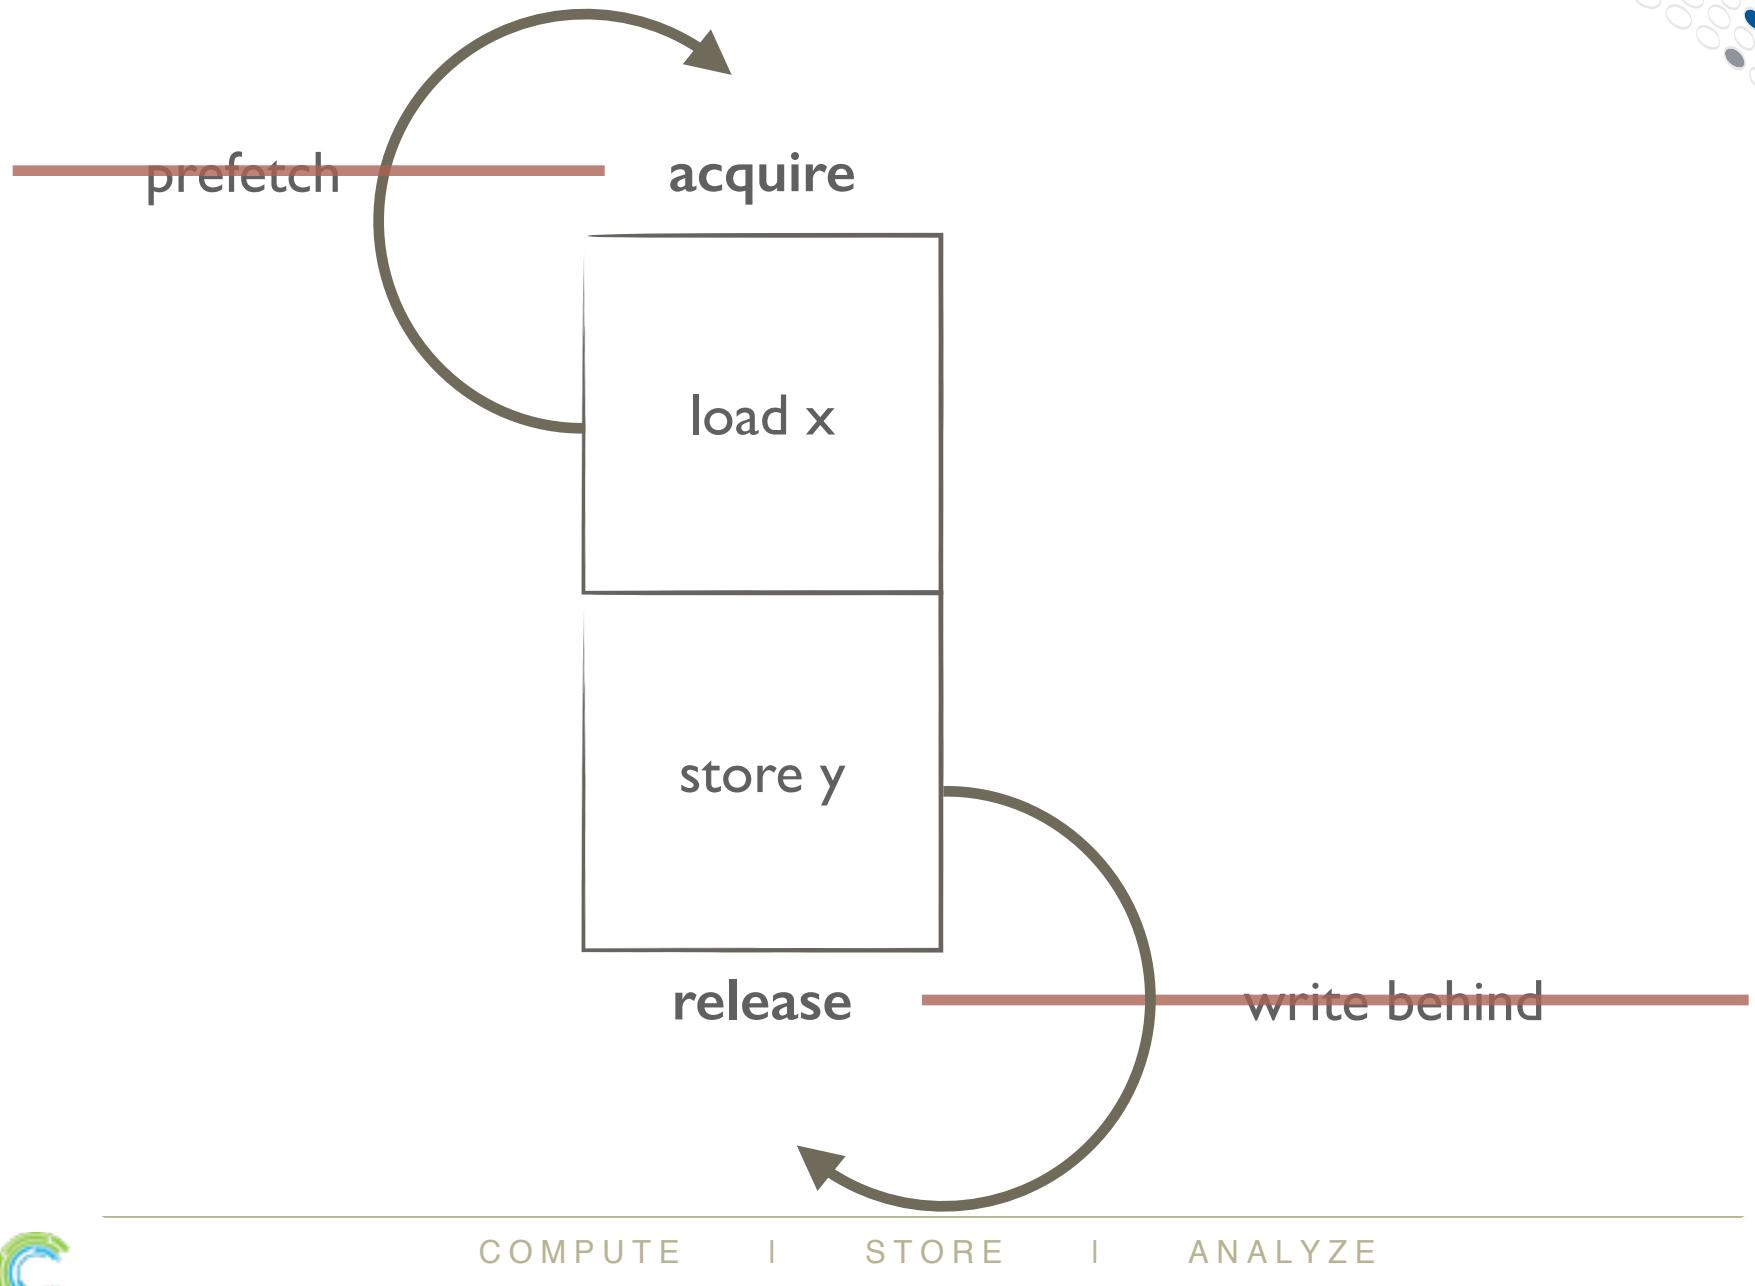
compiler or processor

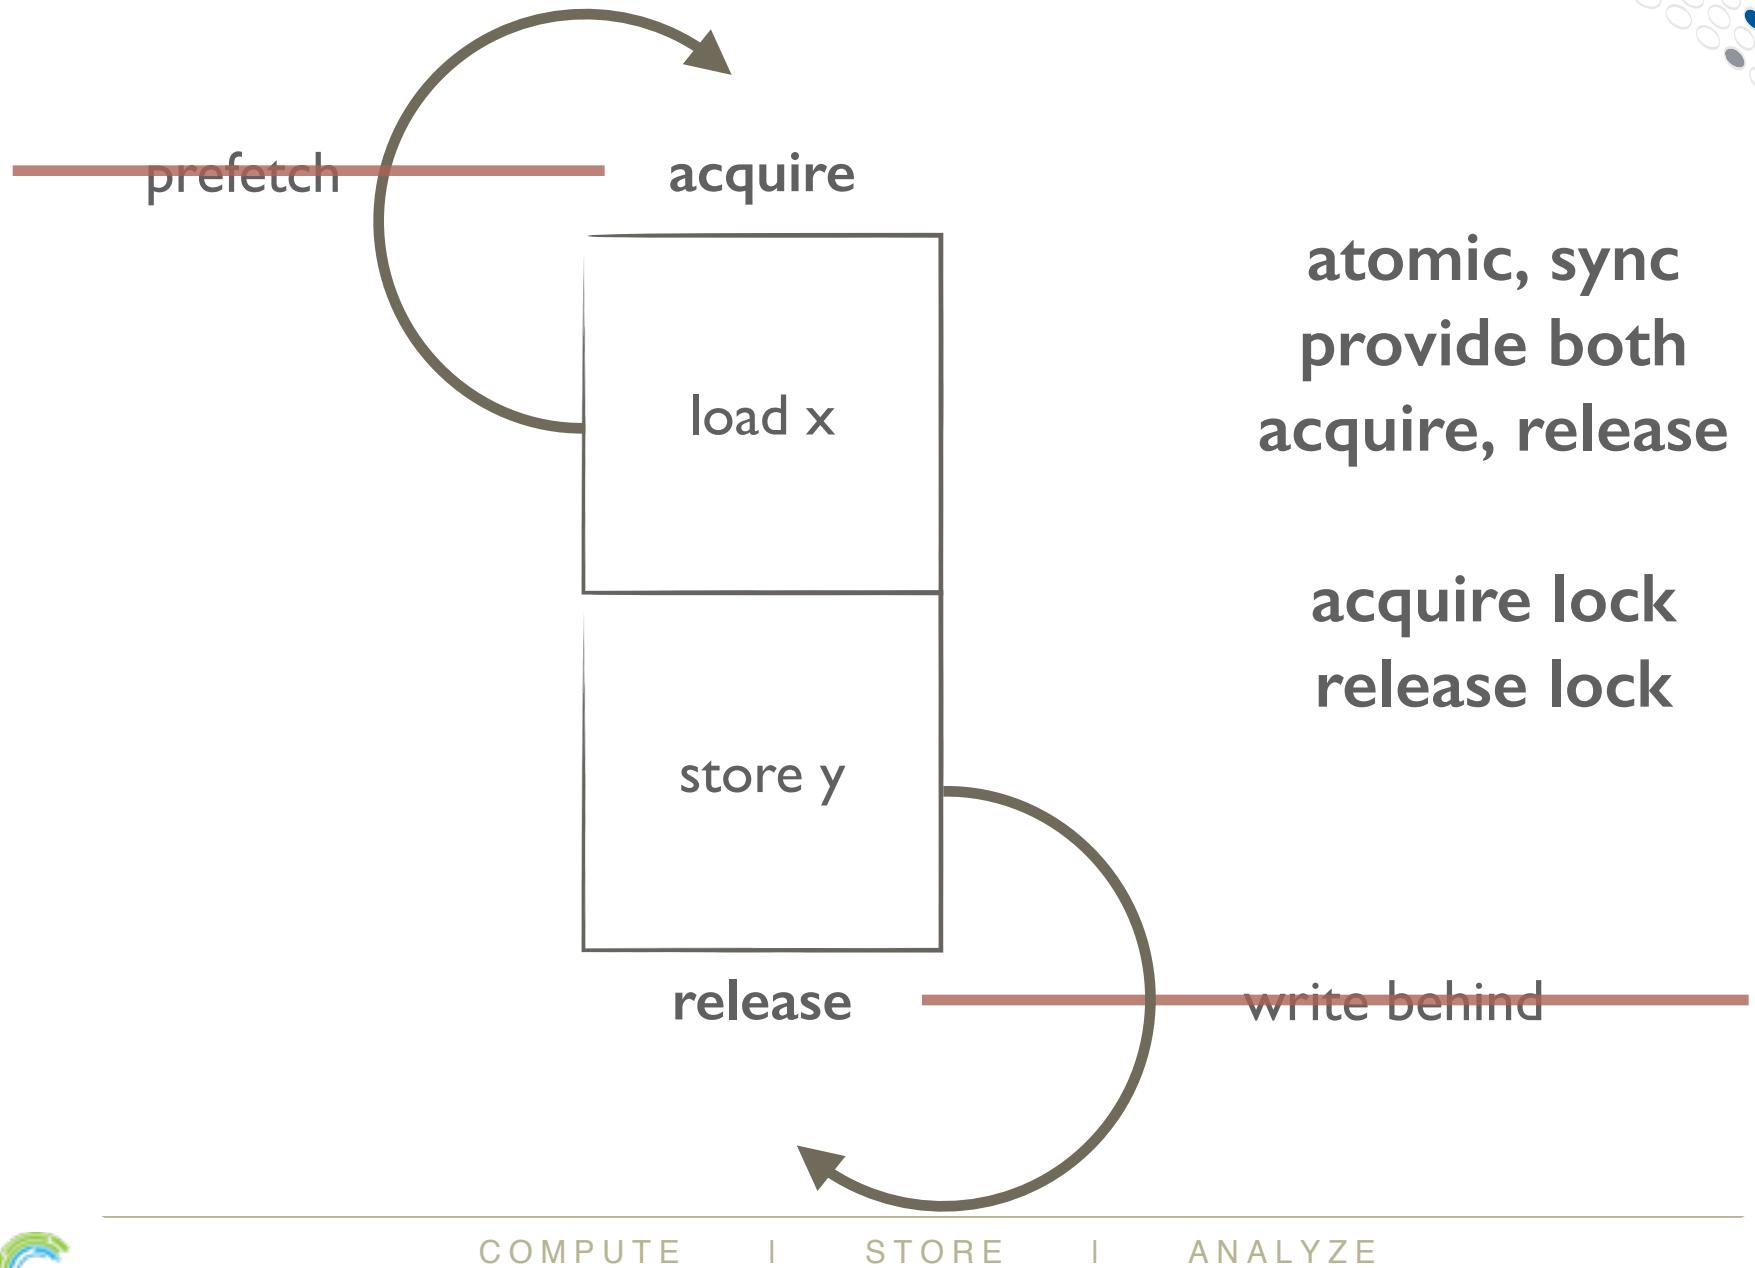


COMPUTE

STORE

ANALYZE





SC FOR DRF



COMPUTE

|

STORE

|

ANALYZE



10170

CASTLE ROCK

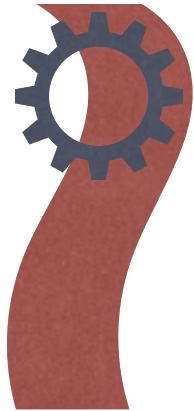
Memory model for **C11, C++11, Chapel:** *data race free programs are sequentially consistent*

- See Adve, S.V., Boehm, H.-J. 2010. Memory models: a case for rethinking parallel languages and hardware. Communications of the ACM 53(8): 90–101. <http://cacm.acm.org/magazines/2010/8/96610-memory-models-a-case-for-rethinking-parallel-languages-and-hardware/fulltext>
- Chapel has a new specification chapter describing the memory consistency model. See <http://chapel.cray.com/spec/spec-0.98.pdf> section 29, page 217.

CONFIGURABLE SC-DRF

- atomic operations in Chapel and C++ support:
 - *memory_order_relaxed* "atomic only"
 - *memory_order_acquire* "acquire"
 - *memory_order_release* "release"
 - *memory_order_seq_cst* "sequentially consistent"
- Beware! No global total order for relaxed, acquire, and release. Instead, the order is per atomic variable.

Program Order



x = 1;



Memory Order



Program Order

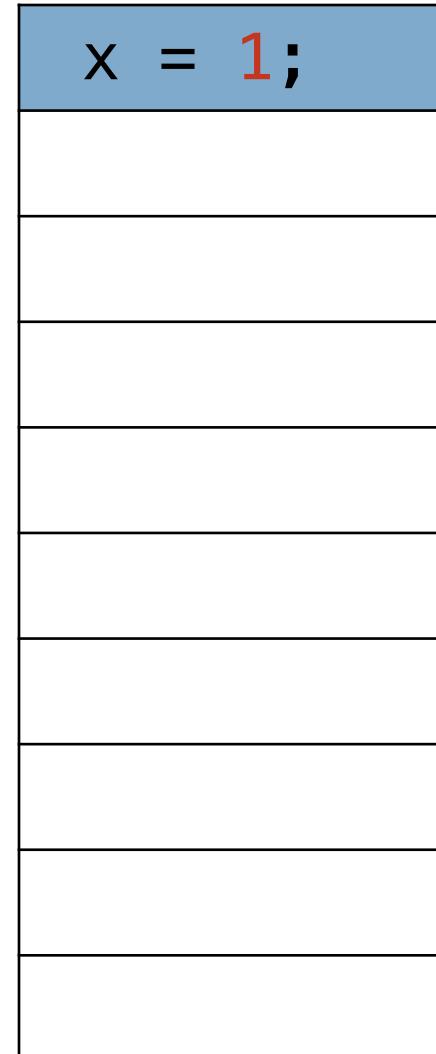


x = 1;



Memory Order

x = 1;



Program Order



```
x = 1;
```

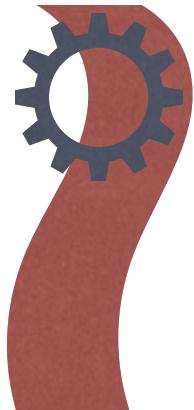


```
y = 3;  
z = 4;
```

Memory Order

x = 1;	

Program Order



```
x = 1;
```

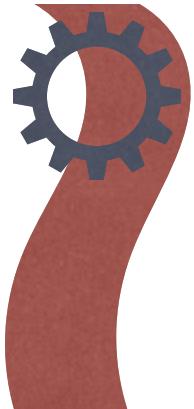


```
y = 3;  
z = 4;
```

Memory Order

x = 1;
y = 3;
z = 4;

Program Order



```
x = 1;
```



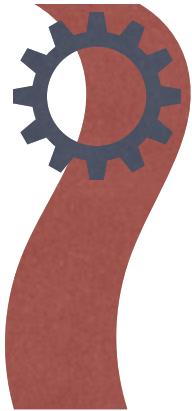
```
y = 3;  
z = 4;
```

Some re-orderings are allowed.

Memory Order

x = 1;
z = 4;
y = 3;

Program Order



```
x = 1;  
a = x;
```

read-after-write order
preserved within tasks

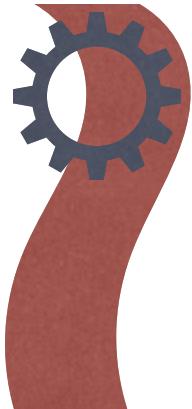


```
y = 3;  
z = 4;  
b = y;
```

Memory Order

x = 1;
z = 4;
y = 3;
b = y;
a = x;

Program Order



```
x = 1;  
x = 2;  
a = x;
```



```
y = 3;  
y = 4;  
b = y;
```

write-after-write order
preserved within tasks

Memory Order

x = 1;
y = 3;
x = 2;
y = 4;
b = y;
a = x;

Program Order

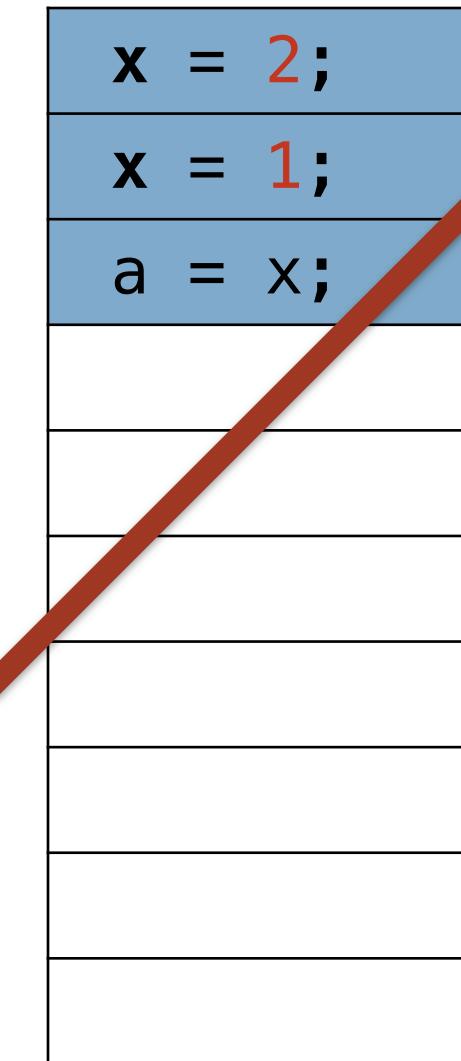


```
x = 1;  
x = 2;  
a = x;
```

Bad reordering! (ie, compiler bug)

Sequential programs must work
as if executed in program order

Memory Order



ASIDE: WEAK MEMORY CONSISTENCY

```
1 x starts at 0;  
  ...  
  if someOption then  
2    x = 2;  
    if someOtherOption then  
3    x = 3;  
4 return x;
```

ASIDE: WEAK MEMORY CONSISTENCY

```
1 x starts at 0;  
...  
...  
2 PUT 2 into x;  
...  
3 PUT 3 into x;  
4 GET x;
```

Chapel

result must be 3

OpenSHMEM

result could be 0, 2, or 3

MORE EXAMPLES: SHARED VARIABLES



COMPUTE

|

STORE

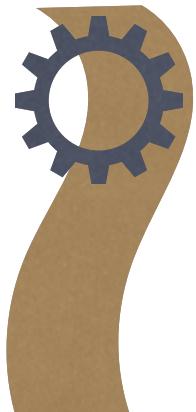
|

ANALYZE

Program Order

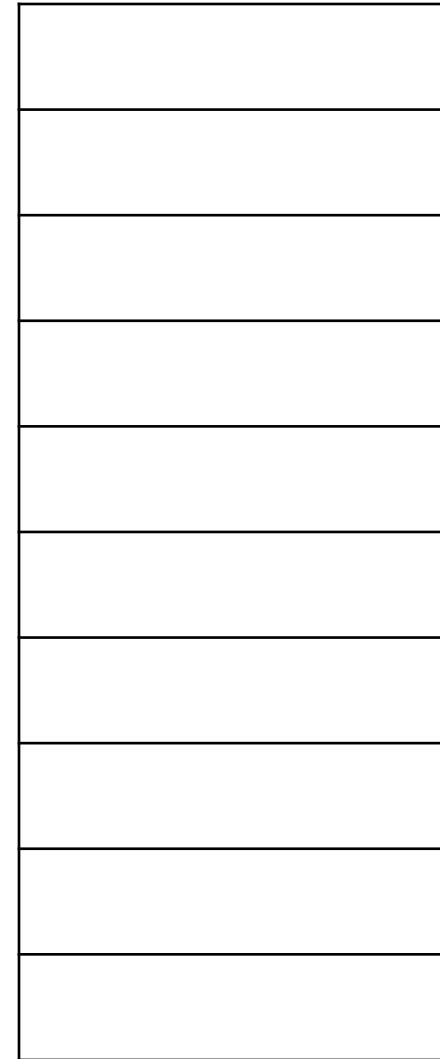


```
x = 0x1234;
```

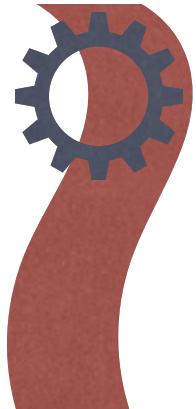


```
x = 0xABCD;  
c = x;
```

Memory Order



Program Order



```
x = 0x1234;
```

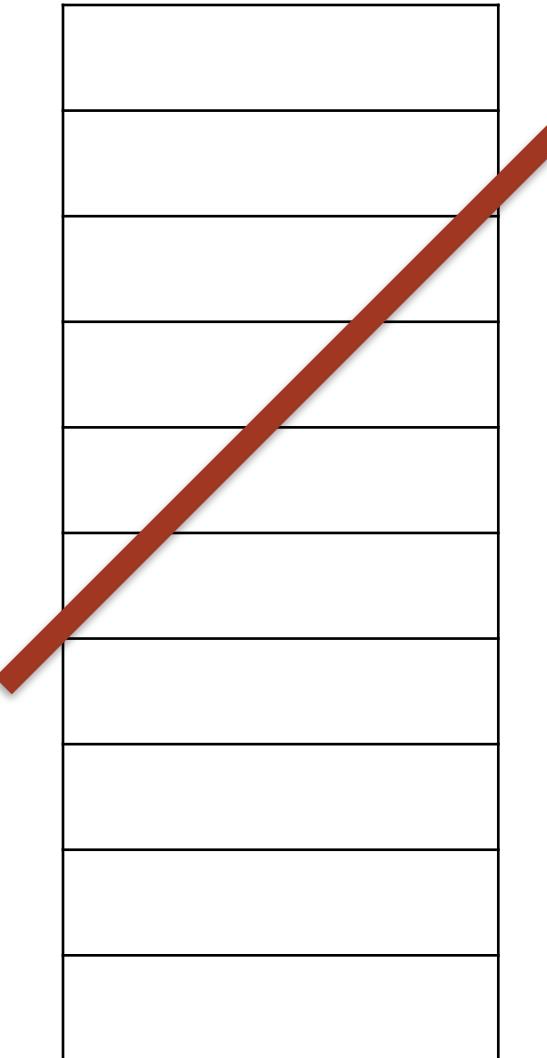


```
x = 0xABCD;  
c = x;
```

Bad program: Data Race.
No global order! This
outcome is possible:

c == 0xAB34

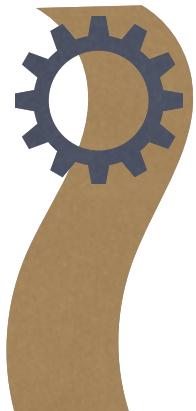
Memory Order



Program Order



```
x = 1;  
ok = 1;
```

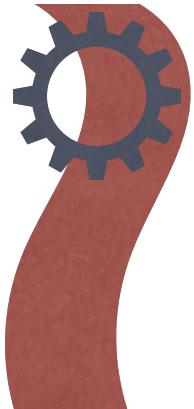


```
b = ok;  
c = x;
```

Memory Order



Program Order



```
x = 1;  
ok = 1;
```

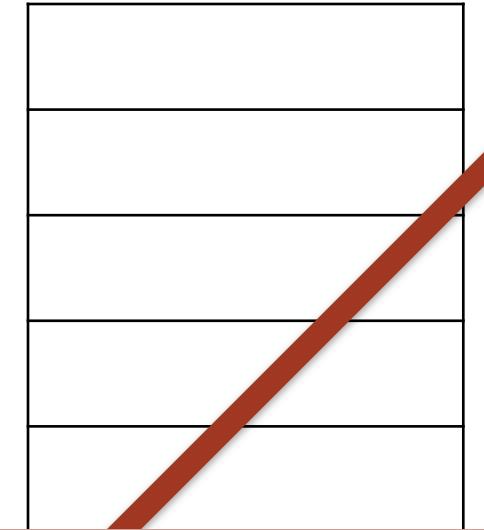


```
b = ok;  
c = x;
```

Bad program: Data Race.
No global order! This
outcome is possible:

```
b == 1  
c == 0
```

Memory Order



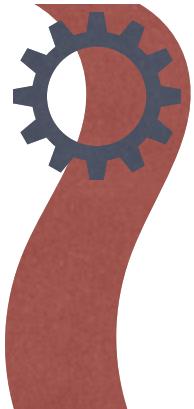
write behind could reorder:

```
ok = 1;  
x = 1;
```

read ahead could reorder:

```
c = x;  
b = ok;
```

Program Order



```
atomic A = 1;  
c = atomic A;
```



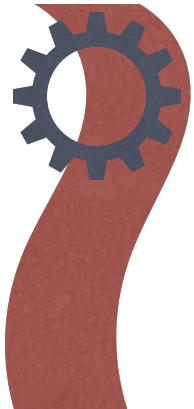
```
b = atomic A;  
atomic A = 2;
```

In Chapel and C++,
atomic vars default to
SC ordering which
includes both acquire
and release

Memory Order

```
atomic A = 1;  
c = atomic A;  
b = atomic A;  
atomic A = 2;
```

Program Order



```
atomic A = 1;  
c = atomic A;
```



```
b = atomic A;  
atomic A = 2;
```

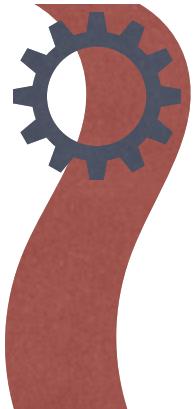
SC atomic vars create a global memory order.
 $c == 2$ $b == 0$ not possible e.g.



Memory Order

```
atomic A = 1;  
c = atomic A;  
b = atomic A;  
atomic A = 2;
```

Program Order



```
atomic A = 1;  
c = atomic A;
```

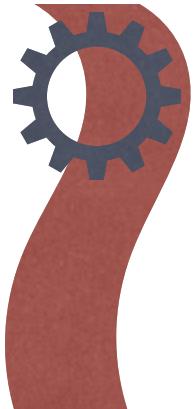


```
b = atomic A;  
atomic A = 2;
```

Memory Order

b = atomic A;
atomic A = 1;
c = atomic A;
atomic A = 2;

Program Order



```
x = 2;  
atomic A = 1;
```



```
b = atomic A;  
c = x;
```

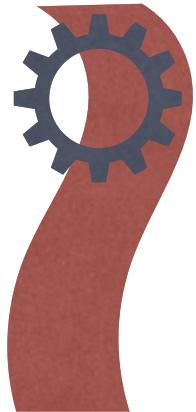
SC atomic ops constrain
the code around them

b == 1 implies
c == 2

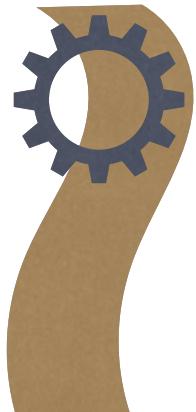
Memory Order

```
x = 2;  
atomic A = 1;  
b = atomic A;  
c = x;
```

Program Order



```
atomic A = 1;  
atomic B = 2;
```



```
c = -1;  
if (atomic B == 2)  
    c = atomic A;
```

SC atomic vars create a global total memory order.

c == -1 || c == 1

Memory Order

atomic A = 1;
c = -2;
atomic B = 2;
(atomic B == 2);
c = atomic A;



Program Order



```
relaxed atomic A = 1;  
relaxed atomic B = 2;
```



```
waitFor	atomic B == 2);  
c = atomic A;
```

Is an outcome of $c == 0$ possible?

Program Order



```
relaxed atomic A = 1;  
relaxed atomic B = 2;
```



```
waitFor	atomic B == 2);  
c = atomic A;
```

Is an outcome of $c == 0$ possible?

Yes. While atomic vars cannot create race conditions, relaxed atomics don't create a total order.

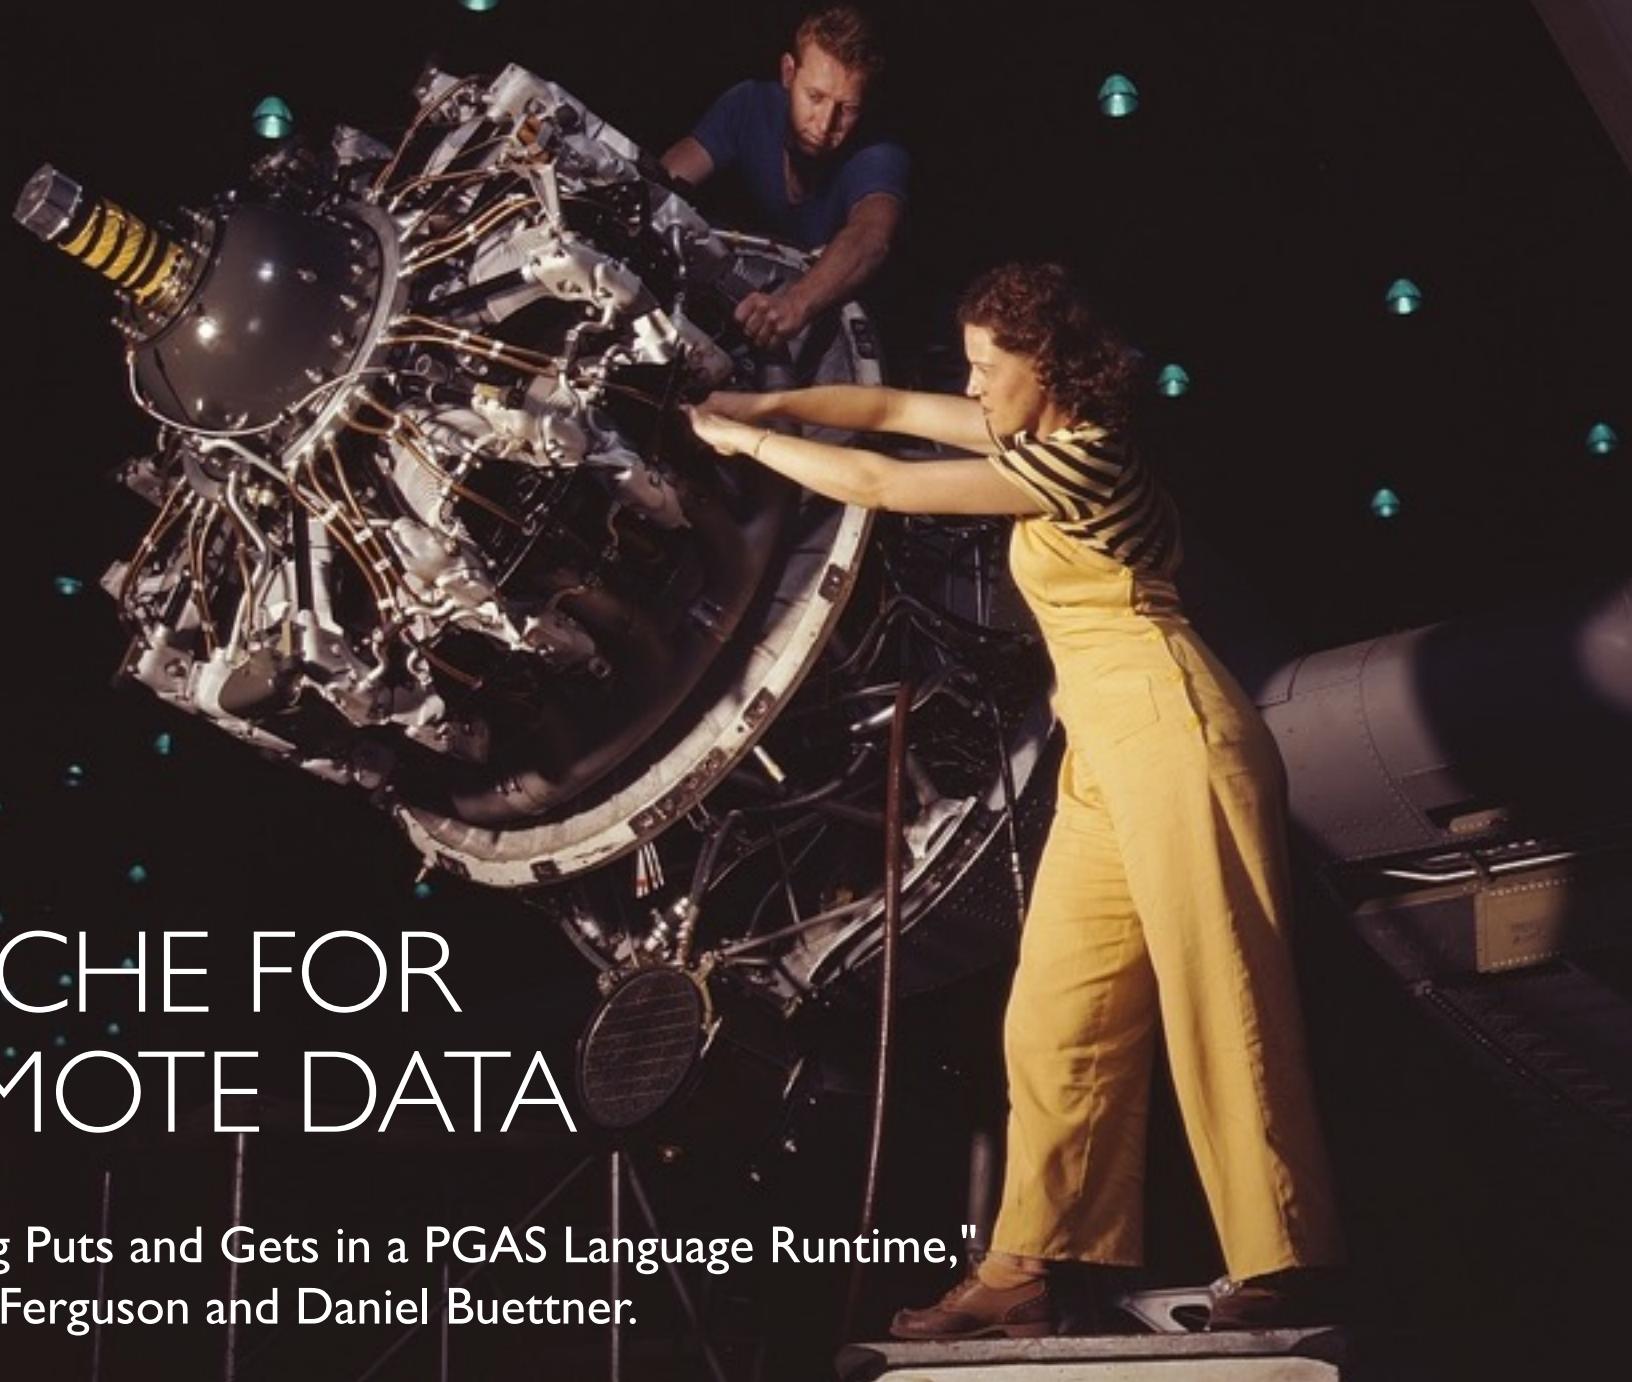
e.g. write behind could reorder:

```
relaxed atomic B = 2;  
relaxed atomic A = 1;
```



CACHE FOR REMOTE DATA

"Caching Puts and Gets in a PGAS Language Runtime,"
Michael Ferguson and Daniel Buettner.



CACHE FOR REMOTE DATA

- **Goal: communication aggregation and overlap**
- **Bonus points: avoiding repeated communication**
- **Software cache in Chapel's runtime**
- **One cache per pthread**
- **Write-back cache with dirty bits**

CACHE COHERENCY

- Simple, local coherency
- Discard all cached data on *acquire*
- Wait for pending operations on a *release*
- Strategy used in related work with UPC

CACHE FEATURES

	<i>Overlap</i>		<i>Aggregation</i>	
	<i>GET</i>	<i>PUT</i>	<i>GET</i>	<i>PUT</i>
<i>Do PUTs in background</i>		X		
<i>Start one PUT per contiguous written region</i>				X
<i>Round GETs up to 64-byte cache lines</i>			X	
<i>Sequential read-ahead</i>	X		X	
<i>Programmer-provided prefetch hints*</i>	X			

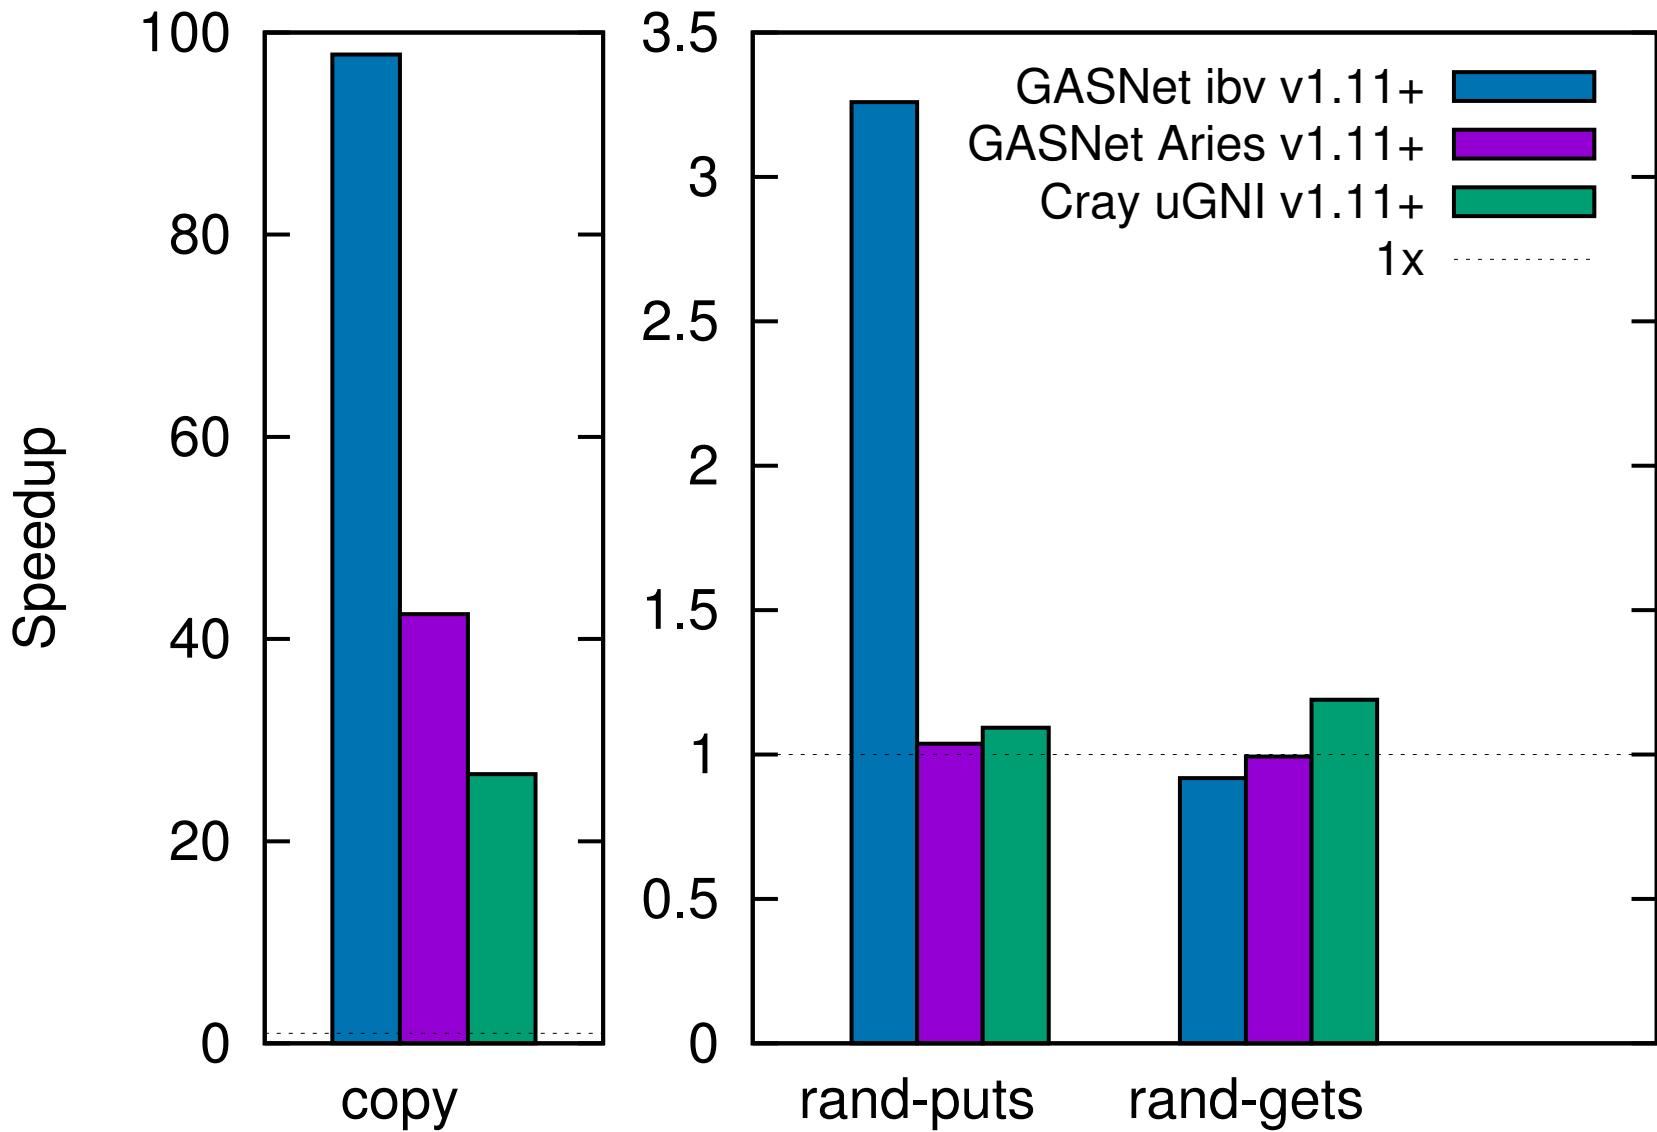


PERFORMANCE

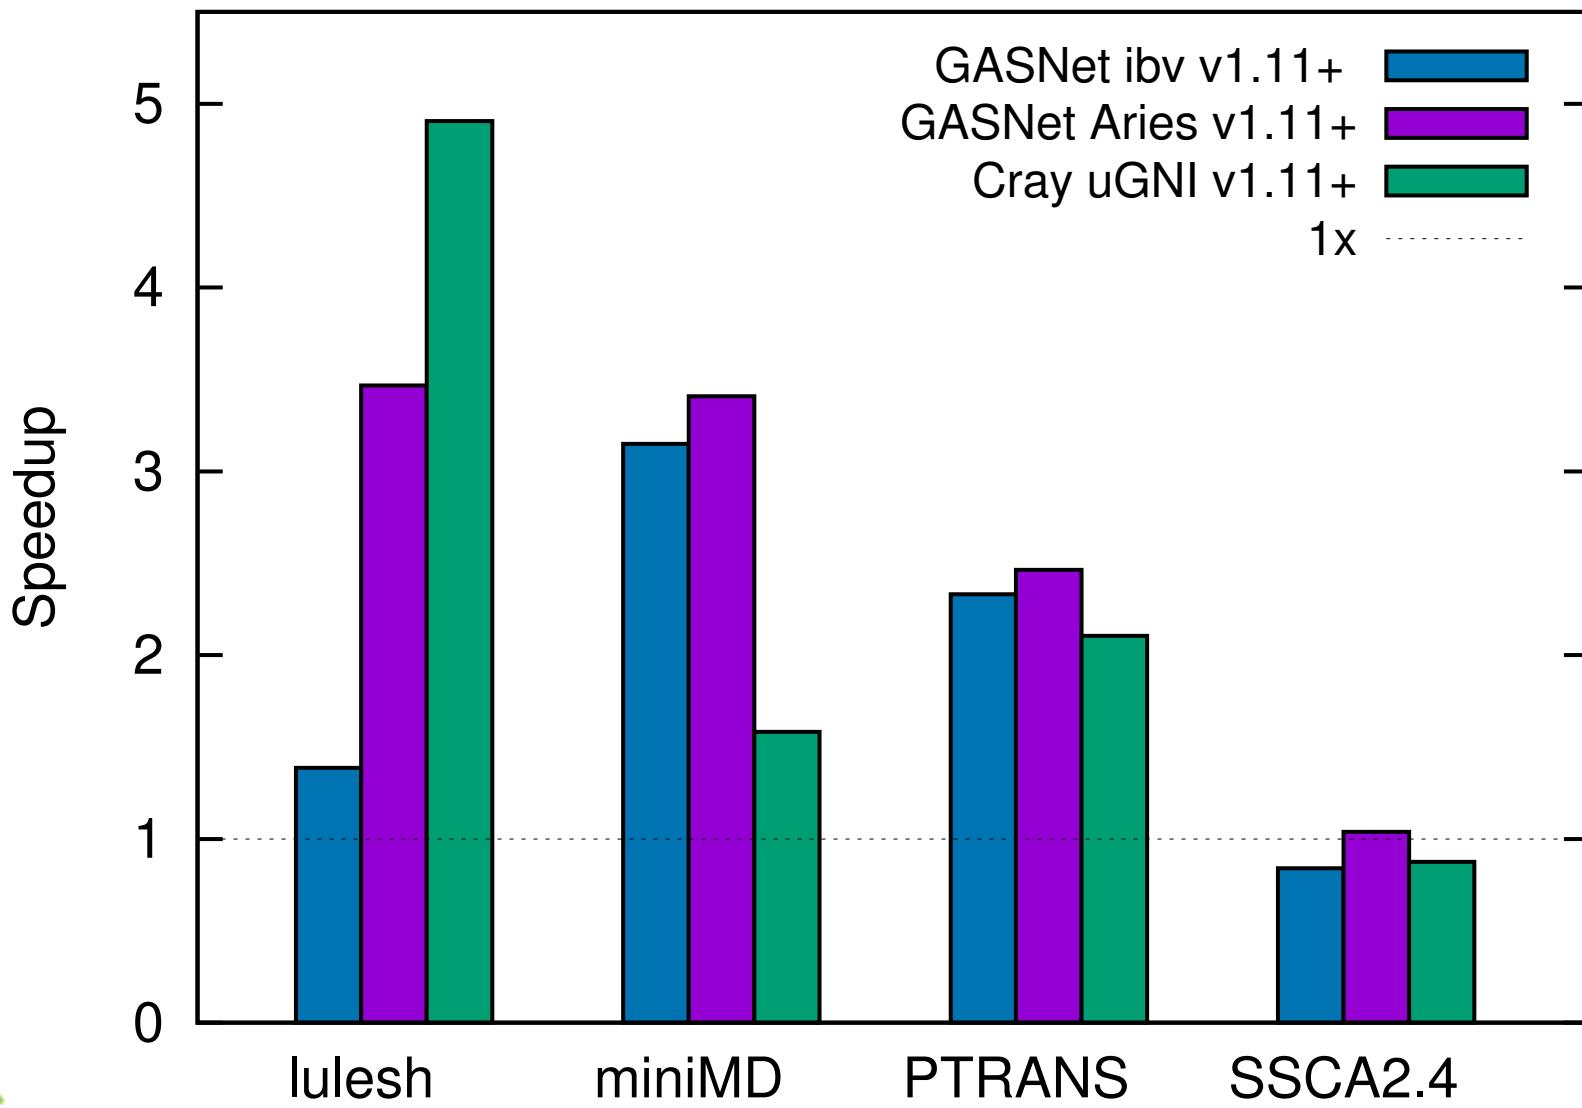


San Diego Air and Space Museum

SYNTHETIC BENCHMARKS



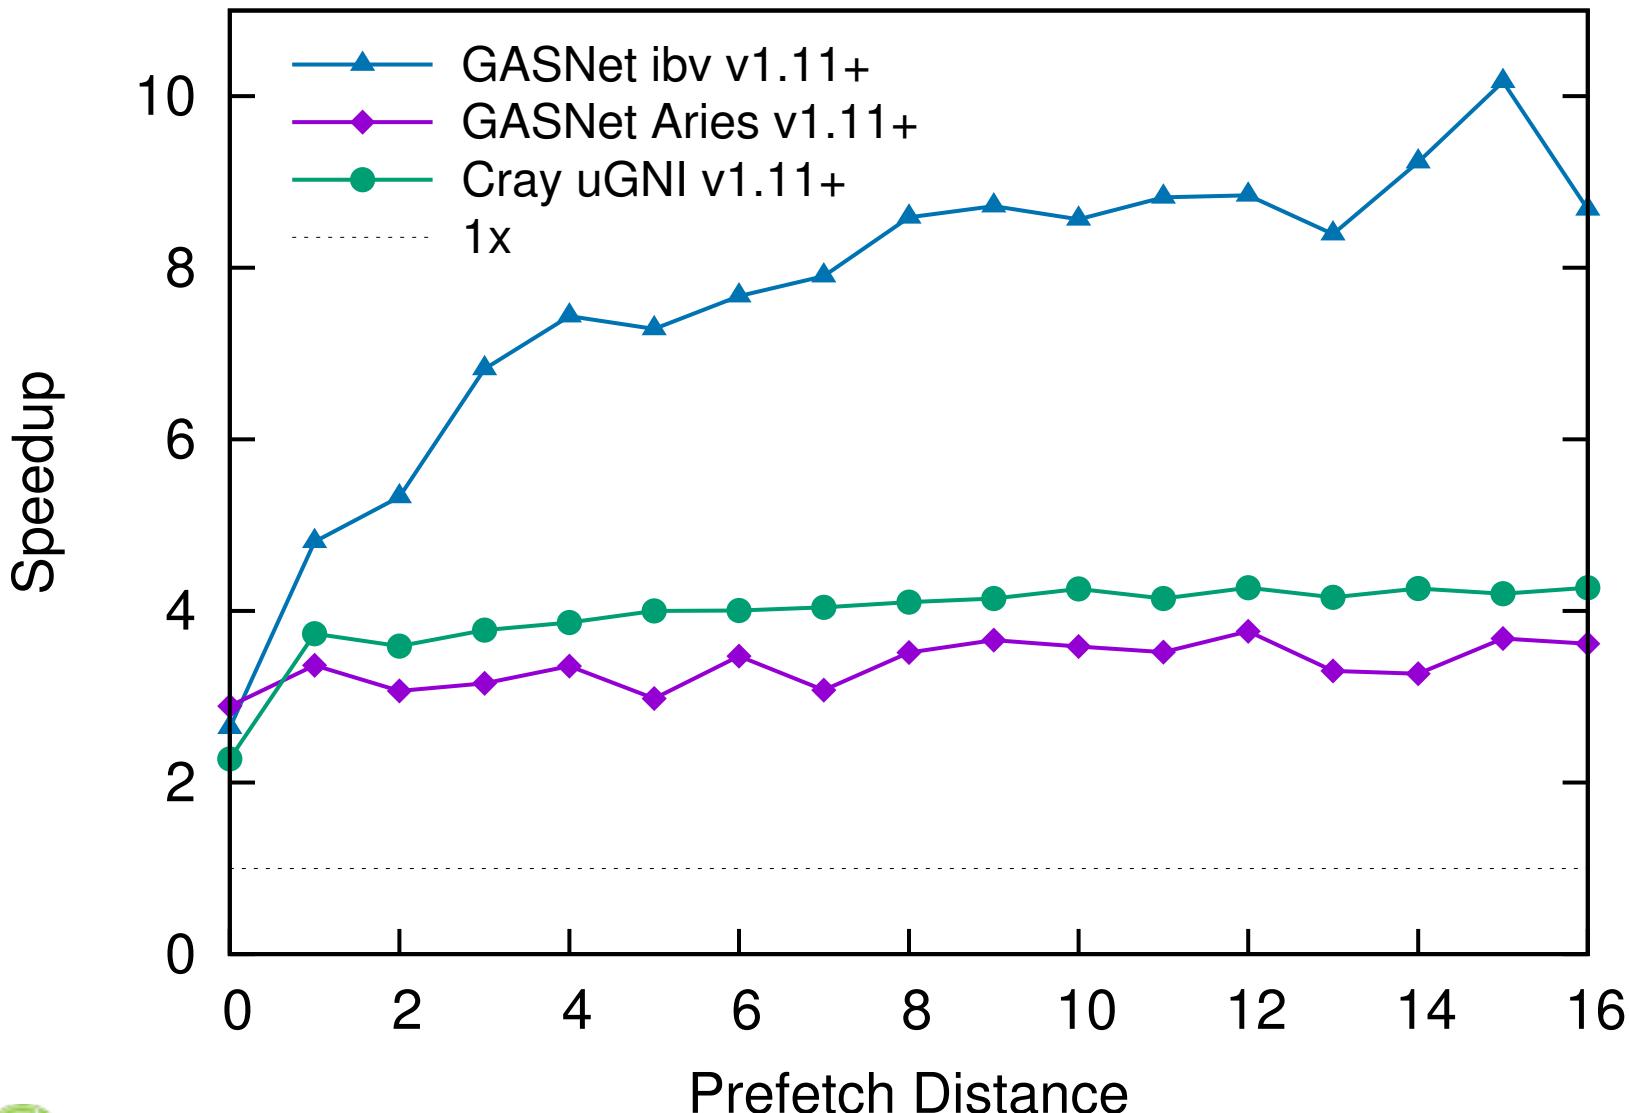
APPLICATION BENCHMARKS



PREFETCH EXAMPLE

```
var A:[1..n] int;
on Locales[1] {
    var sum:int;
    // Optional warm up
    for i in 1..k do prefetch(A[f(i)]);
    for i in 1..n {
        if i+k <= n then prefetch(A[f(i+k)]);
        sum += A[f(i)]
    }
}
```

PREFETCH EXAMPLE

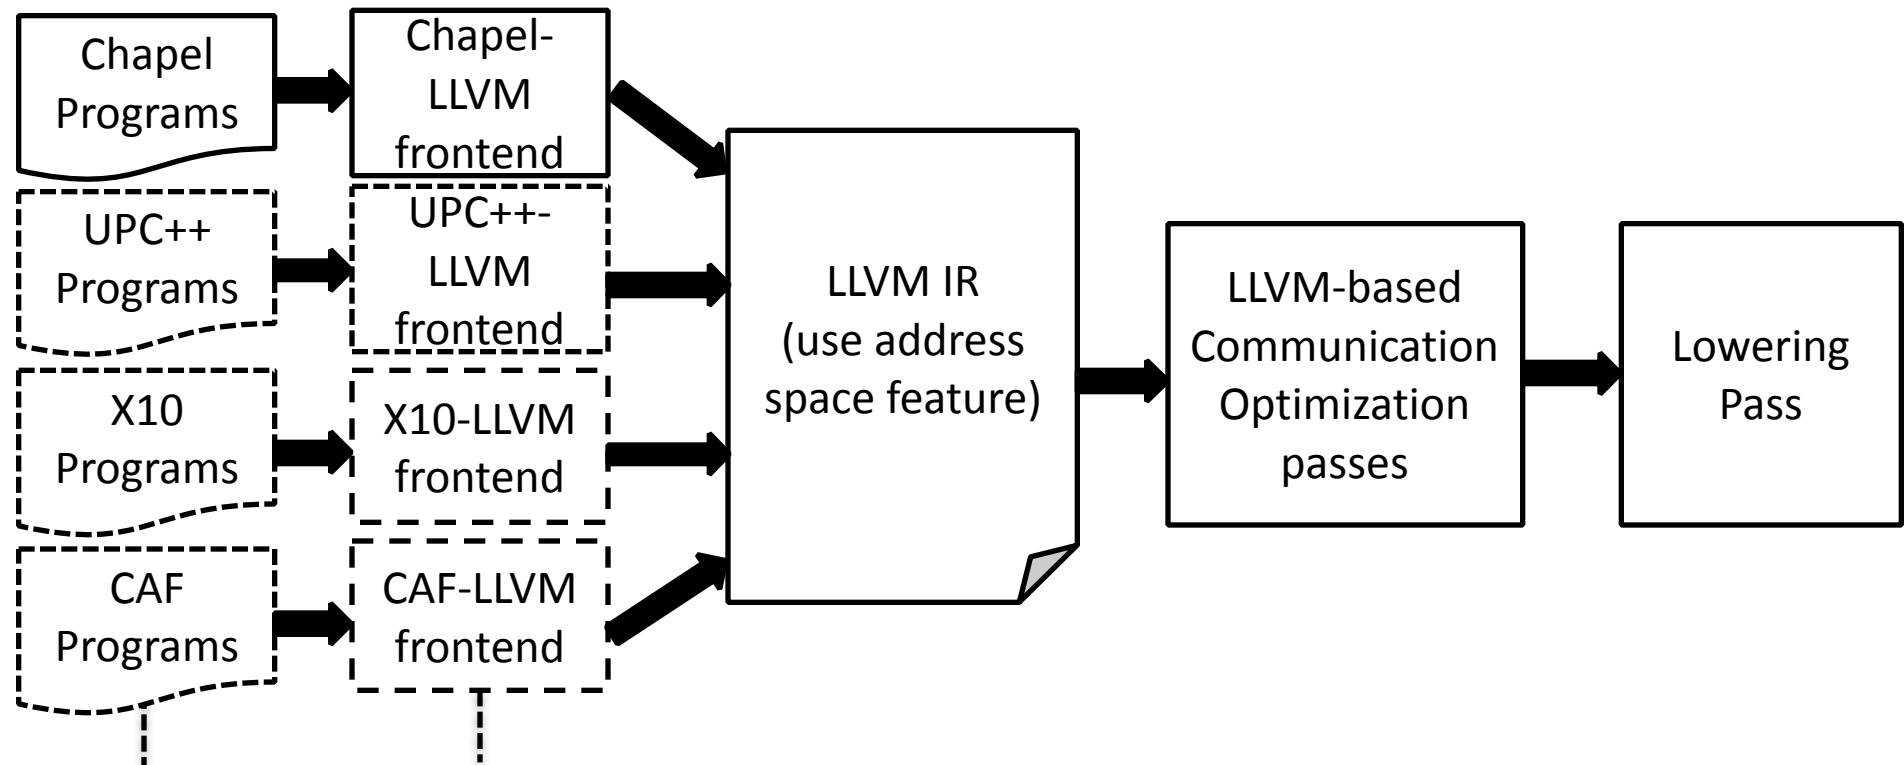




OPTIMIZING COMMUNICATION WITH LLVM

"LLVM-based Communication Optimizations for PGAS Programs"
Akihito Hayashi, Jisheng Zhao, Michael Ferguson, Vivek Sarkar

THE VISION: SHARED PGAS OPTIMIZATION PASSES

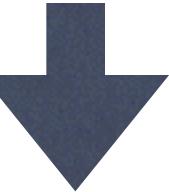


EXAMPLE

```
// x is remote  
var sum = 0;  
for i in 1..100 {  
    sum += get(x);  
}
```

```
// x is possibly remote  
var sum = 0;  
for i in 1..100 {  
    %l = get(x);  
    sum += %l;  
}
```

TO GLOBAL
MEMORY



```
var sum = 0;  
for i in 1..100 {  
    %l = load <100> %x  
    sum += %l;  
}
```

EXISTING LLVM
OPTIMIZATION
LICM

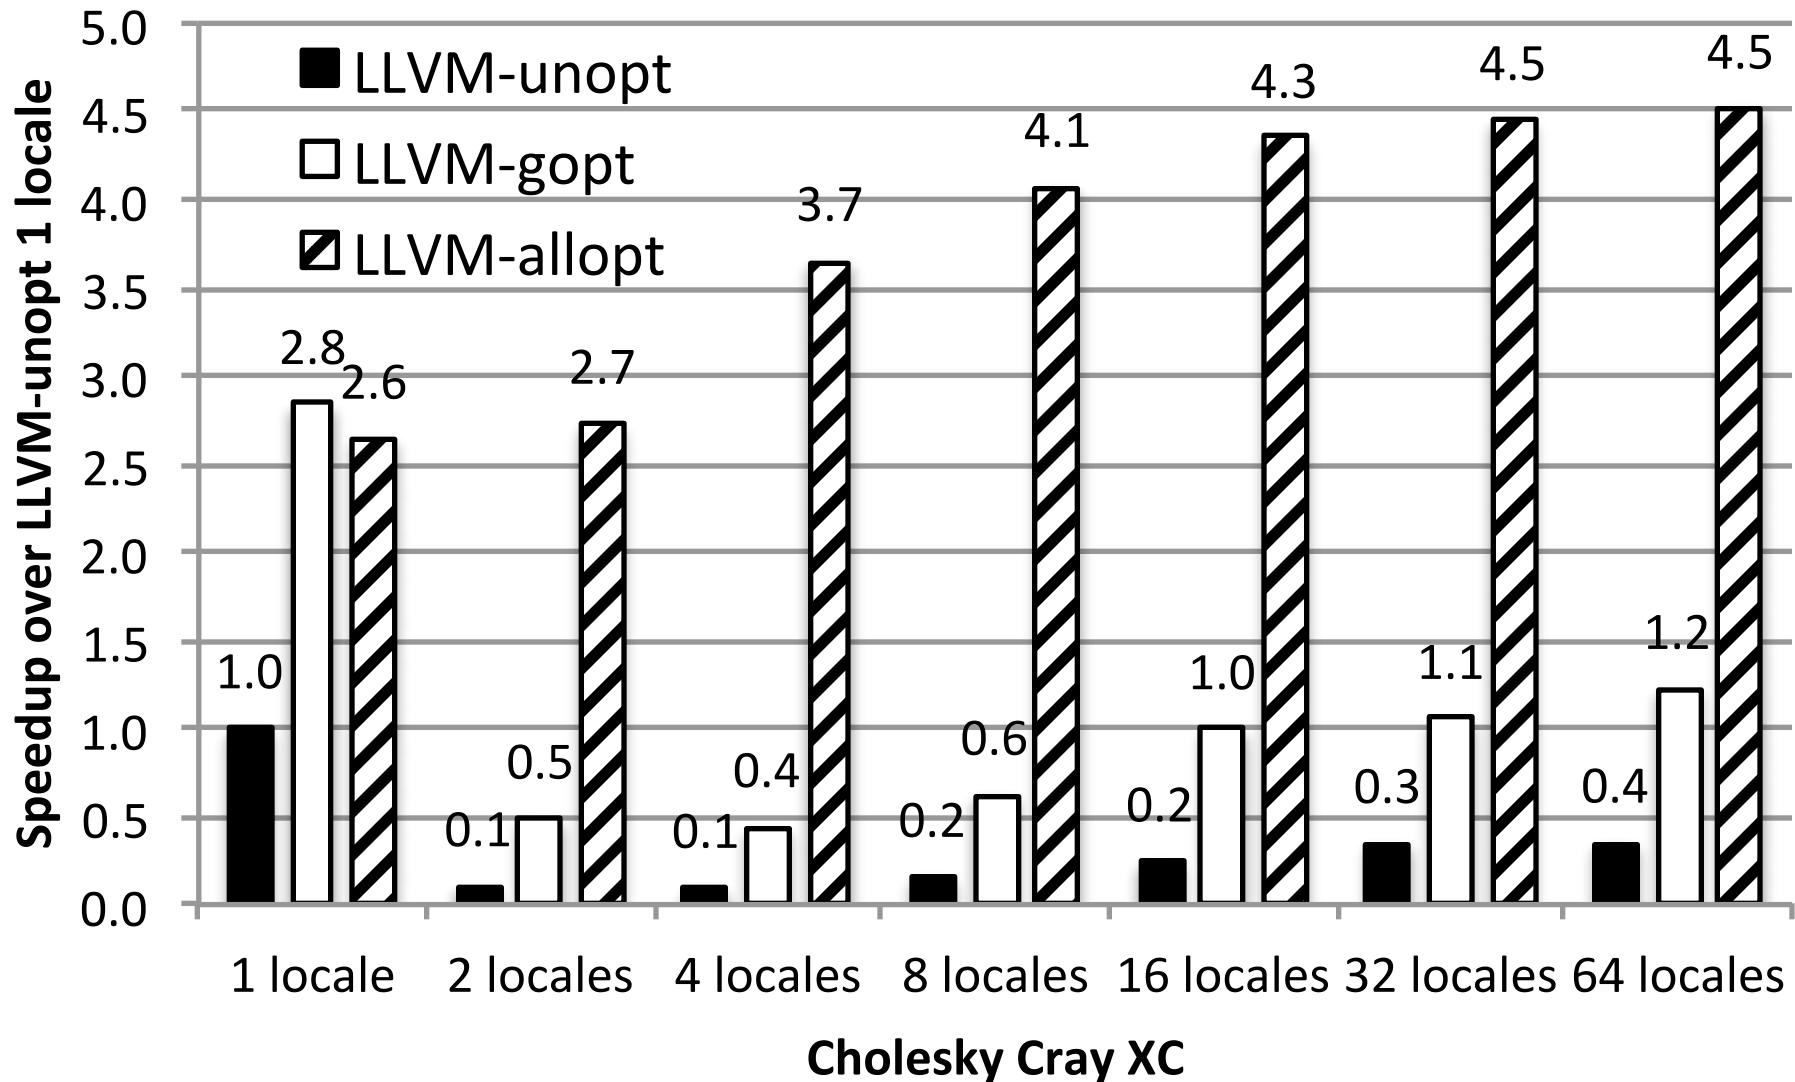
```
var sum = 0;  
%l = get(x);  
for i in 1..100 {  
    sum += %l;  
}
```

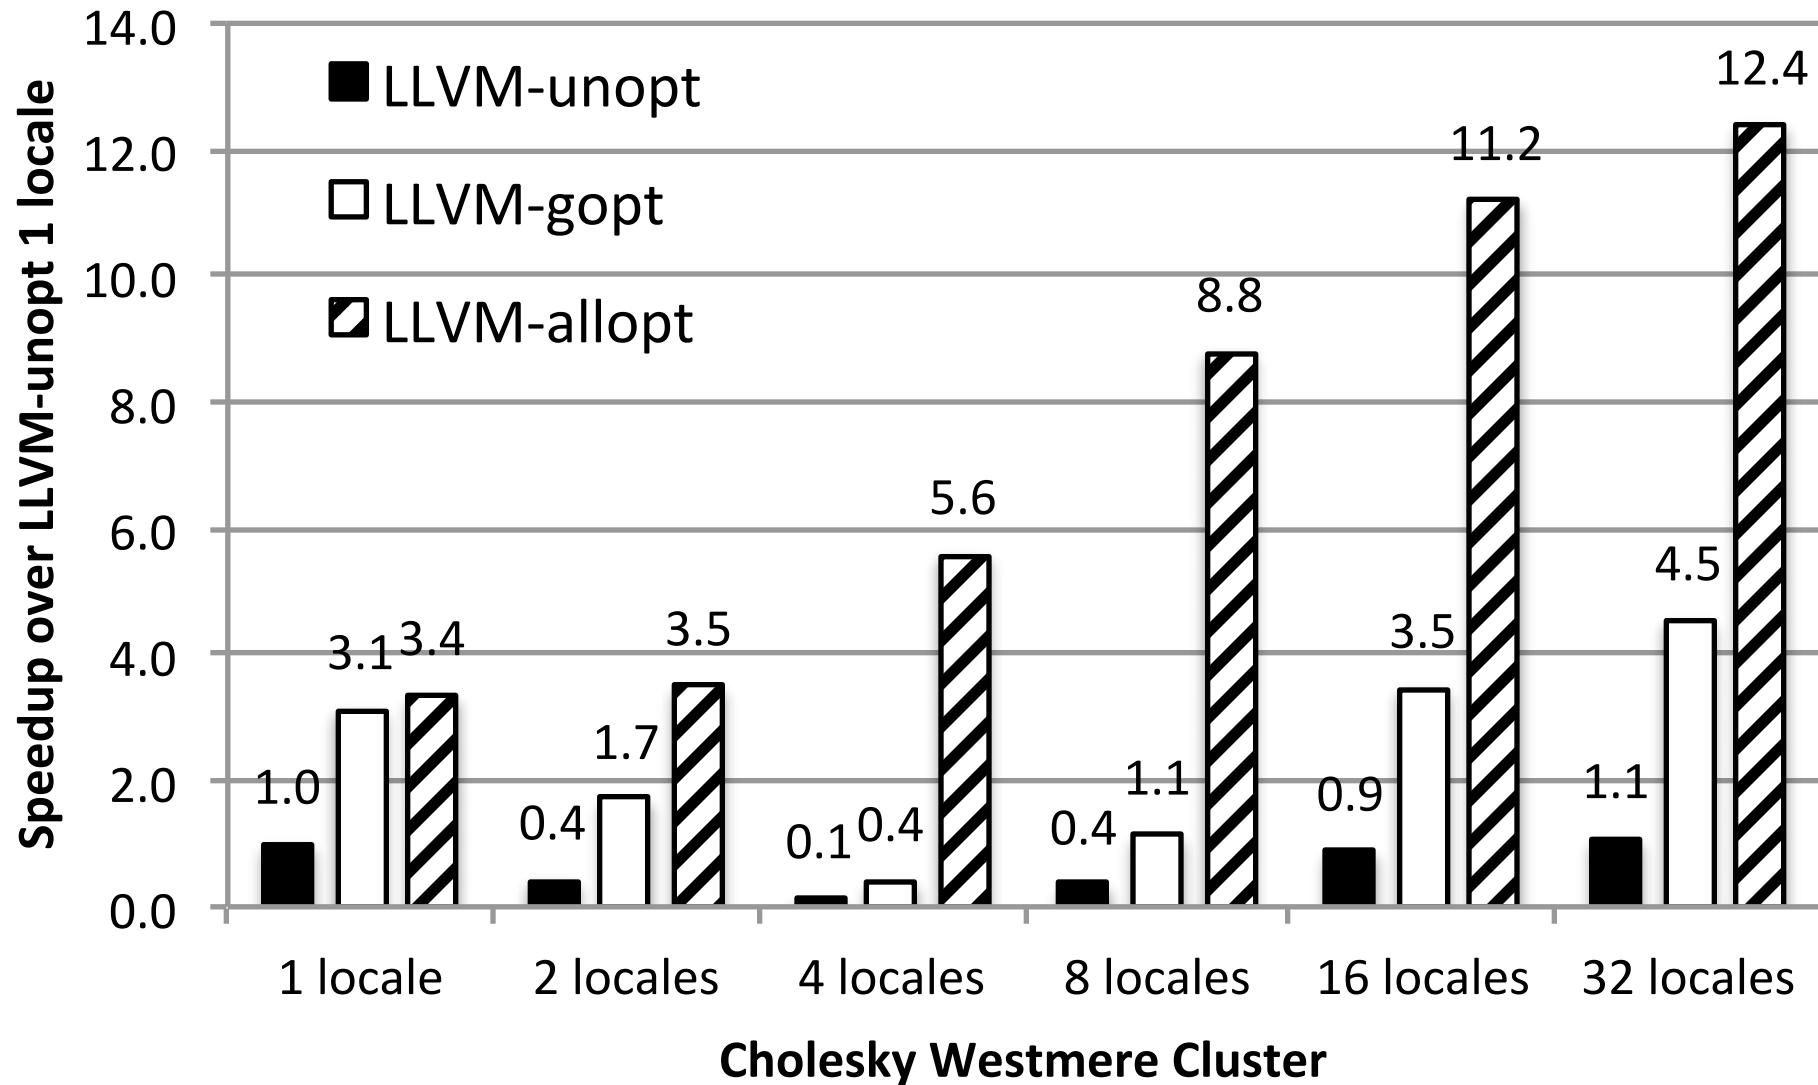
TO DISTRIBUTED
MEMORY

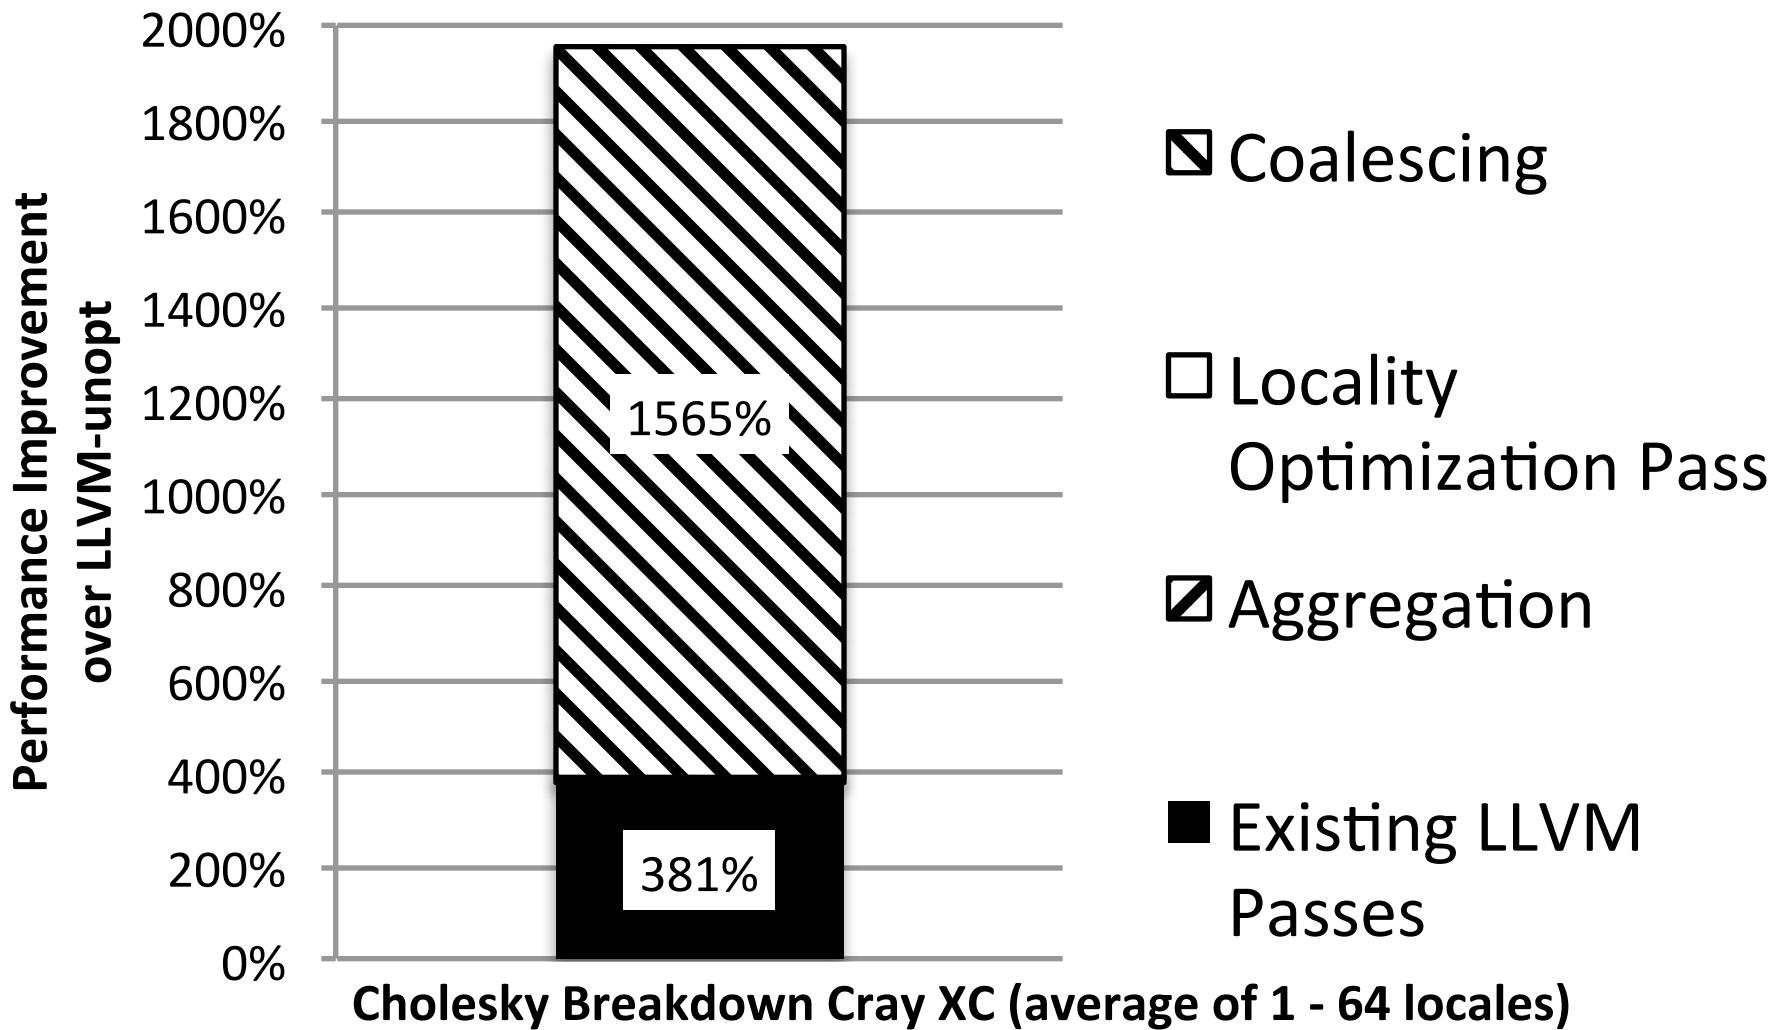
```
// existing LLVM opt  
var sum = 0;  
%l = load <100> %x  
for i in 1..100 {  
    sum += %rl;  
}
```

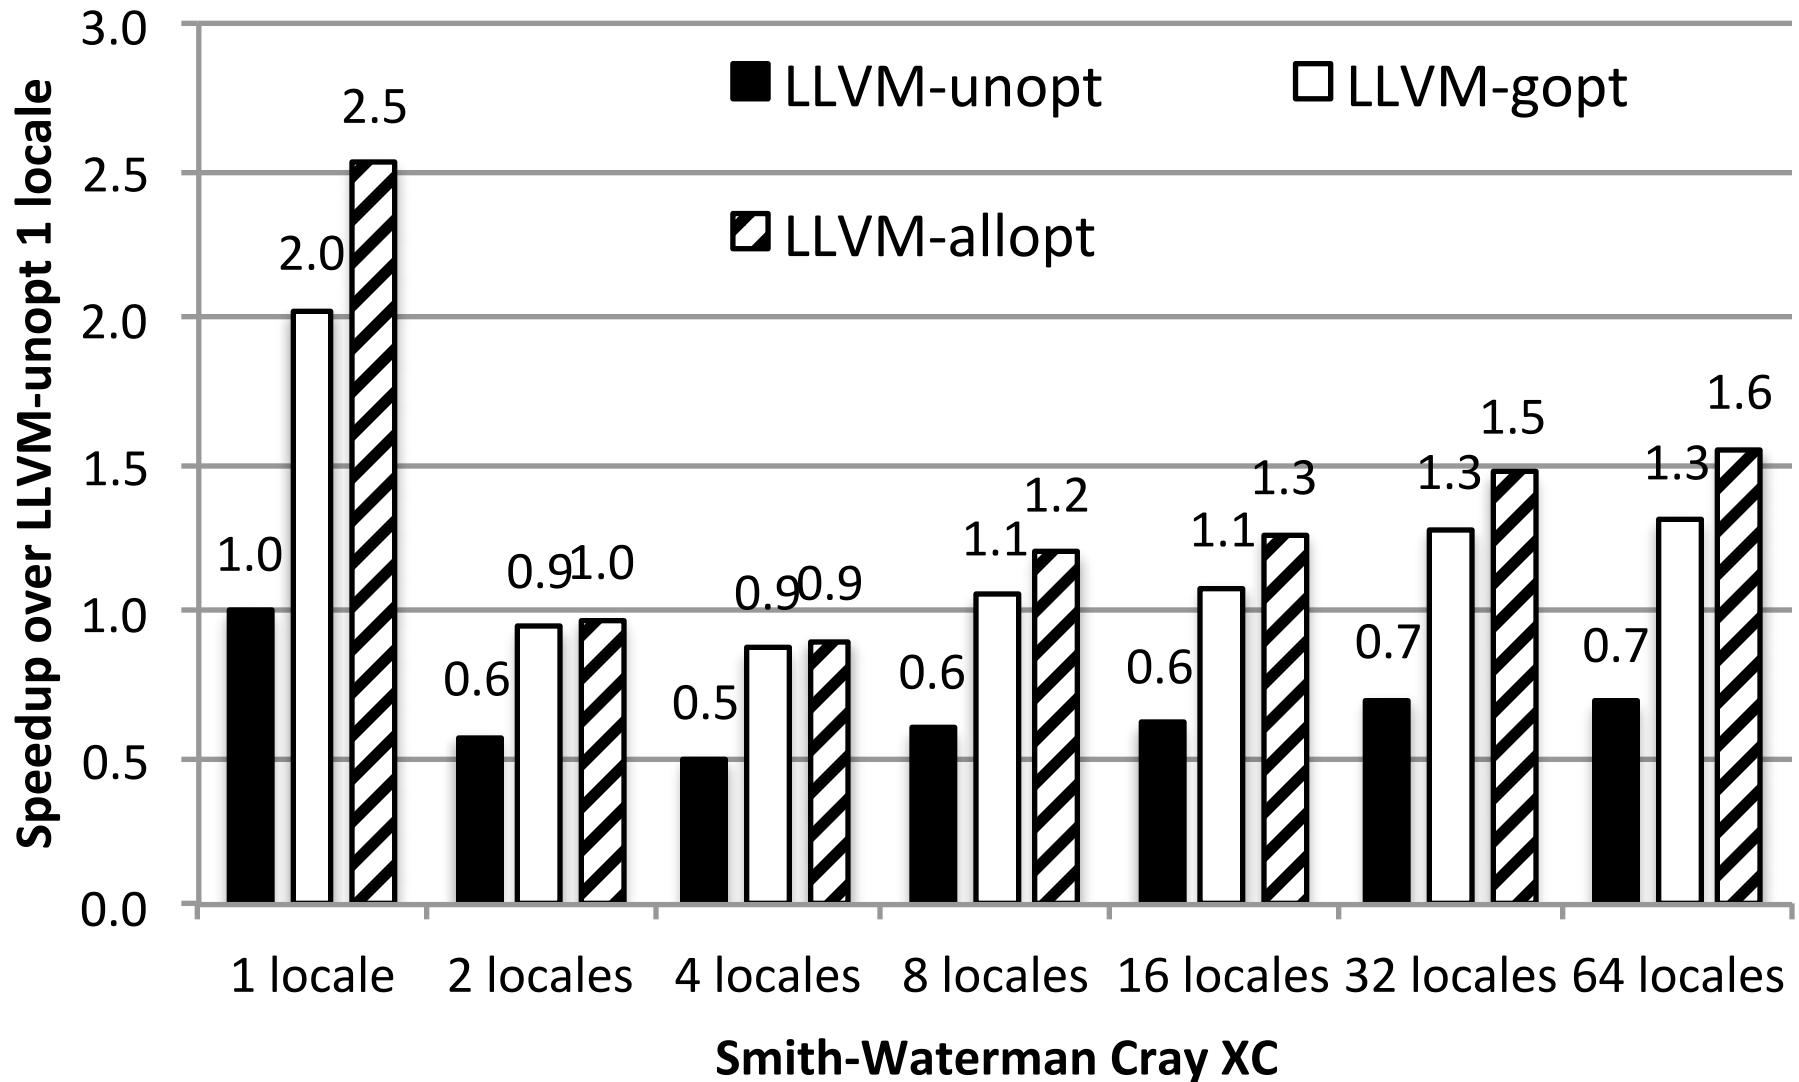
$\text{load } <100> \%x = \text{load i64 addrspace}(100)^* \%x$

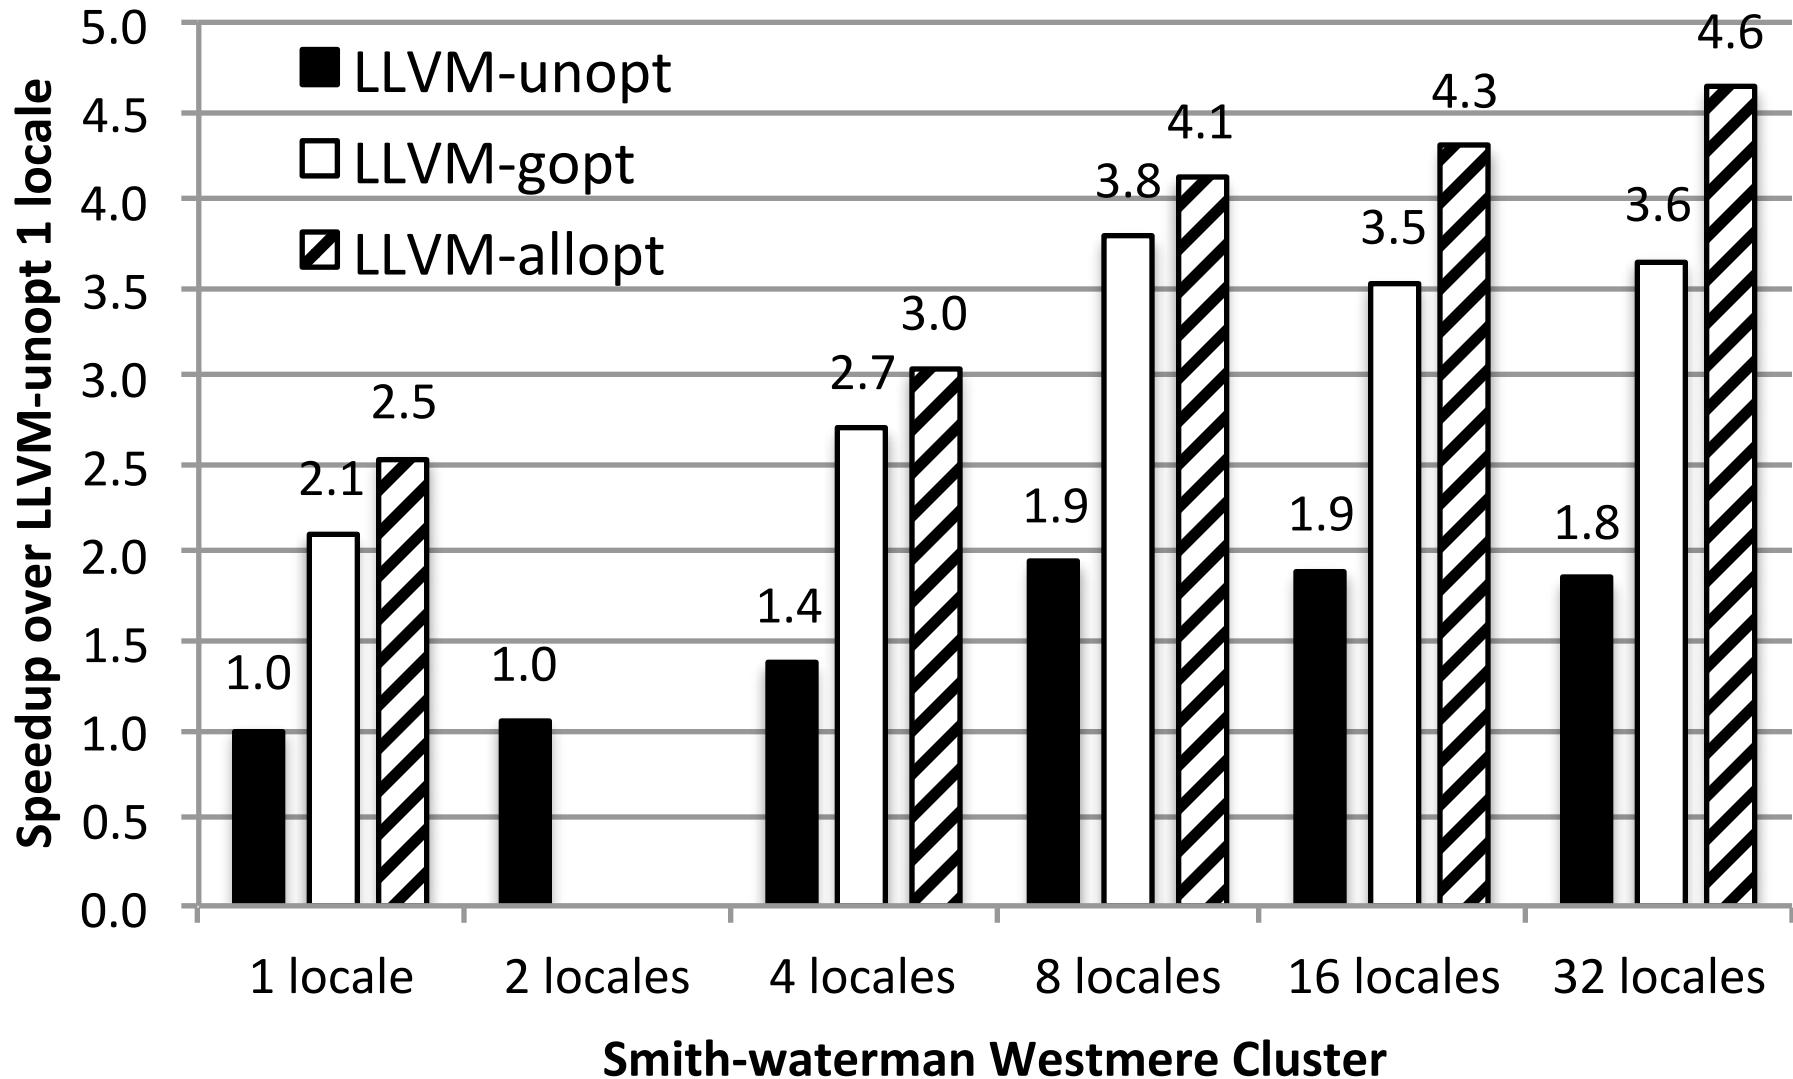


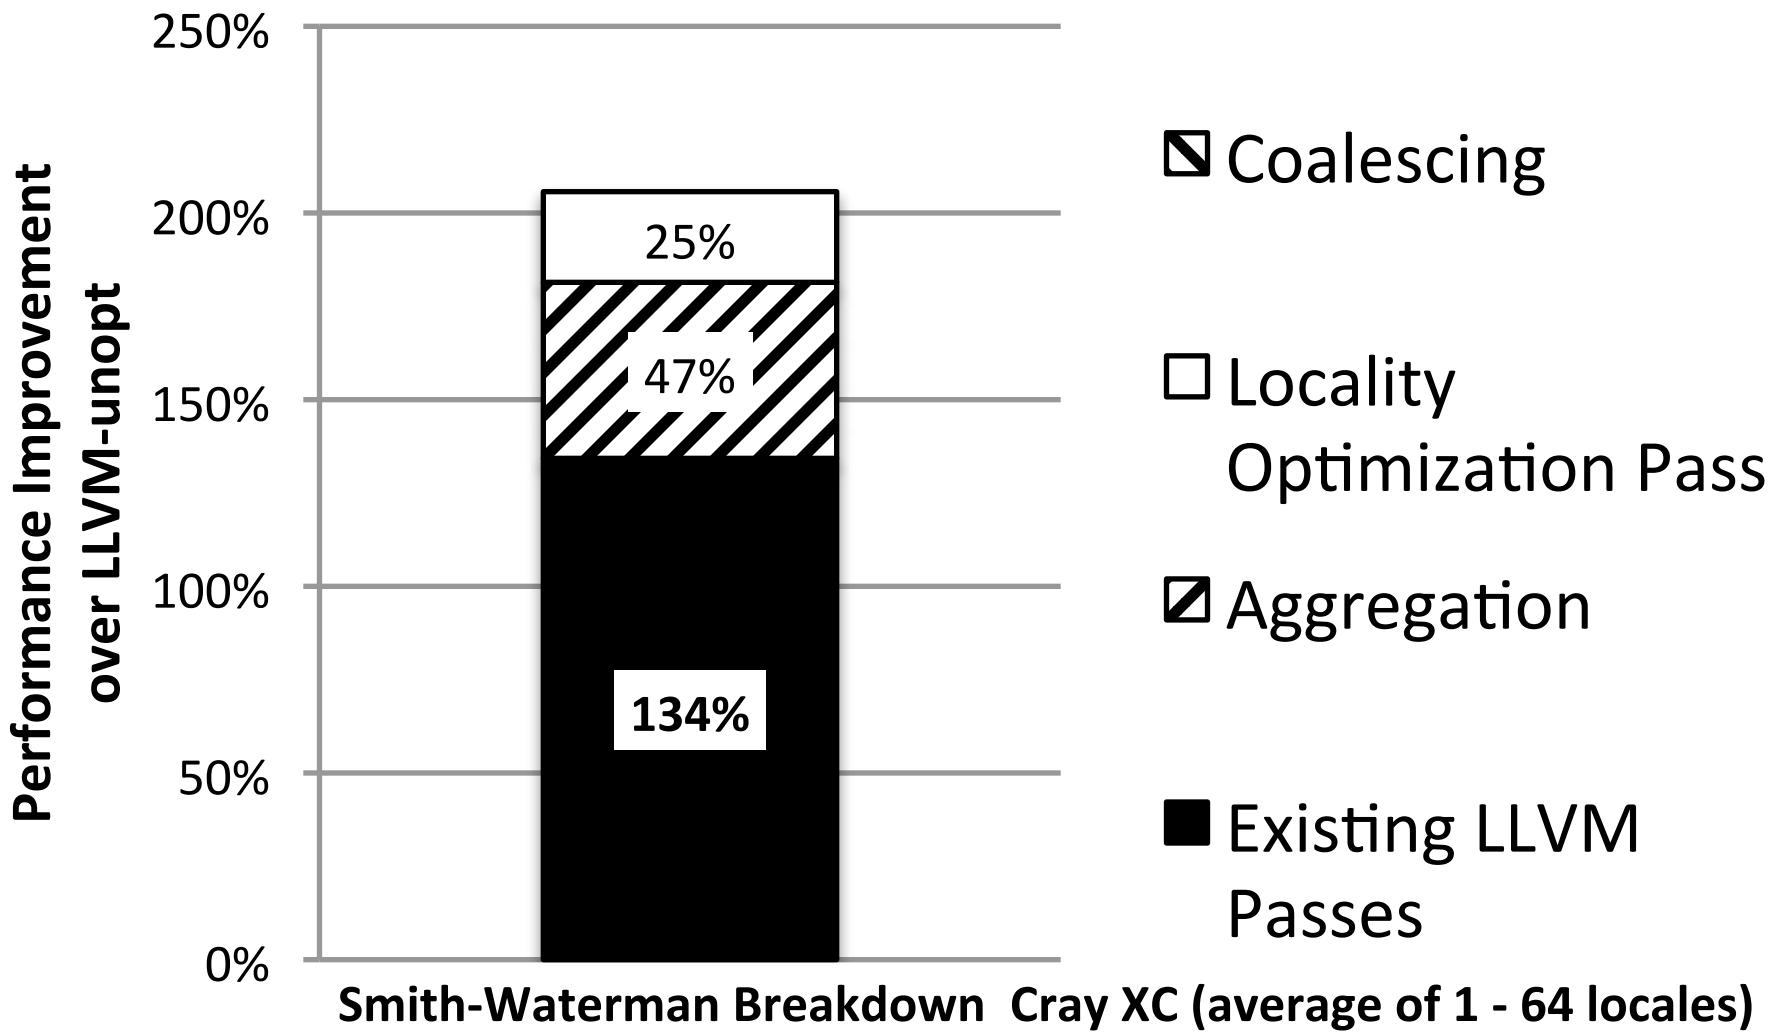


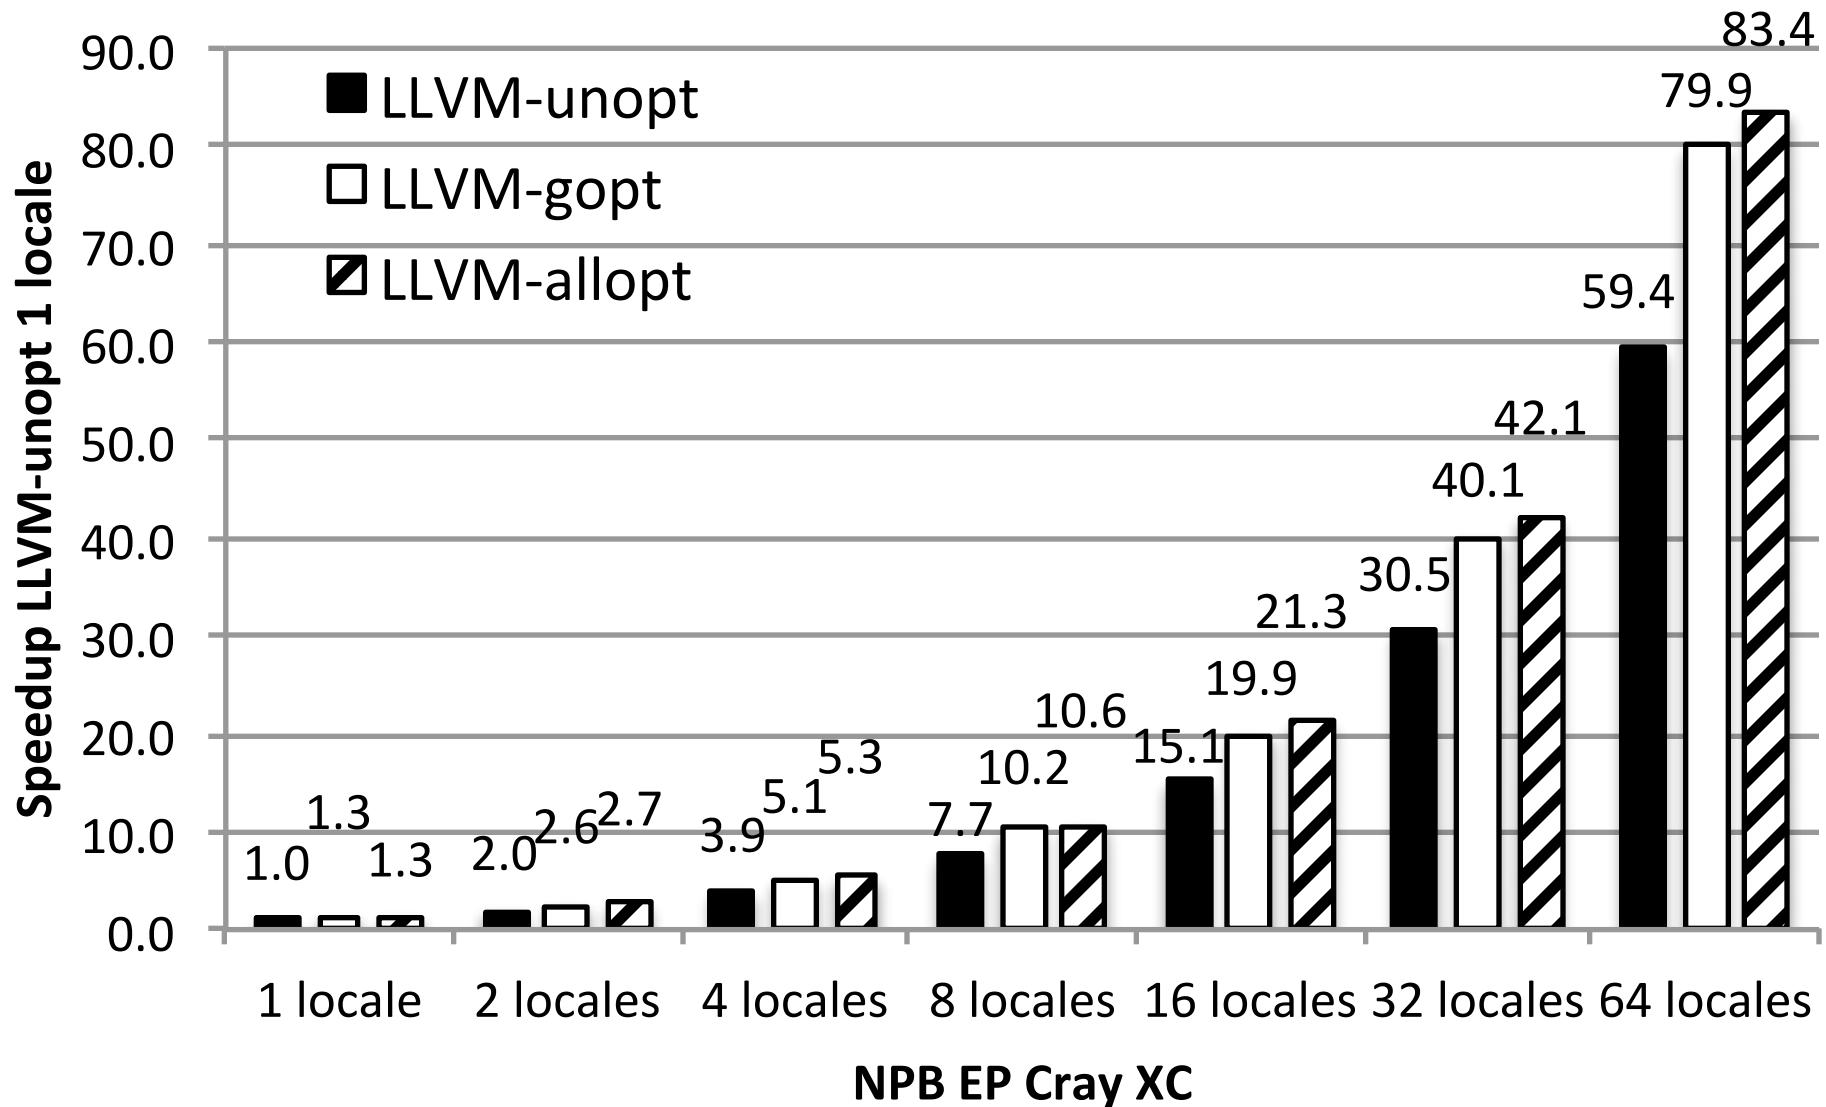


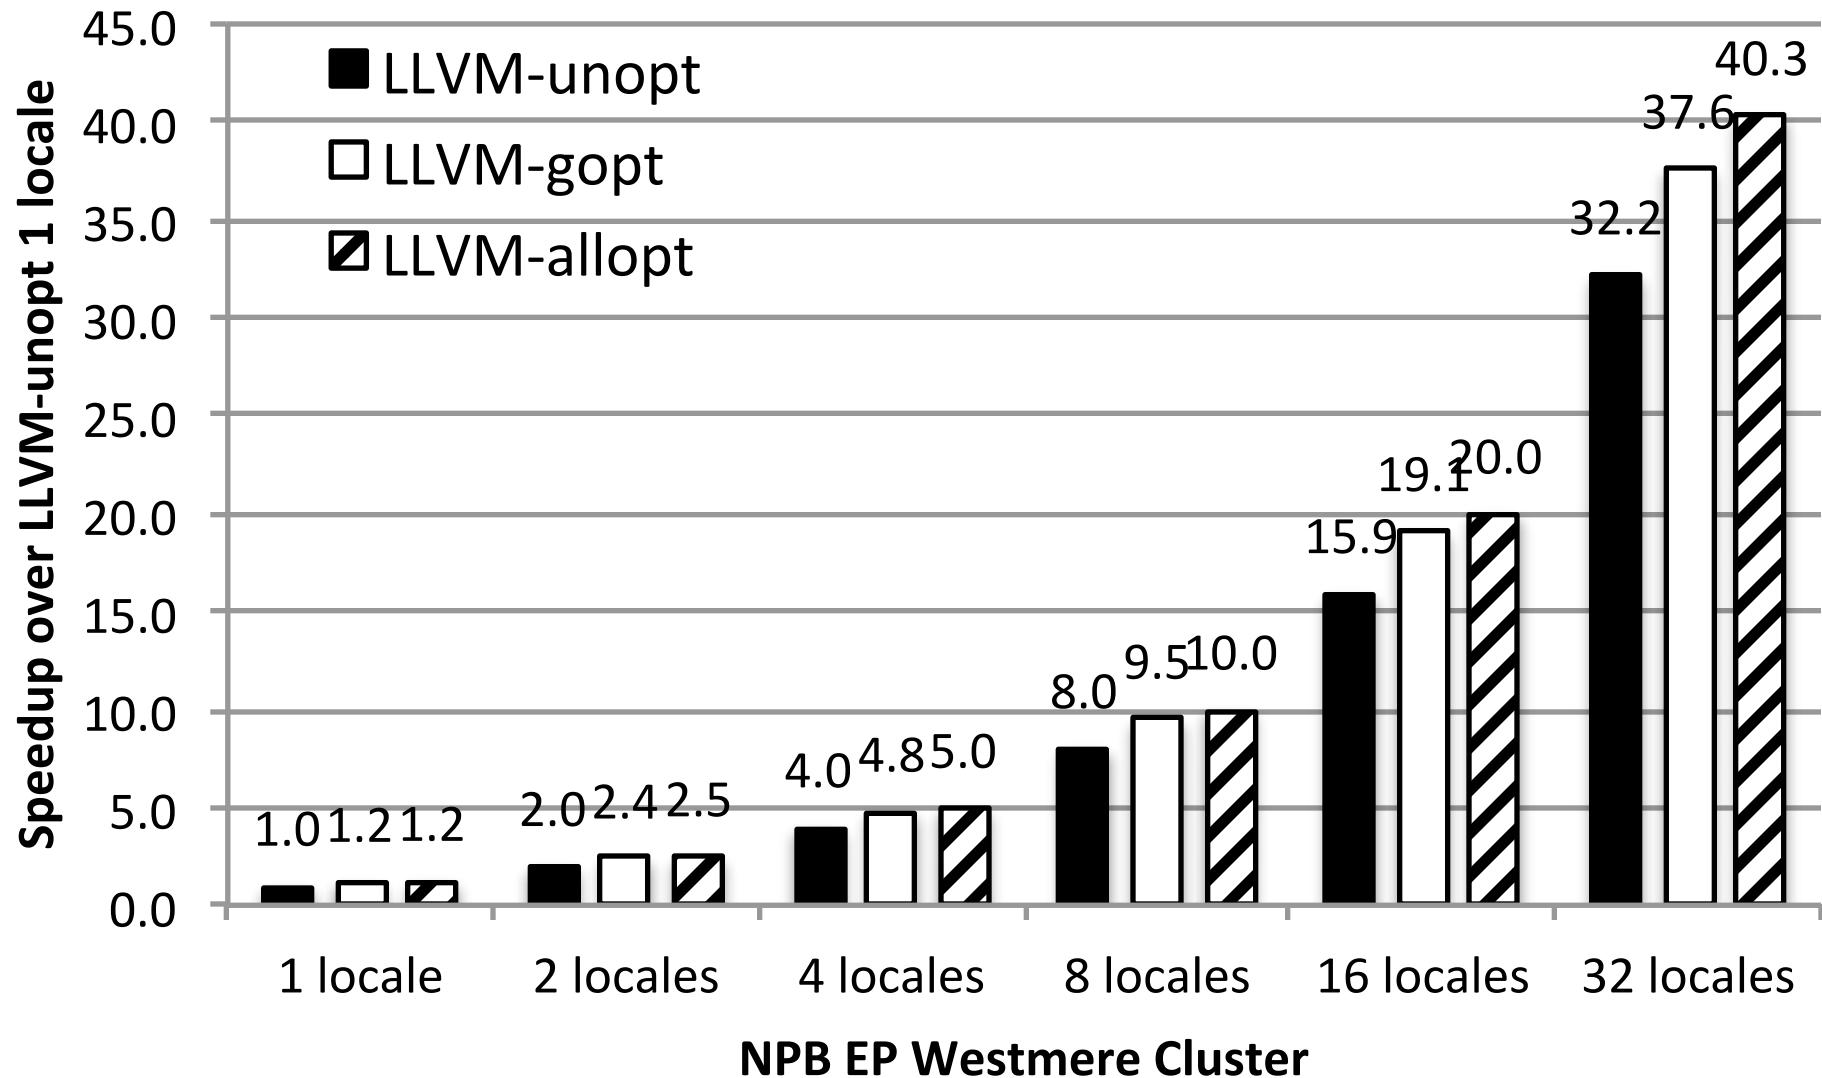


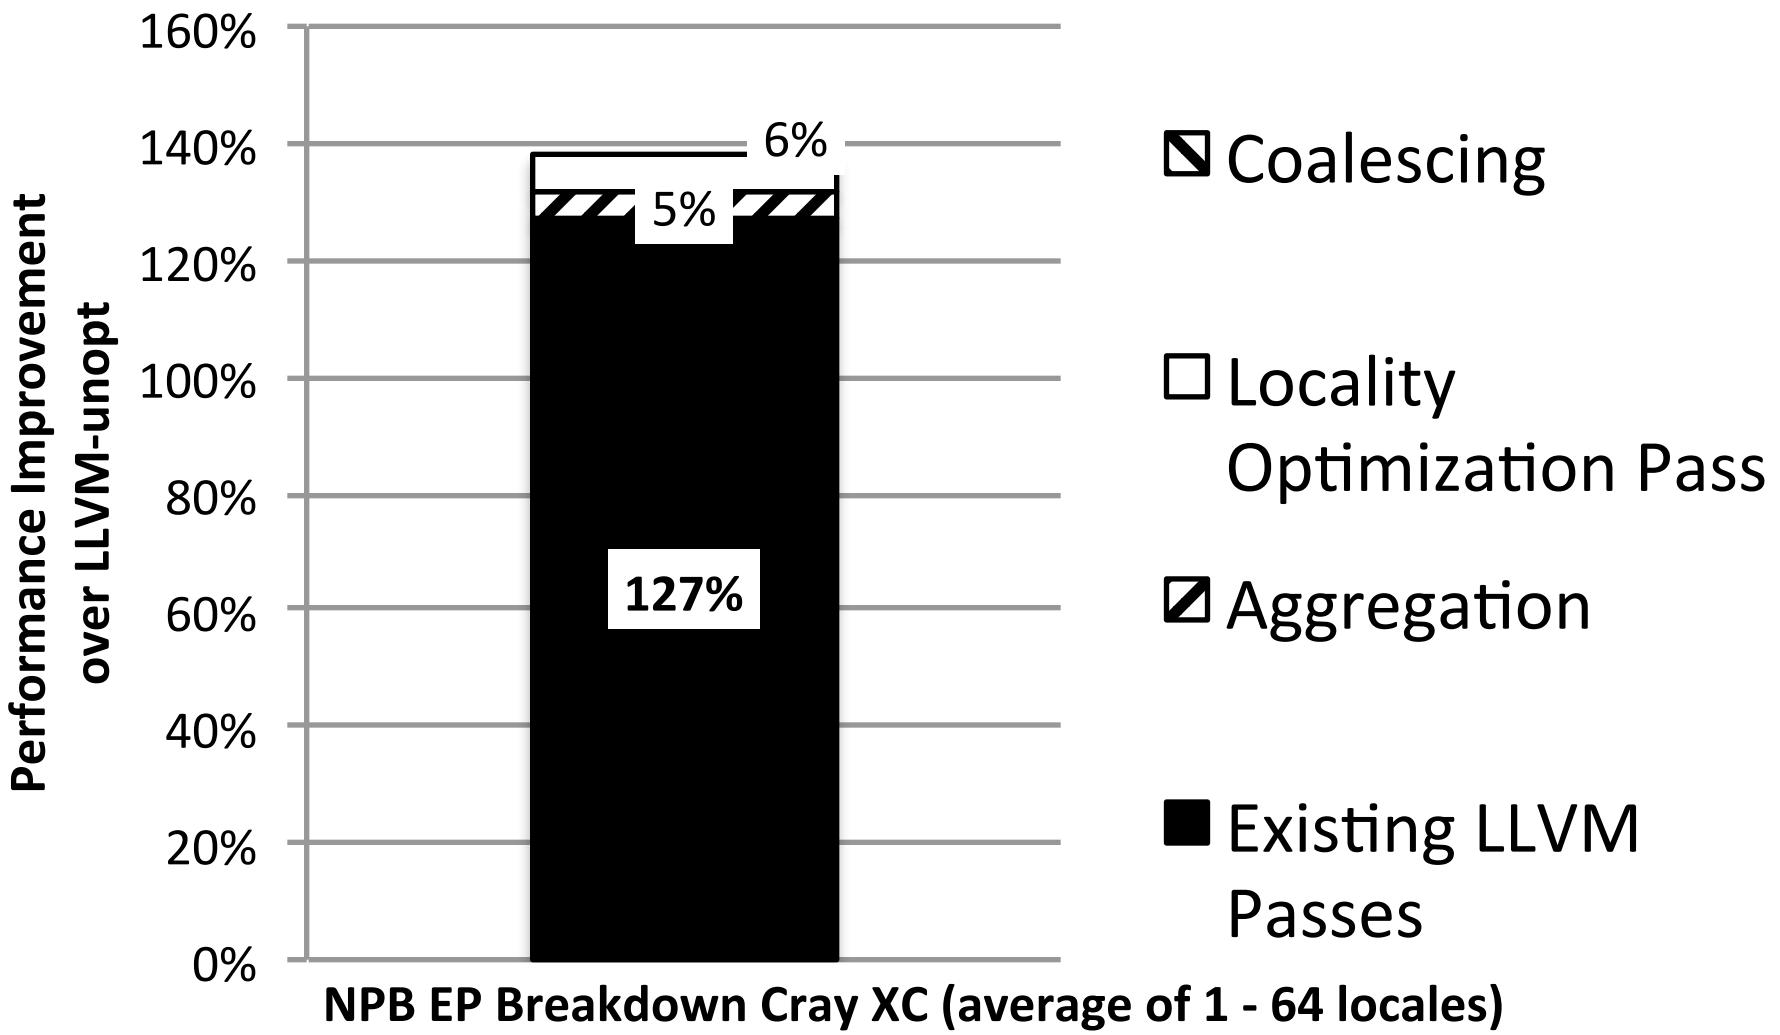






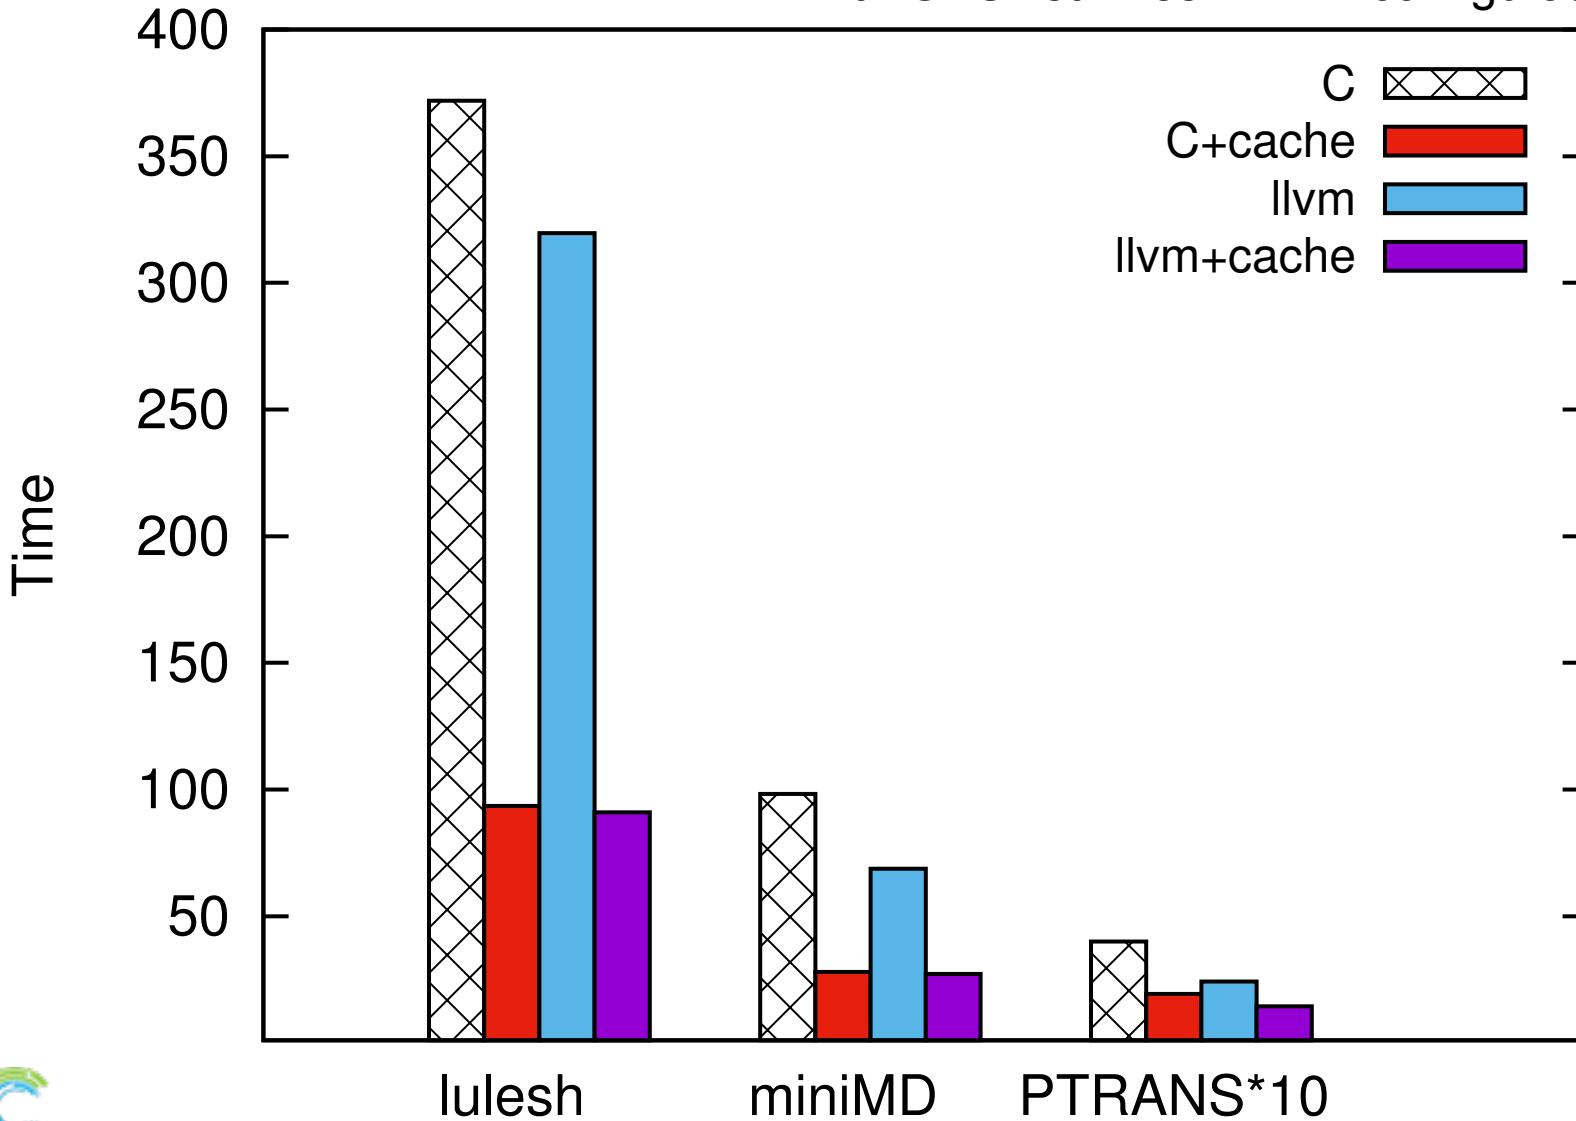






CACHING VS LLVM

* with GASNet Aries v1.11+ configuration



Questions?



10170

CASTLE ROCK

Legal Disclaimer

Information in this document is provided in connection with Cray Inc. products. No license, express or implied, to any intellectual property rights is granted by this document.

Cray Inc. may make changes to specifications and product descriptions at any time, without notice.

All products, dates and figures specified are preliminary based on current expectations, and are subject to change without notice.

Cray hardware and software products may contain design defects or errors known as errata, which may cause the product to deviate from published specifications. Current characterized errata are available on request.

Cray uses codenames internally to identify products that are in development and not yet publically announced for release. Customers and other third parties are not authorized by Cray Inc. to use codenames in advertising, promotion or marketing and any use of Cray Inc. internal codenames is at the sole risk of the user.

Performance tests and ratings are measured using specific systems and/or components and reflect the approximate performance of Cray Inc. products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance.

The following are trademarks of Cray Inc. and are registered in the United States and other countries: CRAY and design, SONEXION, URIKA, and YARCDATA. The following are trademarks of Cray Inc.: ACE, APPRENTICE2, CHAPEL, CLUSTER CONNECT, CRAYPAT, CRAYPORT, ECOPHLEX, LIBSCI, NODEKARE, THREADSTORM. The following system family marks, and associated model number marks, are trademarks of Cray Inc.: CS, CX, XC, XE, XK, XMT, and XT. The registered trademark LINUX is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a worldwide basis. Other trademarks used in this document are the property of their respective owners.

Copyright 2014 Cray Inc.





CRAY
THE SUPERCOMPUTER COMPANY

<http://chapel.cray.com>

chapel_info@cray.com

<https://github.com/chapel-lang/chapel/>