

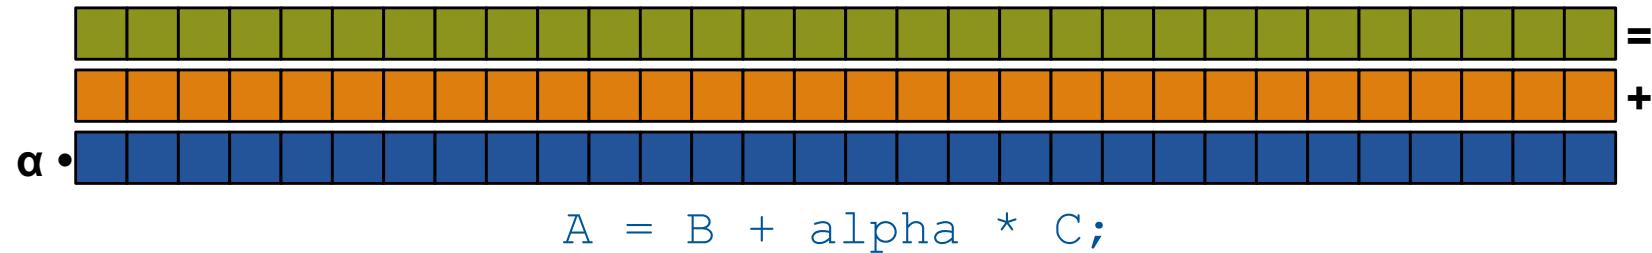
Data Parallelism with Locality: Domain Maps / Distributions (4x3 slides)



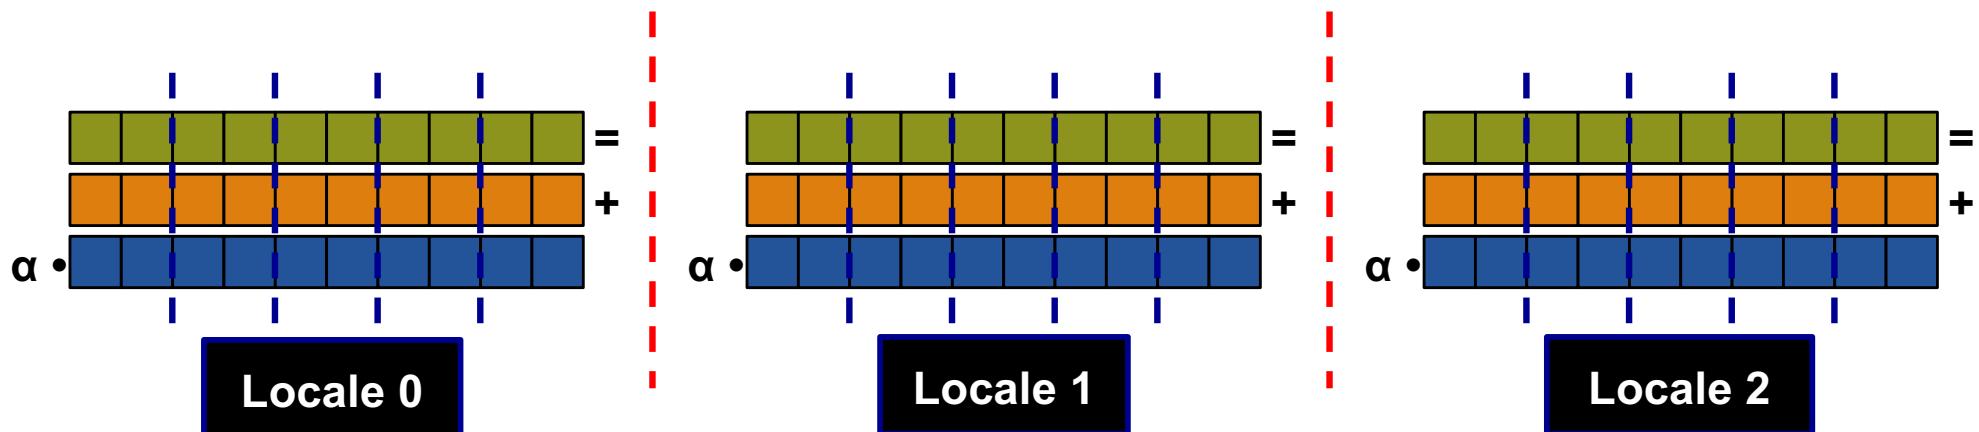
COMPUTE | STORE | ANALYZE

Domain Maps

Domain maps are “recipes” that instruct the compiler how to map the global view of a computation...



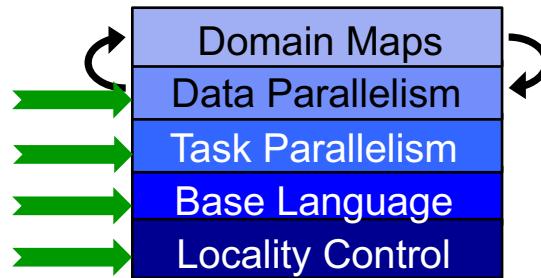
...to the target locales' memory and processors:



Chapel's Domain Map Philosophy

- 1. Chapel provides a library of standard domain maps**
 - to support common array implementations effortlessly

- 2. Expert users can write their own domain maps in Chapel**
 - to cope with any shortcomings in our standard library



- 3. Chapel's standard domain maps are written using the same end-user framework**
 - to avoid a performance cliff between “built-in” and user-defined cases



Domain Map Roles

They define data storage:

- Mapping of domain indices and array elements to locales
- Layout of arrays and index sets in each locale's memory

...as well as operations:

- random access, iteration, slicing, reindexing, rank change,
...
- the Chapel compiler generates calls to these methods to implement the user's array operations



COMPUTE | STORE | ANALYZE

Copyright 2018 Cray Inc.

Layouts and Distributions

Domain Maps fall into two major categories:

layouts:

- e.g., a desktop machine or multicore node
- **examples:** row- and column-major order, tilings, compressed sparse row, space-filling curves

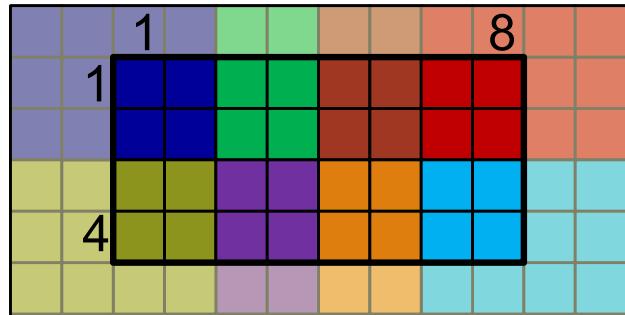
distributions:

- e.g., a distributed memory cluster or supercomputer
- **examples:** Block, Cyclic, Block-Cyclic, Recursive Bisection, ...



Sample Distributions: Block and Cyclic

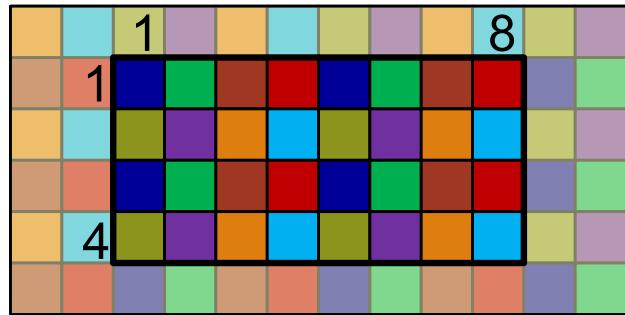
```
var Dom = {1..4, 1..8} dmapped Block( {1..4, 1..8} );
```



distributed to



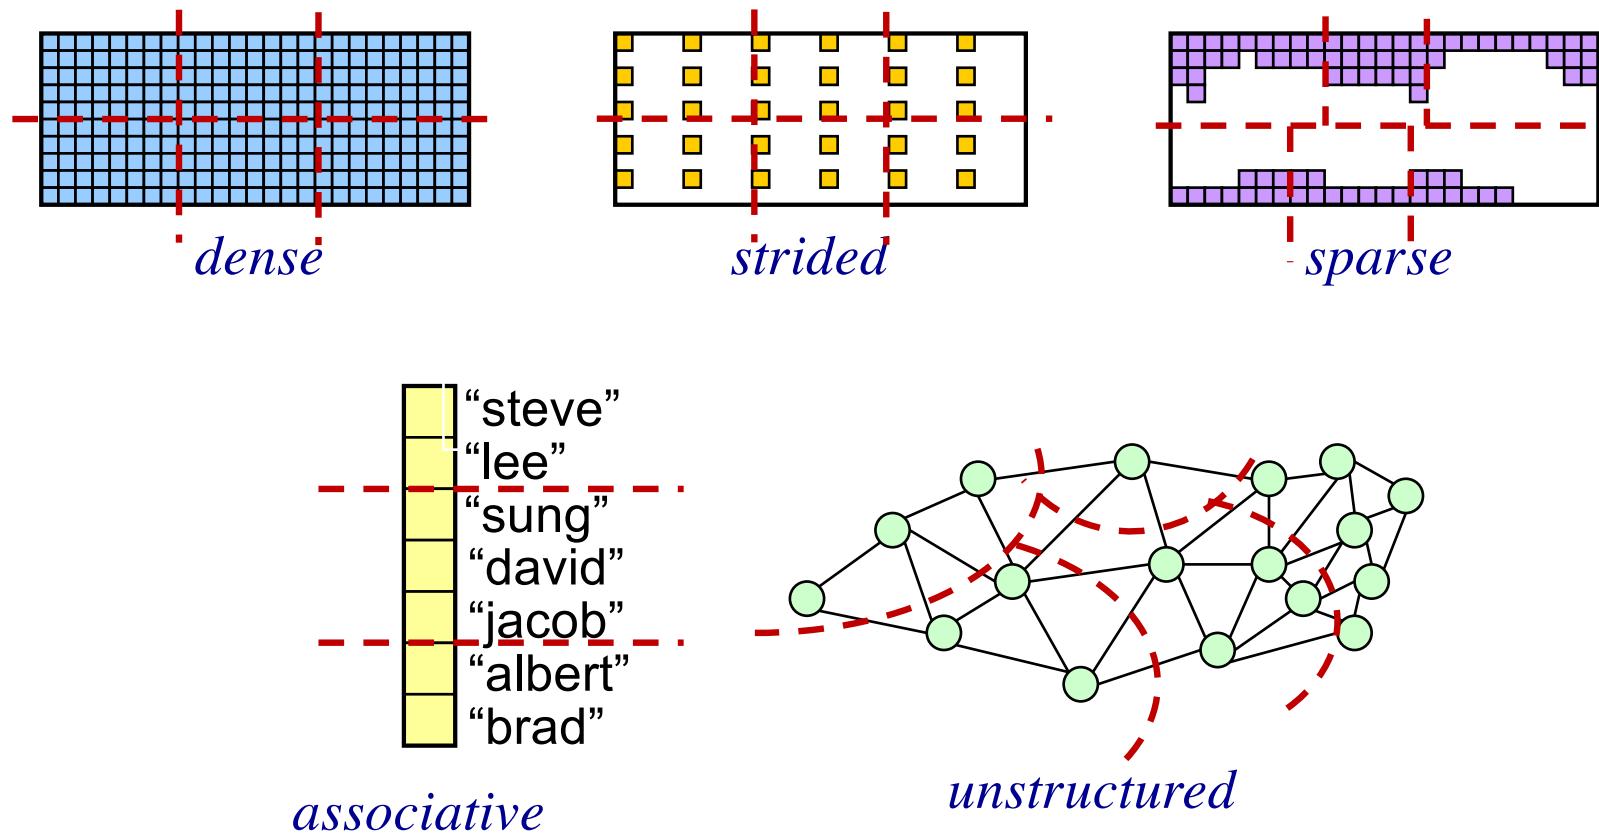
```
var Dom = {1..4, 1..8} dmapped Cyclic( startIdx=(1,1) );
```



distributed to



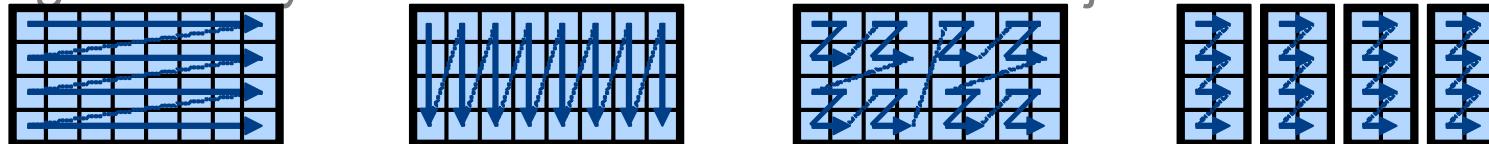
All Domain Types Support Domain Maps



Data Parallelism Implementation Qs

Q1: How are arrays laid out in memory?

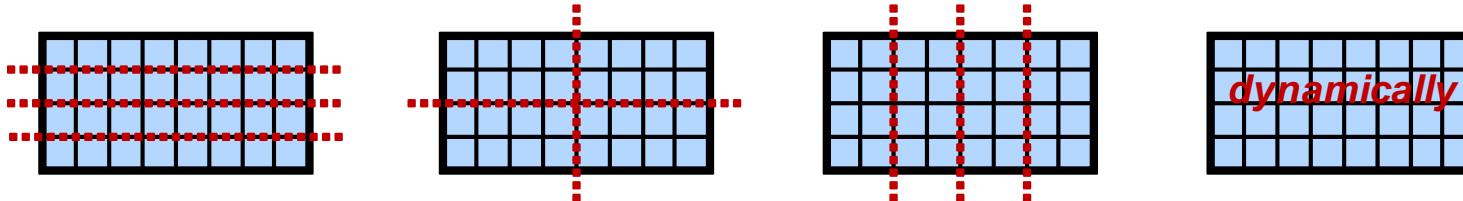
- Are regular arrays laid out in row- or column-major order? Or...?



- How are sparse arrays stored? (COO, CSR, CSC, block-structured, ...?)

Q2: How are arrays stored by the locales?

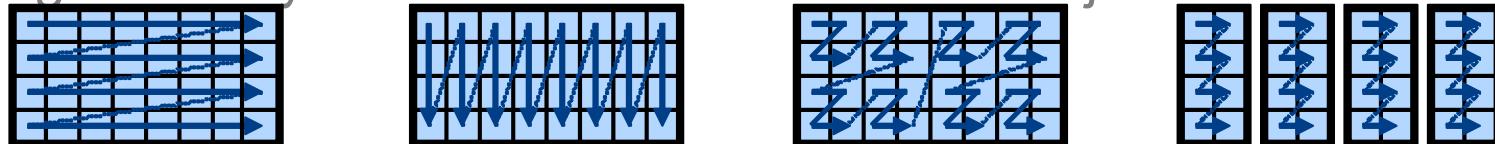
- Completely local to one locale? Or distributed?
- If distributed... In a blocked manner? cyclically? block-cyclically?
recursively bisected? dynamically rebalanced? ...?



Data Parallelism Implementation Qs

Q1: How are arrays laid out in memory?

- Are regular arrays laid out in row- or column-major order? Or...?



- How are sparse arrays stored? (COO, CSR, CSC, block-structured, ...?)

Q2: How are arrays stored by the locales?

- Completely local to one locale? Or distributed?
- If distributed... In a blocked manner? cyclically? block-cyclically?
recursively bisected? dynamically rebalanced? ...?

A: Chapel's *domain maps* are designed to give the user full control over such decisions



Jacobi Iteration in Chapel

```
config const n = 6,  
        epsilon = 1.0e-5;  
  
const BigD = {0..n+1, 0..n+1},  
            D = BigD[1..n, 1..n],  
            LastRow = D.exterior(1,0);  
  
var A, Temp : [BigD] real;
```

By default, domains and their arrays are mapped to a single locale.
Any data parallelism over such domains/ arrays will be executed by the cores on that locale.
Thus, this is a shared-memory parallel program.

```
Temp[i,j] = (A[i-1,j] + A[i+1,j] + A[i,j-1] + A[i,j+1]) / 4;  
  
const delta = max reduce abs(A[D] - Temp[D]);  
A[D] = Temp[D];  
} while (delta > epsilon);  
  
writeln(A);
```



COMPUTE | STORE | ANALYZE

Copyright 2018 Cray Inc.

Jacobi Iteration in Chapel

```

config const n = 6,
      epsilon = 1.0e-5;

const BigD = {0..n+1, 0..n+1} dmapped Block({1..n, 1..n}),
      D = BigD[1..n, 1..n],
      LastRow = D.exterior(1,0);

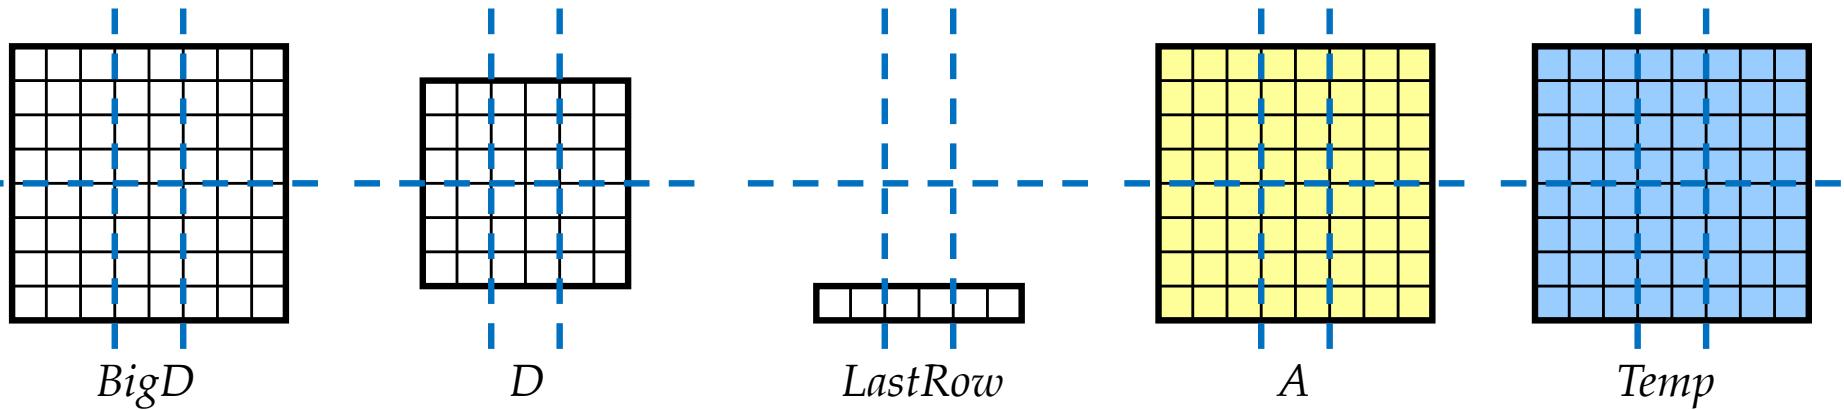
var A, Temp : [BigD] real;

```

With this simple change, we specify a mapping from the domains and arrays to locales
 Domain maps describe the mapping of domain indices and array elements to *locales*

specifies how array data is distributed across locales

specifies how iterations over domains/arrays are mapped to locales





Jacobi Iteration in Chapel

```
config const n = 6,
          epsilon = 1.0e-5;

const BigD = {0..n+1, 0..n+1} dmapped Block({1..n, 1..n}),
      D = BigD[1..n, 1..n],
      LastRow = D.exterior(1,0);

var A, Temp : [BigD] real;

A[LastRow] = 1.0;

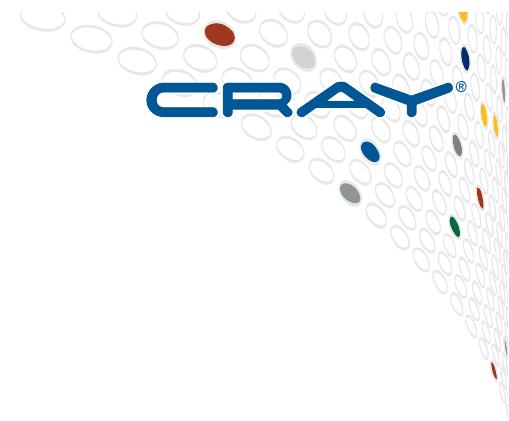
do {
    forall (i,j) in D do
        Temp[i,j] = (A[i-1,j] + A[i+1,j] + A[i,j-1] + A[i,j+1]) / 4;

    const delta = max reduce abs(A[D] - Temp[D]);
    A[D] = Temp[D];
} while (delta > epsilon);

writeln(A);

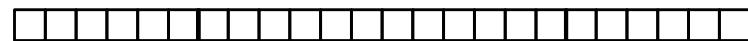
use BlockDist;
```



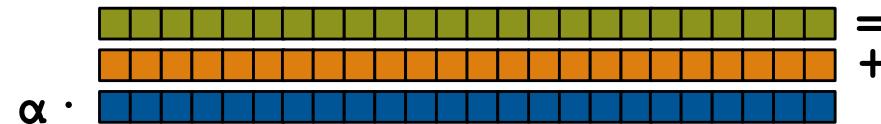


STREAM Triad in Chapel

```
const ProblemSpace = {1..m};
```



```
var A, B, C: [ProblemSpace] real;
```



```
A = B + alpha * C;
```



COMPUTE | STORE | ANALYZE

Copyright 2018 Cray Inc.

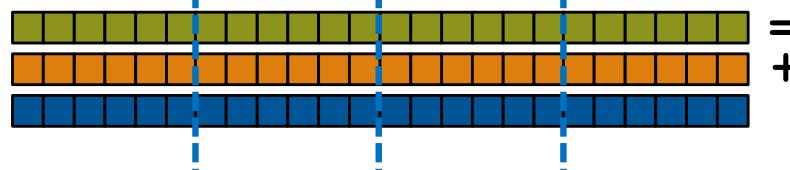


STREAM Triad in Chapel (multicore)

```
const ProblemSpace = {1..m};
```



```
var A, B, C: [ProblemSpace] real;
```



```
A = B + alpha * C;
```

No domain map specified \Rightarrow use default layout

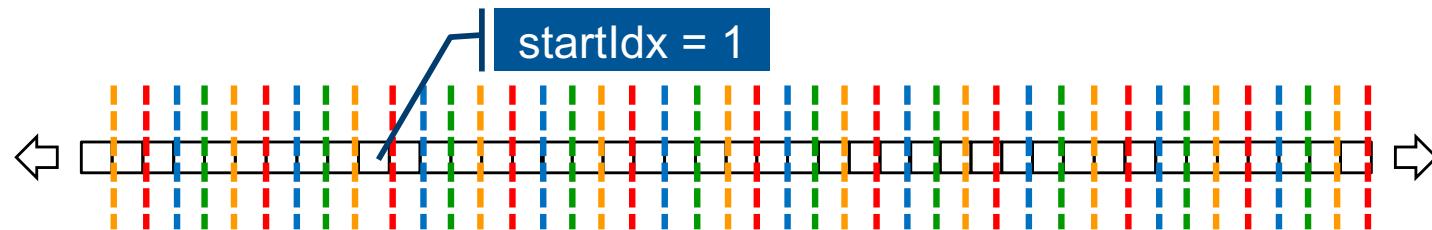
- current locale owns all domain indices and array values
- computation will execute using local processors only



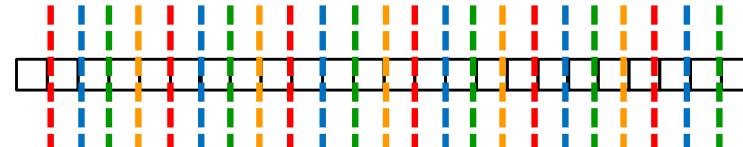
COMPUTE | STORE | ANALYZE

Copyright 2018 Cray Inc.

STREAM Triad in Chapel (multilocale, cyclic)



```
const ProblemSpace = {1..m}
    dmapped Cyclic(startIdx=1);
```

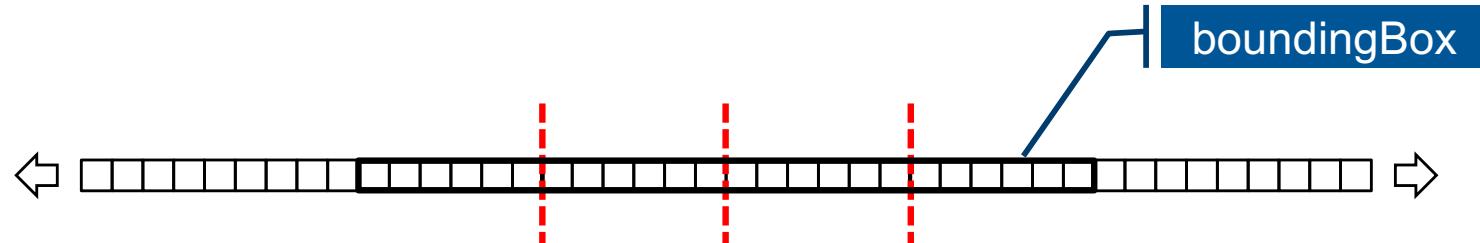


```
var A, B, C: [ProblemSpace] real;
```

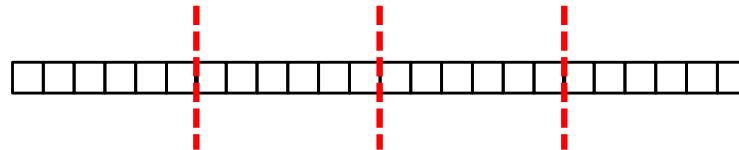


```
A = B + alpha * C;
```

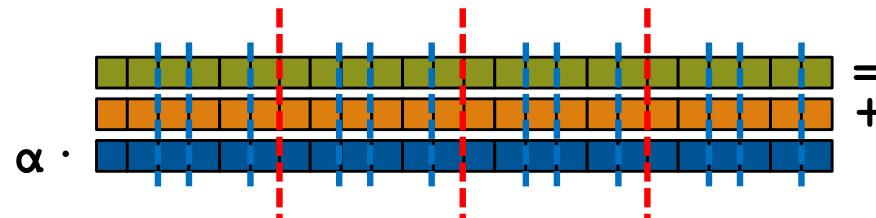
STREAM Triad in Chapel (multilocale, blocked)



```
const ProblemSpace = {1..m}  
                      dmapped Block(boundingBox={1..m});
```



```
var A, B, C: [ProblemSpace] real;
```



```
A = B + alpha * C;
```

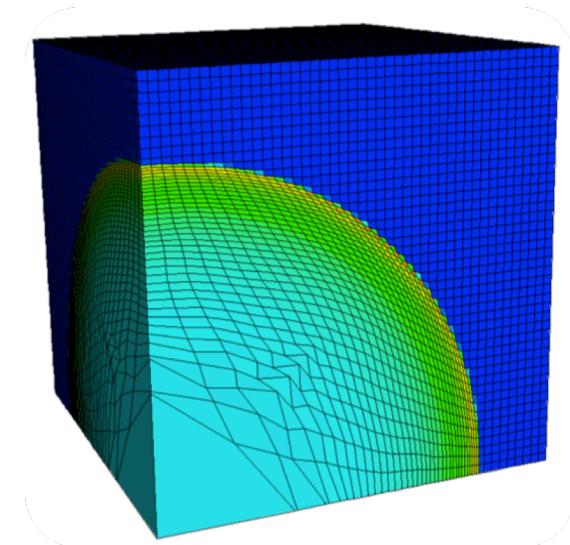
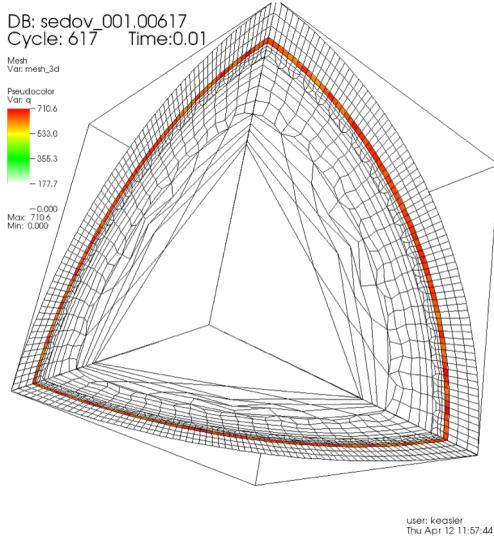
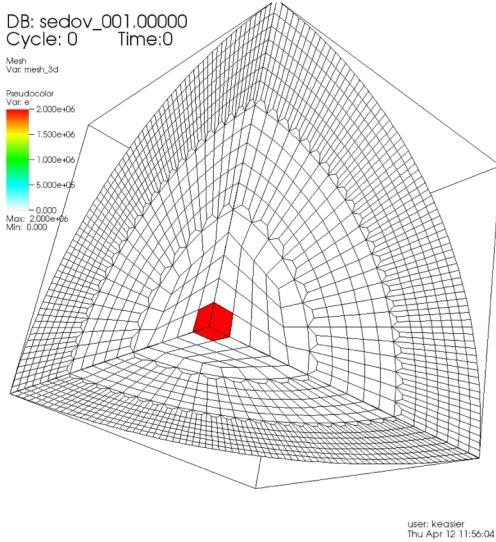


COMPUTE | STORE | ANALYZE

Copyright 2018 Cray Inc.

LULESH: a DOE Proxy Application

Goal: Solve one octant of the spherical Sedov problem (blast wave) using Lagrangian hydrodynamics for a single material



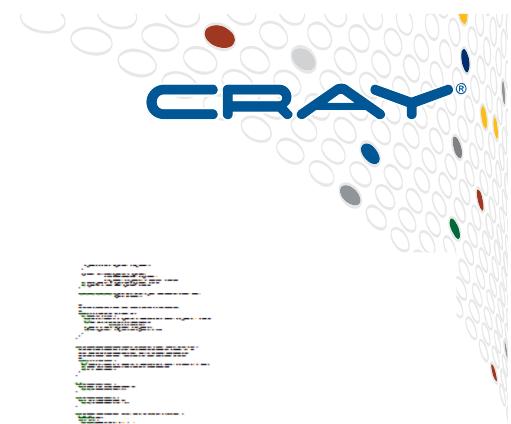
pictures courtesy of Rob Neely, Bert Still, Jeff Keasler, LLNL



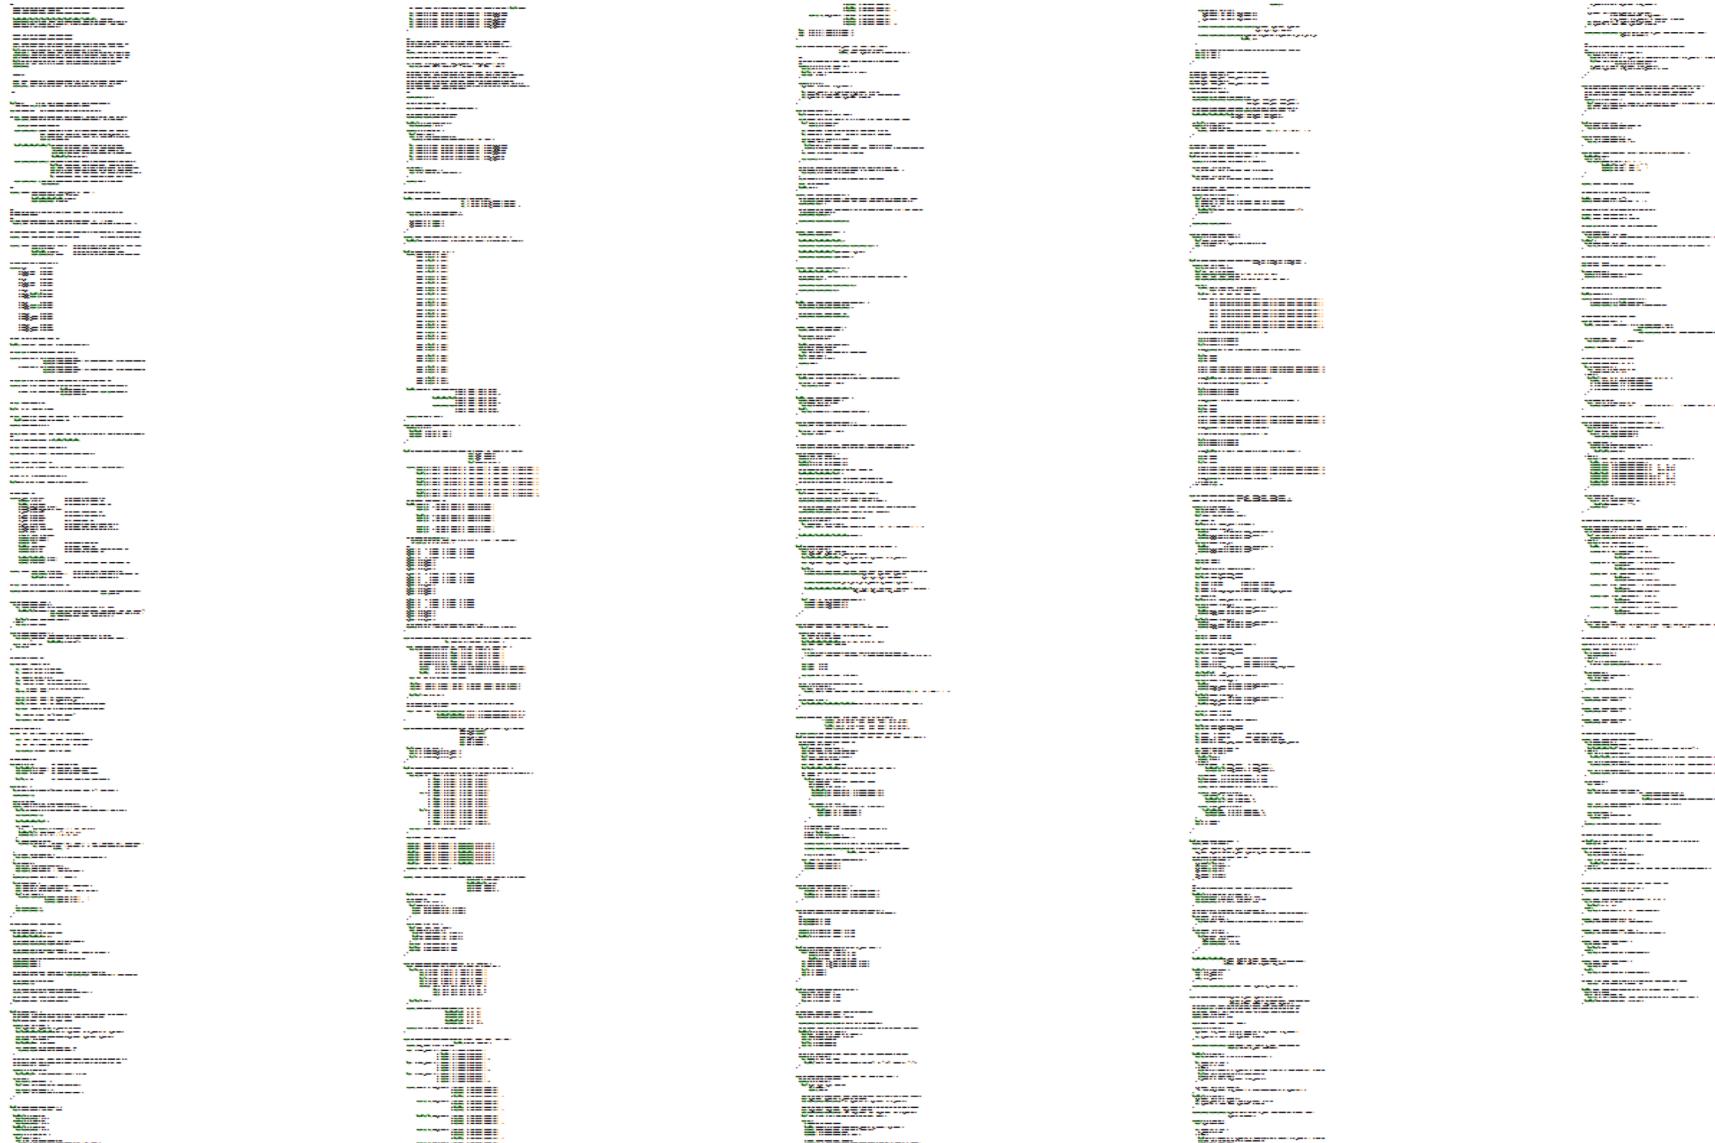
COMPUTE

| STORE

| ANALYZE



LULESH in Chapel



COMPUTE | STORE | ANALYZE

Copyright 2018 Cray Inc.

LULESH in Chapel

1288 lines of source code

plus 266 lines of comments
487 blank lines

(the corresponding C+MPI+OpenMP version is nearly 4x bigger)

This can be found in the Chapel release under examples/benchmarks/lulesh/*.chpl

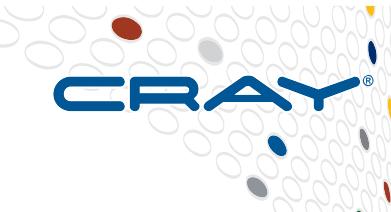


COMPUTE

STORE

ANALYZE

LULESH in Chapel



This is all of the representation dependent code.
It specifies:

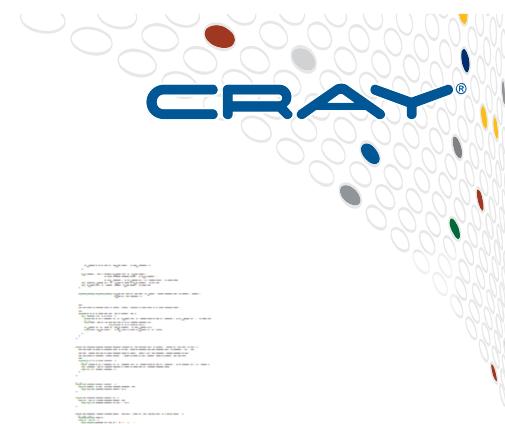
- data structure choices
 - structured vs. unstructured mesh
 - local vs. distributed data
 - sparse vs. dense materials arrays
- a few supporting iterators



COMPUTE

STORE

ANALYZE



LULESH in Chapel



Here is some sample representation-independent code
`IntegrateStressForElems()`
[LULESH spec](#), section 1.5.1.1 (2.)



COMPUTE | STORE | ANALYZE

Representation-Independent Physics

```

proc IntegrateStressForElems(sigxx, sigyy, sigzz, determ) {
    forall k in Elems { ← parallel loop over elements
        var b_x, b_y, b_z: 8*real;
        var x_local, y_local, z_local: 8*real;
        localizeNeighborNodes(k, x, x_local, y, y_local, z, z_local); ← collect nodes neighboring this
        var fx_local, fy_local, fz_local: 8*real; element; localize node fields

        local {
            /* Volume calculation involves extra work for numerical consistency. */
            CalcElemShapeFunctionDerivatives(x_local, y_local, z_local,
                b_x, b_y, b_z, determ[k]);
            CalcElemNodeNormals(b_x, b_y, b_z, x_local, y_local, z_local);
            SumElemStressesToNodeForces(b_x, b_y, b_z, sigxx[k], sigyy[k], sigzz[k],
                fx_local, fy_local, fz_local);
        }
        for (noi, t) in elemToNodesTuple(k) { ← update node forces from
            fx[noi].add(fx_local[t]); element stresses
            fy[noi].add(fy_local[t]);
            fz[noi].add(fz_local[t]);
        }
    }
}

```

Because of domain maps, this code is independent of:

- structured vs. unstructured mesh
- shared vs. distributed data
- sparse vs. dense representation





For More Information on Domain Maps

HotPAR'10: *User-Defined Distributions and Layouts in Chapel: Philosophy and Framework*
Chamberlain, Deitz, Iten, Choi; June 2010

CUG 2011: *Authoring User-Defined Domain Maps in Chapel*
Chamberlain, Choi, Deitz, Iten, Litvinov; May 2011

Chapel release:

- Documentation of current domain maps:
<http://chapel.cray.com/docs/latest/modules/layoutdist.html>
- Technical notes detailing the domain map interface for implementers:
<http://chapel.cray.com/docs/latest/technotes/dsi.html>



COMPUTE | STORE | ANALYZE

Copyright 2018 Cray Inc.

Two Other Thematically Similar Features

- 1) **parallel iterators:** Permit users to specify the parallelism and work decomposition used by forall loops
 - including zippered forall loops

- 2) **locale models:** Permit users to model the target architecture and how Chapel should be implemented on it
 - e.g., how to manage memory, create tasks, communicate, ...

Like domain maps, these are...

- ...written in Chapel by expert users using lower-level features
 - e.g., task parallelism, on-clauses, base language features, ...
- ...available to the end-user via higher-level abstractions
 - e.g., forall loops, on-clauses, lexically scoped PGAS memory, ...





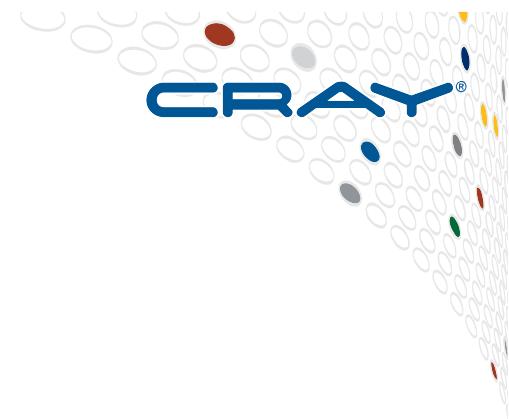
Summary of this Section

- Chapel avoids locking crucial implementation decisions into the language specification
 - local and distributed parallel array implementations
 - parallel loop scheduling policies
 - target architecture models
- Instead, these can be...
 - ...specified in the language by an advanced user
 - ...swapped between with minimal code changes
- The result cleanly separates the roles of domain scientist, parallel programmer, and compiler/runtime



COMPUTE | STORE | ANALYZE

Copyright 2018 Cray Inc.



Any Questions about Domain Maps?



COMPUTE

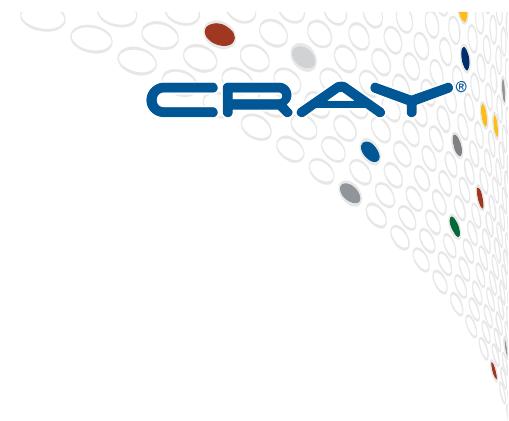
|

STORE

|

ANALYZE

Copyright 2018 Cray Inc.



Overarching Example:

Smith-Waterman Algorithm for Sequence Alignment



COMPUTE | STORE | ANALYZE

Smith-Waterman

Goal: Determine the similarities/differences between two protein sequences/nucleotides.

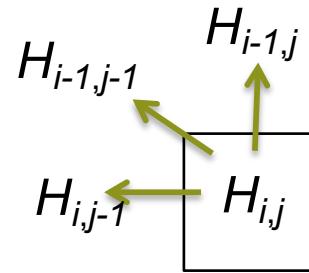
- e.g., ACACACTA and AGCACACA*

Basis of Computation: Defined via a recursive formula:

$$H(i,0) = 0$$

$$H(0,j) = 0$$

$$H(i,j) = f(H(i-1, j-1), H(i-1, j), H(i, j-1))$$



Caveat: This is a classic, rather than cutting-edge sequence alignment algorithm, but it illustrates an important parallel paradigm: *wavefront computation*

*Source of running example: Wikipedia



Smith-Waterman

Naïve Task-Parallel Approach:

```

proc computeH(i, j) {
    if (i == 0 || j == 0) then
        return 0;
    else
        var h_NW, h_N, h_W: int;
        cobegin {
            h_NW = computeH(i-1, j-1);
            h_N   = computeH(i-1, j);
            h_W   = computeH(i,     j-1);
        }
        return f(h_NW, h_N, h_W);
}

```

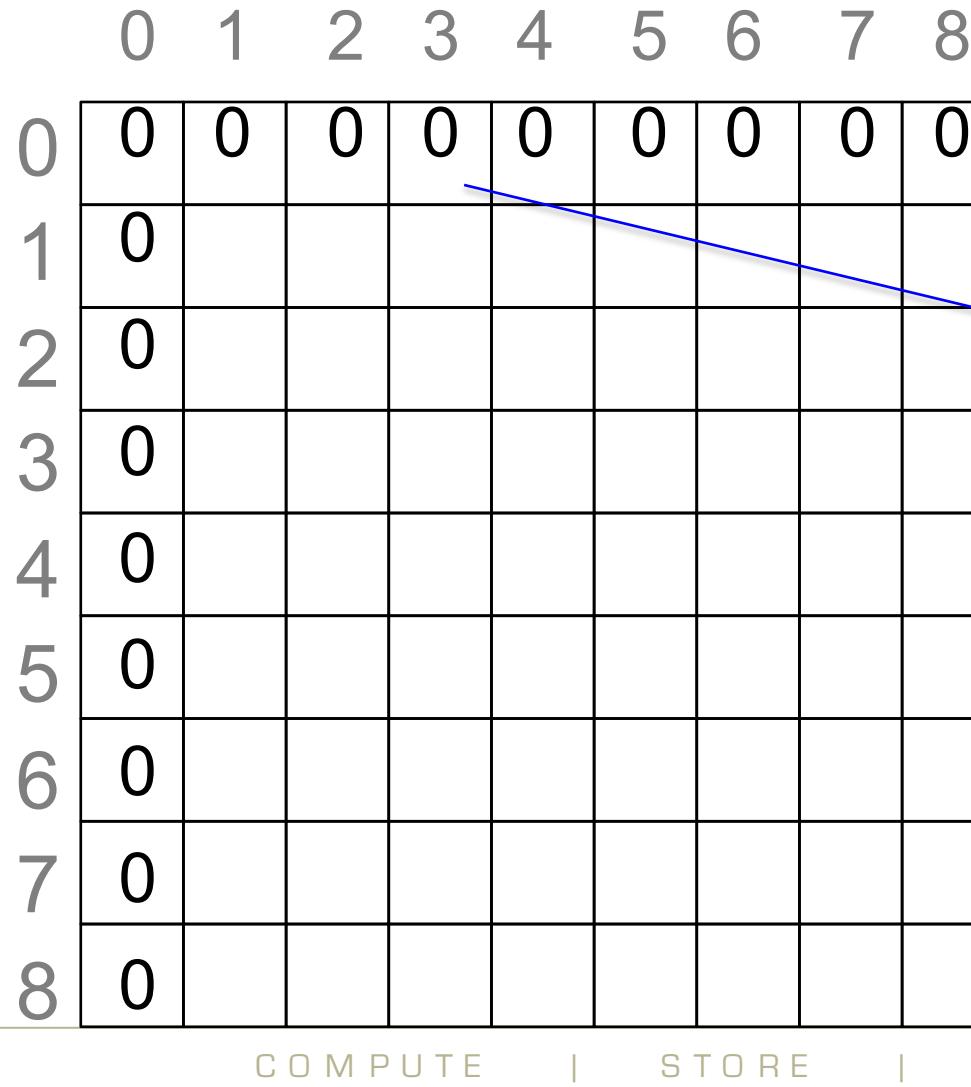
Note: Recomputes most subexpressions redundantly

This is a case for dynamic programming!



Smith-Waterman

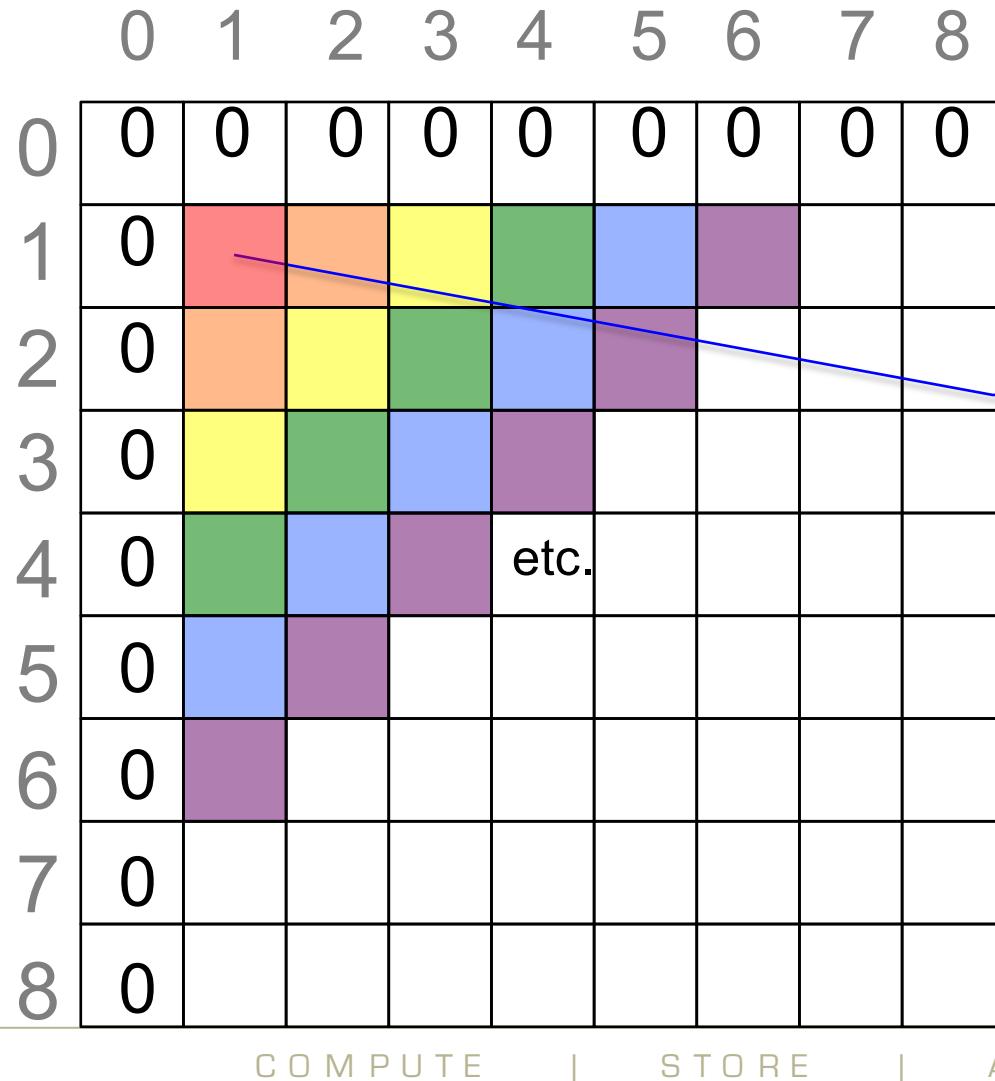
Dynamic Programming Approach:



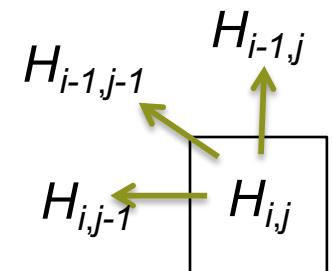
Step 1: Initialize
boundaries to 0

Smith-Waterman

Dynamic Programming Approach:



Step 2: Compute cells when we're able to



Smith-Waterman

Dynamic Programming Approach:

	0	1	2	3	4	5	6	7	8
0	0	0	0	0	0	0	0	0	0
1	0	2	1	2	1	2	1	0	2
2	0	1	1	1	1	1	1	0	1
3	0	0	3	2	3	2	3	2	1
4	0	2	2	5	4	5	4	3	4
5	0	1	4	4	7	6	7	6	5
6	0	2	3	6	6	9	8	7	8
7	0	1	4	5	8	8	11	10	9
8	0	2	3	6	7	10	10	10	12

Step 3: Follow trail of breadcrumbs back



COMPUTE

STORE

ANALYZE

Smith-Waterman

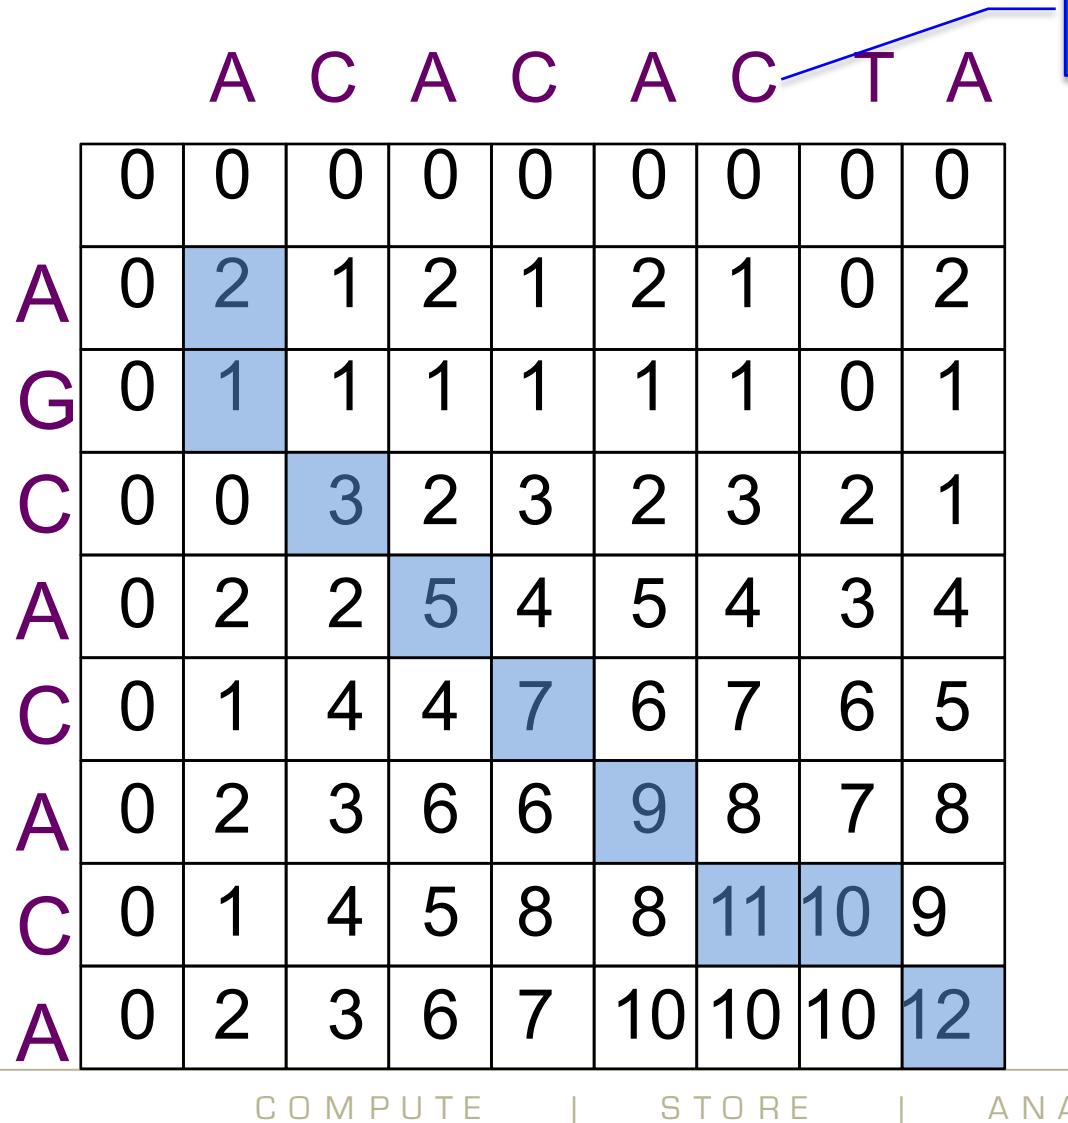
Dynamic Programming Approach:

	0	1	2	3	4	5	6	7	8
0	0	0	0	0	0	0	0	0	0
1	0	2	1	2	1	2	1	0	2
2	0	1	1	1	1	1	1	0	1
3	0	0	3	2	3	2	3	2	1
4	0	2	2	5	4	5	4	3	4
5	0	1	4	4	7	6	7	6	5
6	0	2	3	6	6	9	8	7	8
7	0	1	4	5	8	8	11	10	9
8	0	2	3	6	7	10	10	10	12

Step 3: Follow trail of breadcrumbs back

Smith-Waterman

Dynamic Programming Approach:

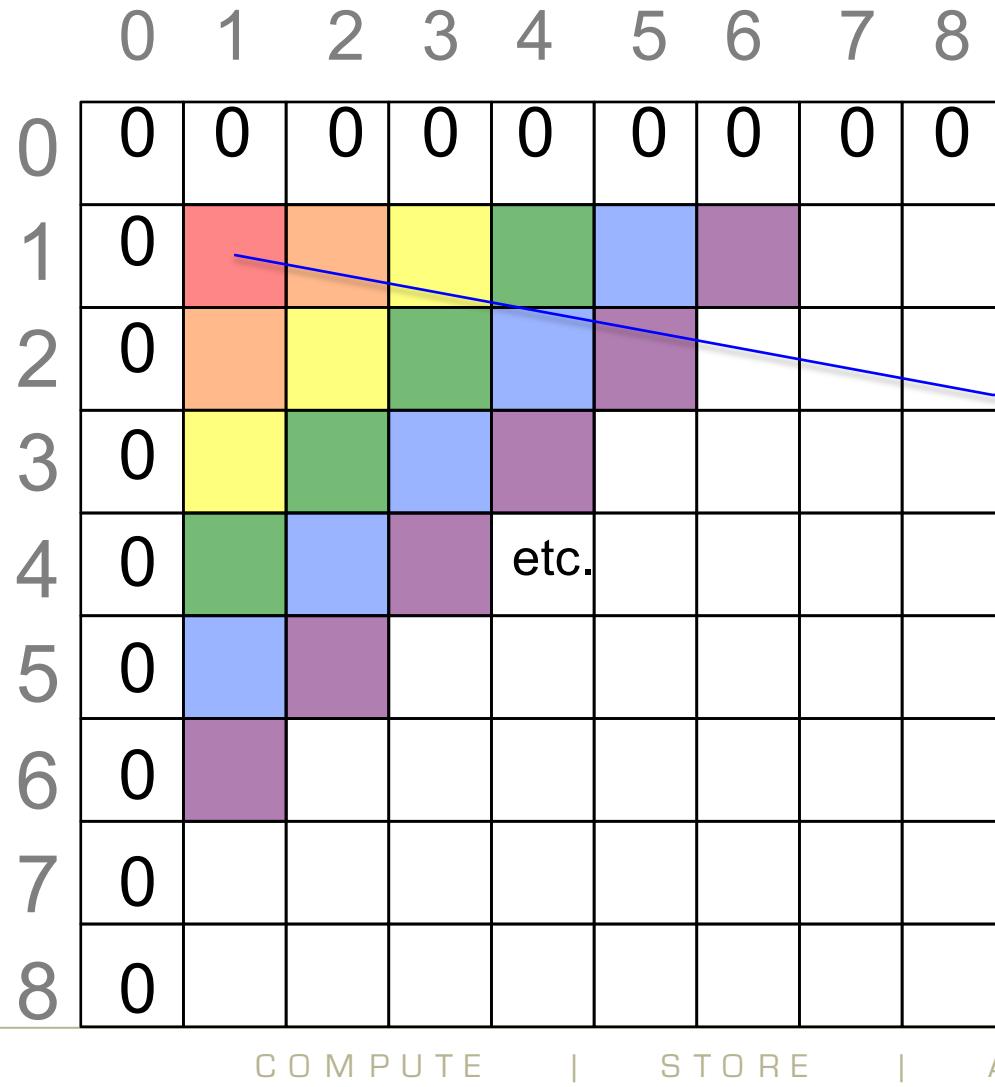


Step 4: Interpret the path against the original sequences

AGCACAC-A
A-CACACTA

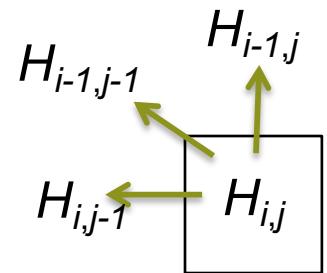
Smith-Waterman

Dynamic Programming Approach:



Step 2: Compute cells
when we're able to

How could we do
this in parallel?



Smith-Waterman

Data-Parallel Approach:

```

proc computeH(H: [0..n, 0..n] int) {
    for upperDiag in 1..n do
        forall diagPos in 0..#upperDiag {
            const (i,j) = (diagPos+1, upperDiag-diagPos);
            H[i,j] = f(H[i-1,j-1], H[i-1,j], H[i,j-1]);
        }
    for lowerDiag in 1..n-1 do
        forall diagPos in lowerDiag..n-1 by -1 {
            const (i,j) = (diagPos+1, lowerDiag+diagPos);
            H[i,j] = f(H[i-1,j-1], H[i-1,j], H[i,j-1]);
        }
}

```

Loop over upper diagonals serially

Process each diagonal in parallel

Repeat for lower diagonals

Advantages:

- Reasonably clean
(if I got my indexing correct)

Disadvantages:

- Not so great in terms of cache use
- A bit fine-grained
 - small number of iterations per task



Naïve Data-Driven Task-Parallel Approach:

```

proc computeH(H: [0..n, 0..n] int) {
    const ProbSpace = H.domain.translate(1,1);
    var NeighborsDone: [ProbSpace] atomic int;
    var Ready$: [ProbSpace] sync int;

    NeighborsDone[1, ...].add(1);
    NeighborsDone[..., 1].add(1);
    NeighborsDone[1, 1].add(1);
    Ready$[1,1] = 1;

    coforall (i,j) in ProbSpace {
        const goNow = Ready$[i,j];
        H[i,j] = f(H[i-1,j-1], H[i-1,j], H[i,j-1]);
        const eastReady = NeighborsDone[i, j+1].fetchAdd(1);
        const seReady = NeighborsDone[i+1,j+1].fetchAdd(1);
        const southReady = NeighborsDone[i+1,j].fetchAdd(1);
        if (eastReady == 2) then Ready$[i, j+1] = 1;
        if (seReady == 2) then Ready$[i+1,j+1] = 1;
        if (southReady == 2) then Ready$[i+1,j] = 1;
    }
}

```

Create a domain describing shifted version of H's domain

Arrays to count how many of our 3 neighbors are done; and to signal when we can compute

Set up boundaries: north and west elements have a neighbor done; top-left is ready

Create a task per matrix element and have it block until ready

Compute our element

Increment our neighbors' counts

Signal our neighbors as ready if we're the third



Smith-Waterman

Naïve Data-Driven Task-Parallel Approach:

```

proc computeH(H: [0..n, 0..n] int) {
    const ProbSpace = H.domain.translate(1,1);
    var NeighborsDone: [ProbSpace] atomic int;
    var Ready$: [ProbSpace] sync int;

    NeighborsDone[1, ...].add(1);
    NeighborsDone[..., 1].add(1);
    NeighborsDone[1, 1].add(1);
    Ready$[1,1] = 1;

    cforall (i,j) in ProbSpace {
        const goNow = Ready$[i,j];
        H[i,j] = f(H[i-1,j-1], H[i-1,j], H[i,j-1]);
        const eastReady = NeighborsDone[i, j+1].fetchAdd(1);
        const seReady = NeighborsDone[i+1,j+1].fetchAdd(1);
        const southReady = NeighborsDone[i+1,j ].fetchAdd(1);
        if (eastReady == 2) then Ready$[i, j+1] = 1;
        if (seReady == 2) then Ready$[i+1,j+1] = 1;
        if (southReady == 2) then Ready$[i+1,j ] = 1;
    }
}

```

Disadvantages:

- Still not great in cache use
- Uses n^2 tasks
- Most spend most of their time blocking



Smith-Waterman

Slightly Less Naïve Data-Driven Task-Parallel Approach:

```

proc computeH(H: [0..n, 0..n] int) {
    const ProbSpace = H.domain.translate(1,1);
    var NeighborsDone: [ProbSpace] atomic int;
    NeighborsDone[1, ...].add(1);
    NeighborsDone[..., 1].add(1);
    NeighborsDone[1, 1].add(1);
    sync { computeHHelp(1,1); }
}

```

sync to ensure they're all done before we go on

```

proc computeHHelp(i,j) {
    H[i,j] = f(H[i-1,j-1], H[i-1,j], H[i,j-1]);
    const eastReady = NeighborsDone[i, j+1].fetchAdd(1);
    const seReady = NeighborsDone[i+1,j+1].fetchAdd(1);
    const southReady = NeighborsDone[i+1,j].fetchAdd(1);
    if (eastReady == 2) then begin computeHHelp(i, j+1);
    if (seReady == 2) then begin computeHHelp(i+1,j+1);
    if (southReady == 2) then begin computeHHelp(i+1,j);
}

```

Rather than create the tasks *a priori*, fire them off once we know they're ready to compute



COMPUTE | STORE | ANALYZE

Smith-Waterman

Slightly Less Naïve Data-Driven Task-Parallel Approach:

```
proc computeH(H: [0..n, 0..n] int) {
    const ProbSpace = H.domain.translate(1,1);
    var NeighborsDone: [ProbSpace] atomic int;
```

```
NeighborsDone[1, ...].add(1);
NeighborsDone[..., 1].add(1);
NeighborsDone[1, 1].add(1);
sync { computeHHelp(1,1); }
```

Disadvantages:

- Still uses a lot of tasks
- Each task is very fine-grained

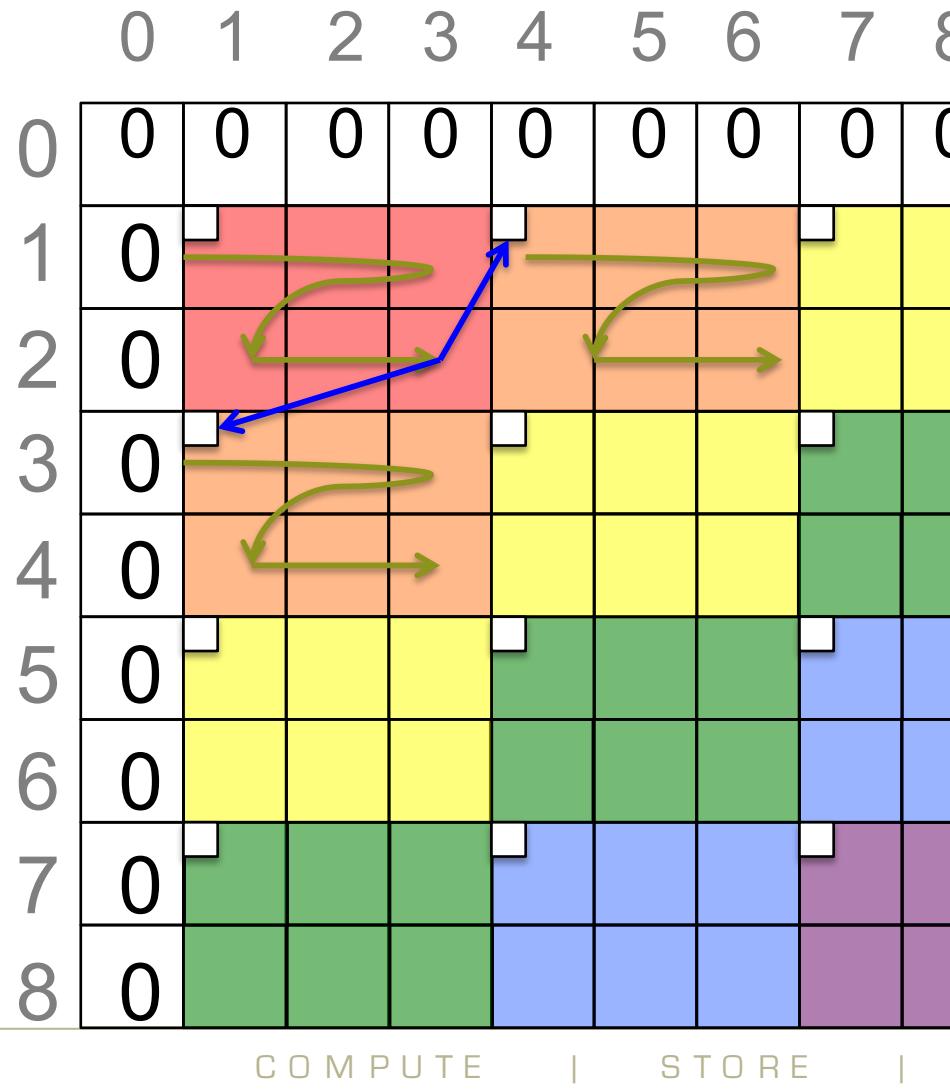
```
proc computeHHelp(i,j) {
    H[i,j] = f(H[i-1,j-1], H[i-1,j], H[i,j-1]);
    const eastReady = NeighborsDone[i, j+1].fetchAdd(1);
    const seReady = NeighborsDone[i+1,j+1].fetchAdd(1);
    const southReady = NeighborsDone[i+1,j].fetchAdd(1);
    if (eastReady == 2) then begin computeHHelp(i, j+1);
    if (seReady == 2) then begin computeHHelp(i+1,j+1);
    if (southReady == 2) then begin computeHHelp(i+1,j);
}
```

COMPUTE | STORE | ANALYZE



Smith-Waterman

Coarsening the Parallelism into Chunks:



Chunked Data-Driven Task-Parallel Approach:

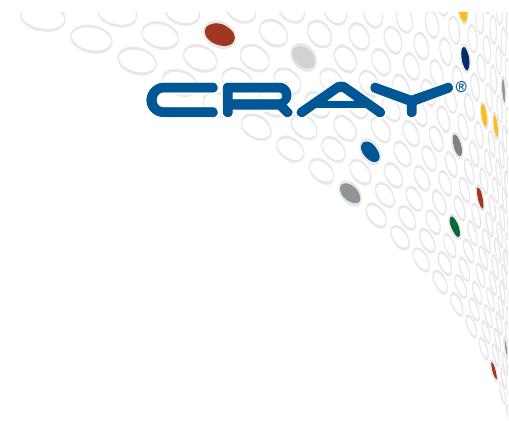
```

proc computeH(H: [0..n, 0..n] int) {
    const ProbSpace = H.domain.translate(1,1);
    const StrProbSpace = ProbSpace by (rowsPerChunk, colsPerChunk);
    var NeighborsDone: [StrProbSpace] atomic int; → Use strided array for atomics
        NeighborsDone[1, ...].add(1);
        NeighborsDone[..., 1].add(1);
        NeighborsDone[1, 1].add(1);
    sync { computeHHelp(1,1); } → Change helper to iterate over a chunk serially

proc computeHHelp(x,y) {
    for (i,j) in ProbSpace[x..#rowsPerChunk, y..#colsPerChunk] do
        H[i,j] = f(H[i-1,j-1], H[i-1,j], H[i,j-1]);
    const eastReady = NeighborsDone[x, y+colsPerChunk].fetchAdd(1);
    const seReady = NeighborsDone[x+rowsPerChunk, y+colsPerChunk].fetchAdd(1);
    const southReady = NeighborsDone[x+rowsPerChunk, y].fetchAdd(1);
    if (eastReady == 2) then begin computeHHelp(x, y+colsPerChunk);
    if (seReady == 2) then begin computeHHelp(x+rowsPerChunk, y+colsPerChunk);
    if (southReady == 2) then begin computeHHelp(x+rowsPerChunk, y);
} → Stride indices to get to next chunk's origin
    }
```



COMPUTE | STORE | ANALYZE



Distributed Smith-Waterman

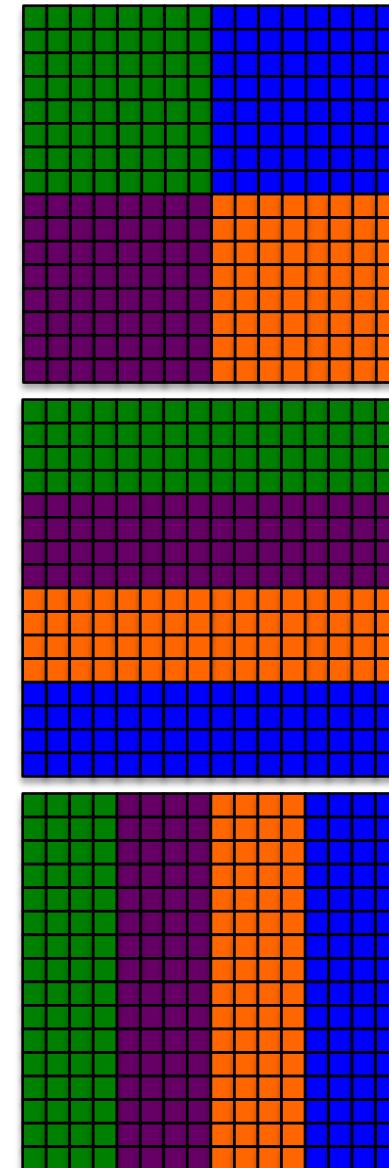
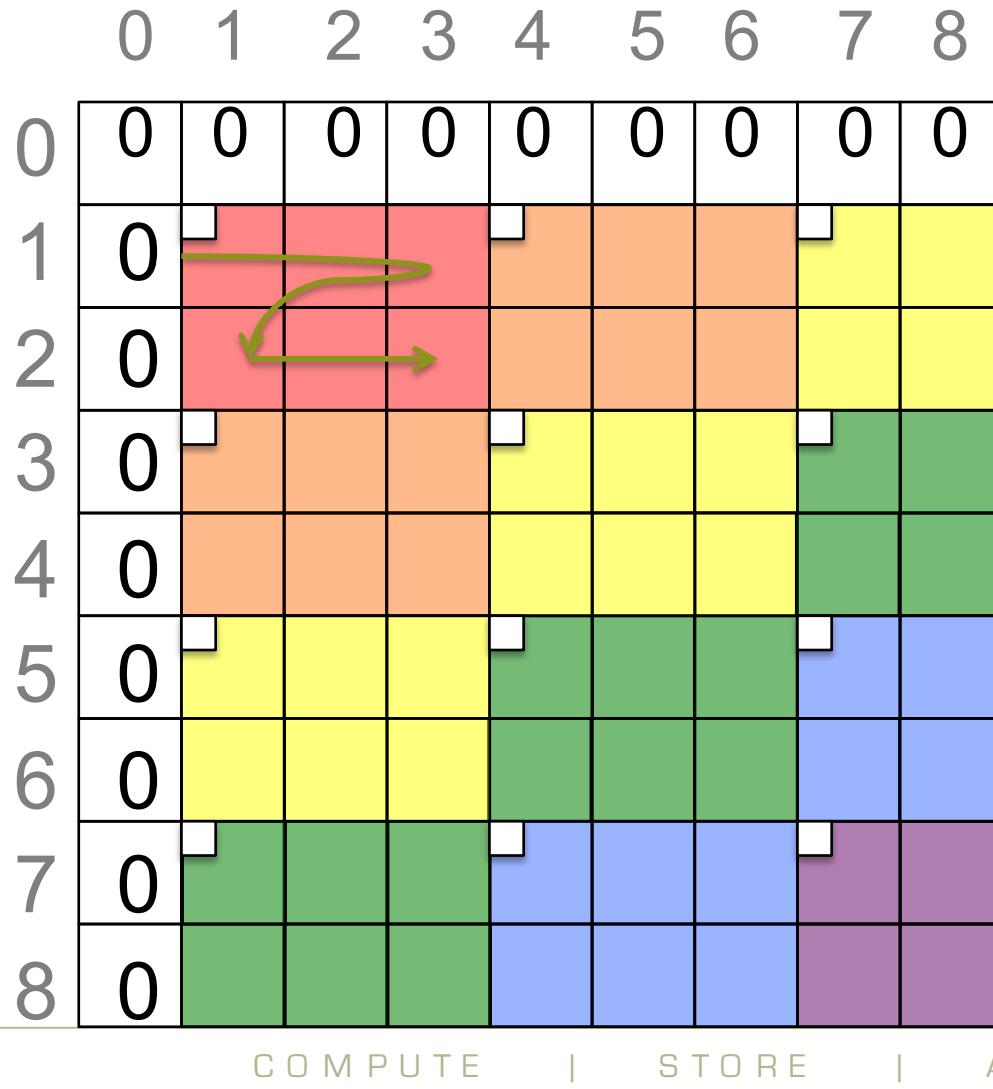


COMPUTE | STORE | ANALYZE

Copyright 2018 Cray Inc.

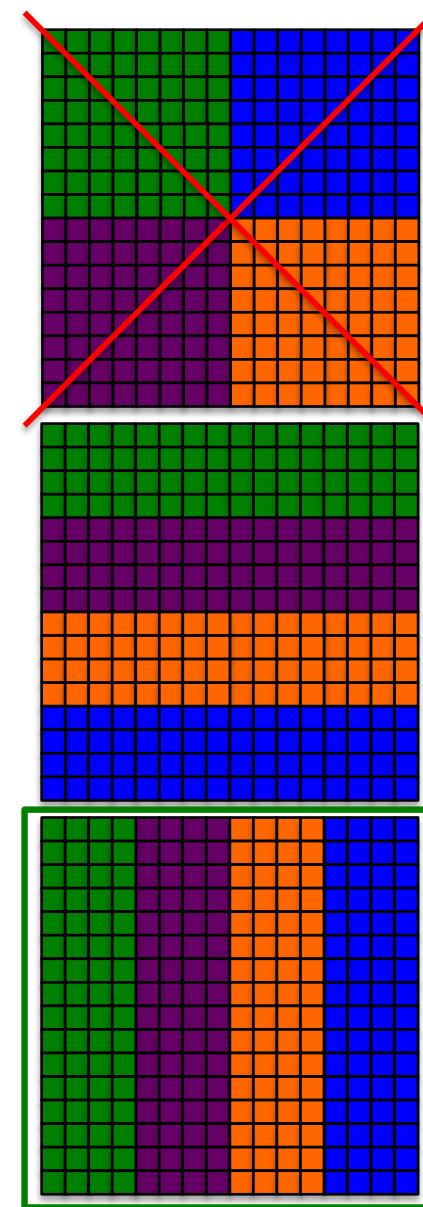
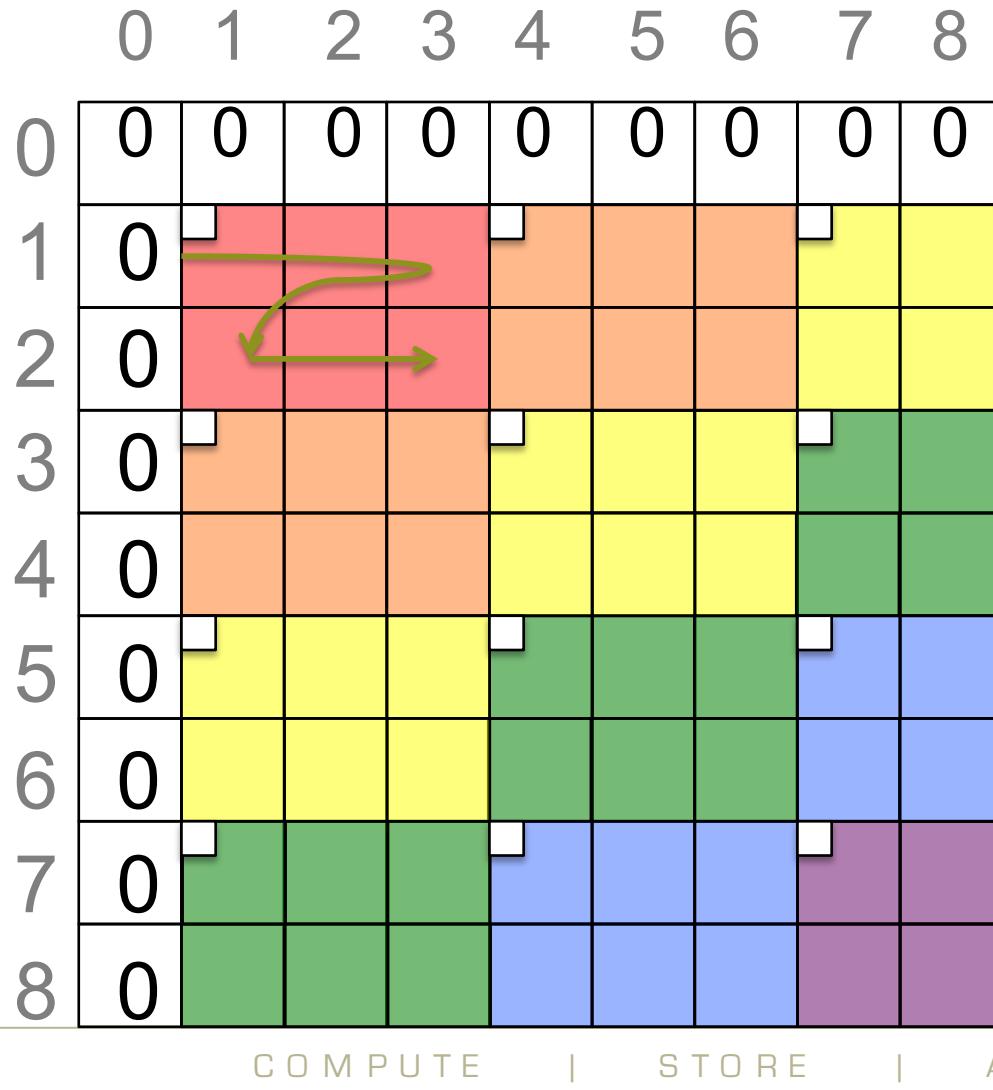
Distributed Smith-Waterman

Now, what about distributed memory?



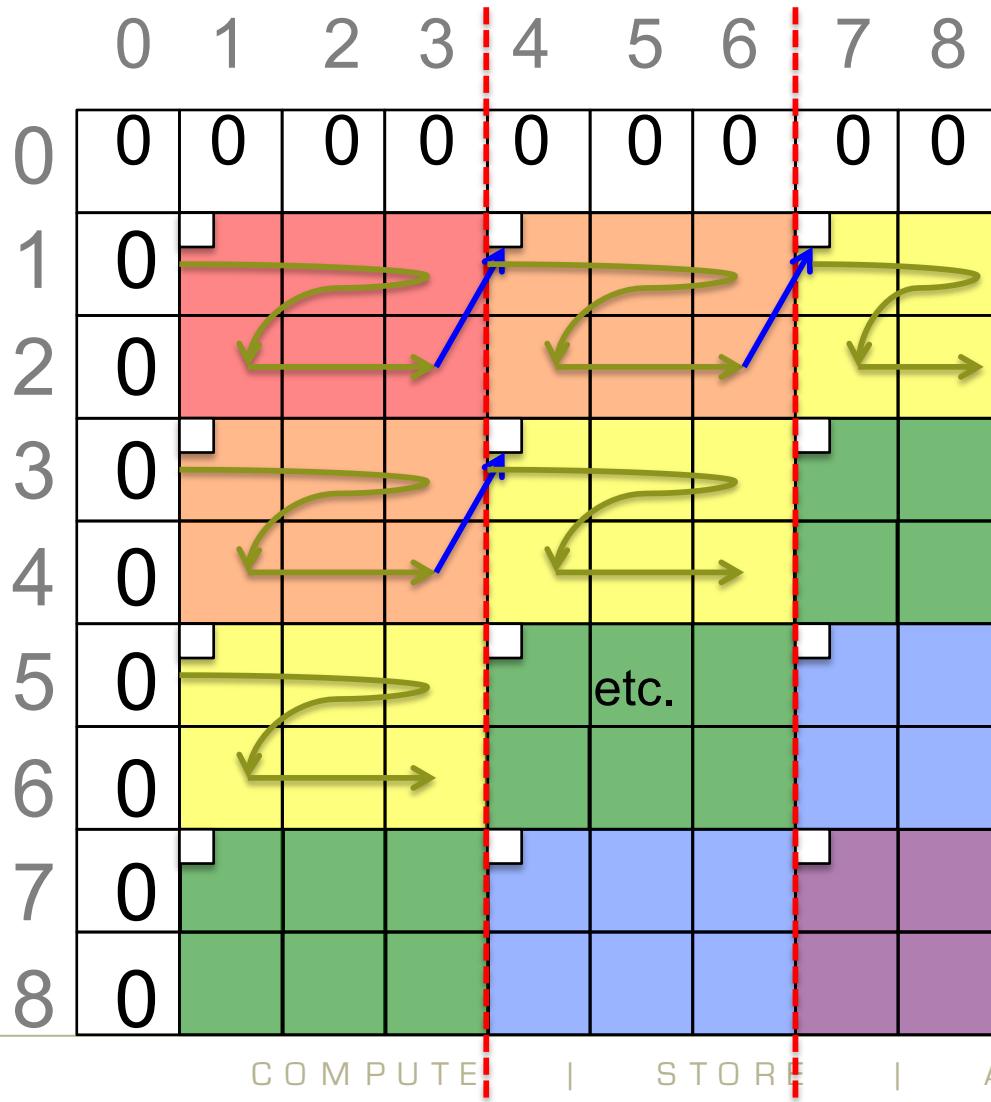
Distributed Smith-Waterman

Now, what about distributed memory?



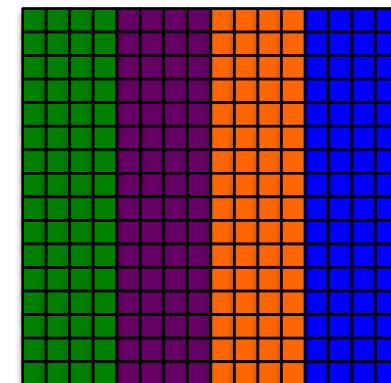
Distributed Smith-Waterman

Now, what about distributed memory?



Advantages:

- Good cache behavior: Nice fat blocks of data touchable in memory order
- Pipeline parallelism: Good utilization once pipeline is filled



Distributed Smith-Waterman

Distributed Chunked Data-Driven Task-Parallel Approach:

```

const Hspace = {0..n, 0..n};
const LocaleGrid = Locales.reshape({0..#numLocales, 0..0});
const DistHSpace = Hspace dmapped Block(Hspace, LocaleGrid);
var H: [DistHSpace] int;

proc computeH(H: [] int) {
    const ProbSpace = H.domain.translate(1,1);
    const StrProbSpace = ProbSpace by (rowsPerChunk, colsPerChunk):
    var NeighborsDone: [StrProbSpace] atomic int;
...
proc computeHHelp(x,y) {
    on H[x,y] {
        for (i,j) in ProbSpace[x..#rowsPerChunk, y..#colsPerChunk] do
            H[i,j] = f(H[i-1,j-1], H[i-1,j], H[i,j-1]);
        const eastReady = NeighborsDone[x, y+colsPerChunk].fetchAdd(1);
        ...etc...
        if (eastReady == 2) then begin computeHHelp(x, y+colsPerChunk);
        ...etc...
    }
}

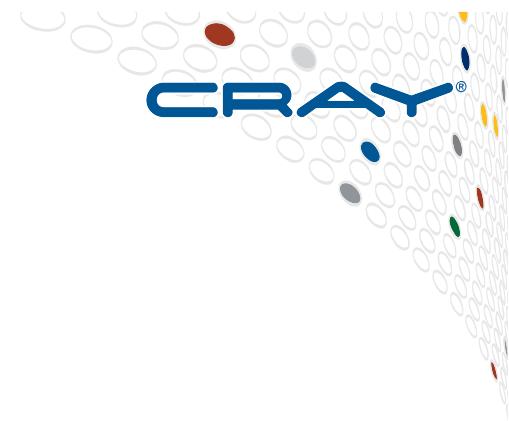
```

Reshape the 1D Locales array into a 2D column

Block-distribute the data space across the column of locales

Compute each chunk on the locale that owns its initial element





Any Questions about Smith-Waterman?



COMPUTE

|

STORE

|

ANALYZE

Copyright 2018 Cray Inc.



Legal Disclaimer

Information in this document is provided in connection with Cray Inc. products. No license, express or implied, to any intellectual property rights is granted by this document.

Cray Inc. may make changes to specifications and product descriptions at any time, without notice.

All products, dates and figures specified are preliminary based on current expectations, and are subject to change without notice.

Cray hardware and software products may contain design defects or errors known as errata, which may cause the product to deviate from published specifications. Current characterized errata are available on request.

Cray uses codenames internally to identify products that are in development and not yet publically announced for release. Customers and other third parties are not authorized by Cray Inc. to use codenames in advertising, promotion or marketing and any use of Cray Inc. internal codenames is at the sole risk of the user.

Performance tests and ratings are measured using specific systems and/or components and reflect the approximate performance of Cray Inc. products as measured by those tests. Any difference in system hardware or software design or configuration may affect actual performance.

The following are trademarks of Cray Inc. and are registered in the United States and other countries: CRAY and design, SONEXION, and URIKA. The following are trademarks of Cray Inc.: ACE, APPRENTICE2, CHAPEL, CLUSTER CONNECT, CRAYPAT, CRAYPORT, ECOPHLEX, LIBSCI, NODEKARE, THREADSTORM. The following system family marks, and associated model number marks, are trademarks of Cray Inc.: CS, CX, XC, XE, XK, XMT, and XT. The registered trademark LINUX is used pursuant to a sublicense from LMI, the exclusive licensee of Linus Torvalds, owner of the mark on a worldwide basis. Other trademarks used in this document are the property of their respective owners.

